

Understanding the Foundations of TensorFlow

INTRODUCING TENSORFLOW

Overview

Introduce TensorFlow(TF), a language for numerical computations

Understand the basics of machine learning, deep learning and neural networks

Learn why TF is slowly becoming the default library for ML

Install and set up TensorFlow on your local machine

What You Need in Your Toolkit

Prerequisites



Familiarity with the command line on a Mac, Linux or Windows machine

Comfortable with writing programs in Python

Install and Setup



The latest version of TensorFlow 1.2rc0

A compatible version of Python, version 2.7 and 3.x

- Only works with versions >3.5 on Windows

A Mac, Linux or Windows machine on which TensorFlow can be installed



Course Overview

Introduction to TensorFlow, install and set up

Basics of TensorFlow, computation graphs, tensors, sessions and TensorBoard

Fundamentals of TensorFlow, placeholders, variables, the feed dictionary

Working with images, representing RGB and grayscale images, image operations

Machine Learning with TensorFlow, identifying handwritten digits in the MNIST dataset using the nearest neighbors algorithm

Understanding Machine Learning


Machine Learning



Work with a huge maze
of data



Find patterns



Make intelligent
decisions

A machine learning algorithm is an algorithm that is able to learn from data

Machine Learning



Emails on a server




Spam or Ham?



Trash or Inbox

Machine Learning



Images represented as
pixels



Identify edges, colors,
shapes



A photo of a little
bird


Types of Machine Learning Problems



Classification



Regression



Clustering



Rule-extraction

Types of Machine Learning Problems



Classification



Regression




Clustering




Rule-extraction

Whales: Fish or Mammals?



Mammals

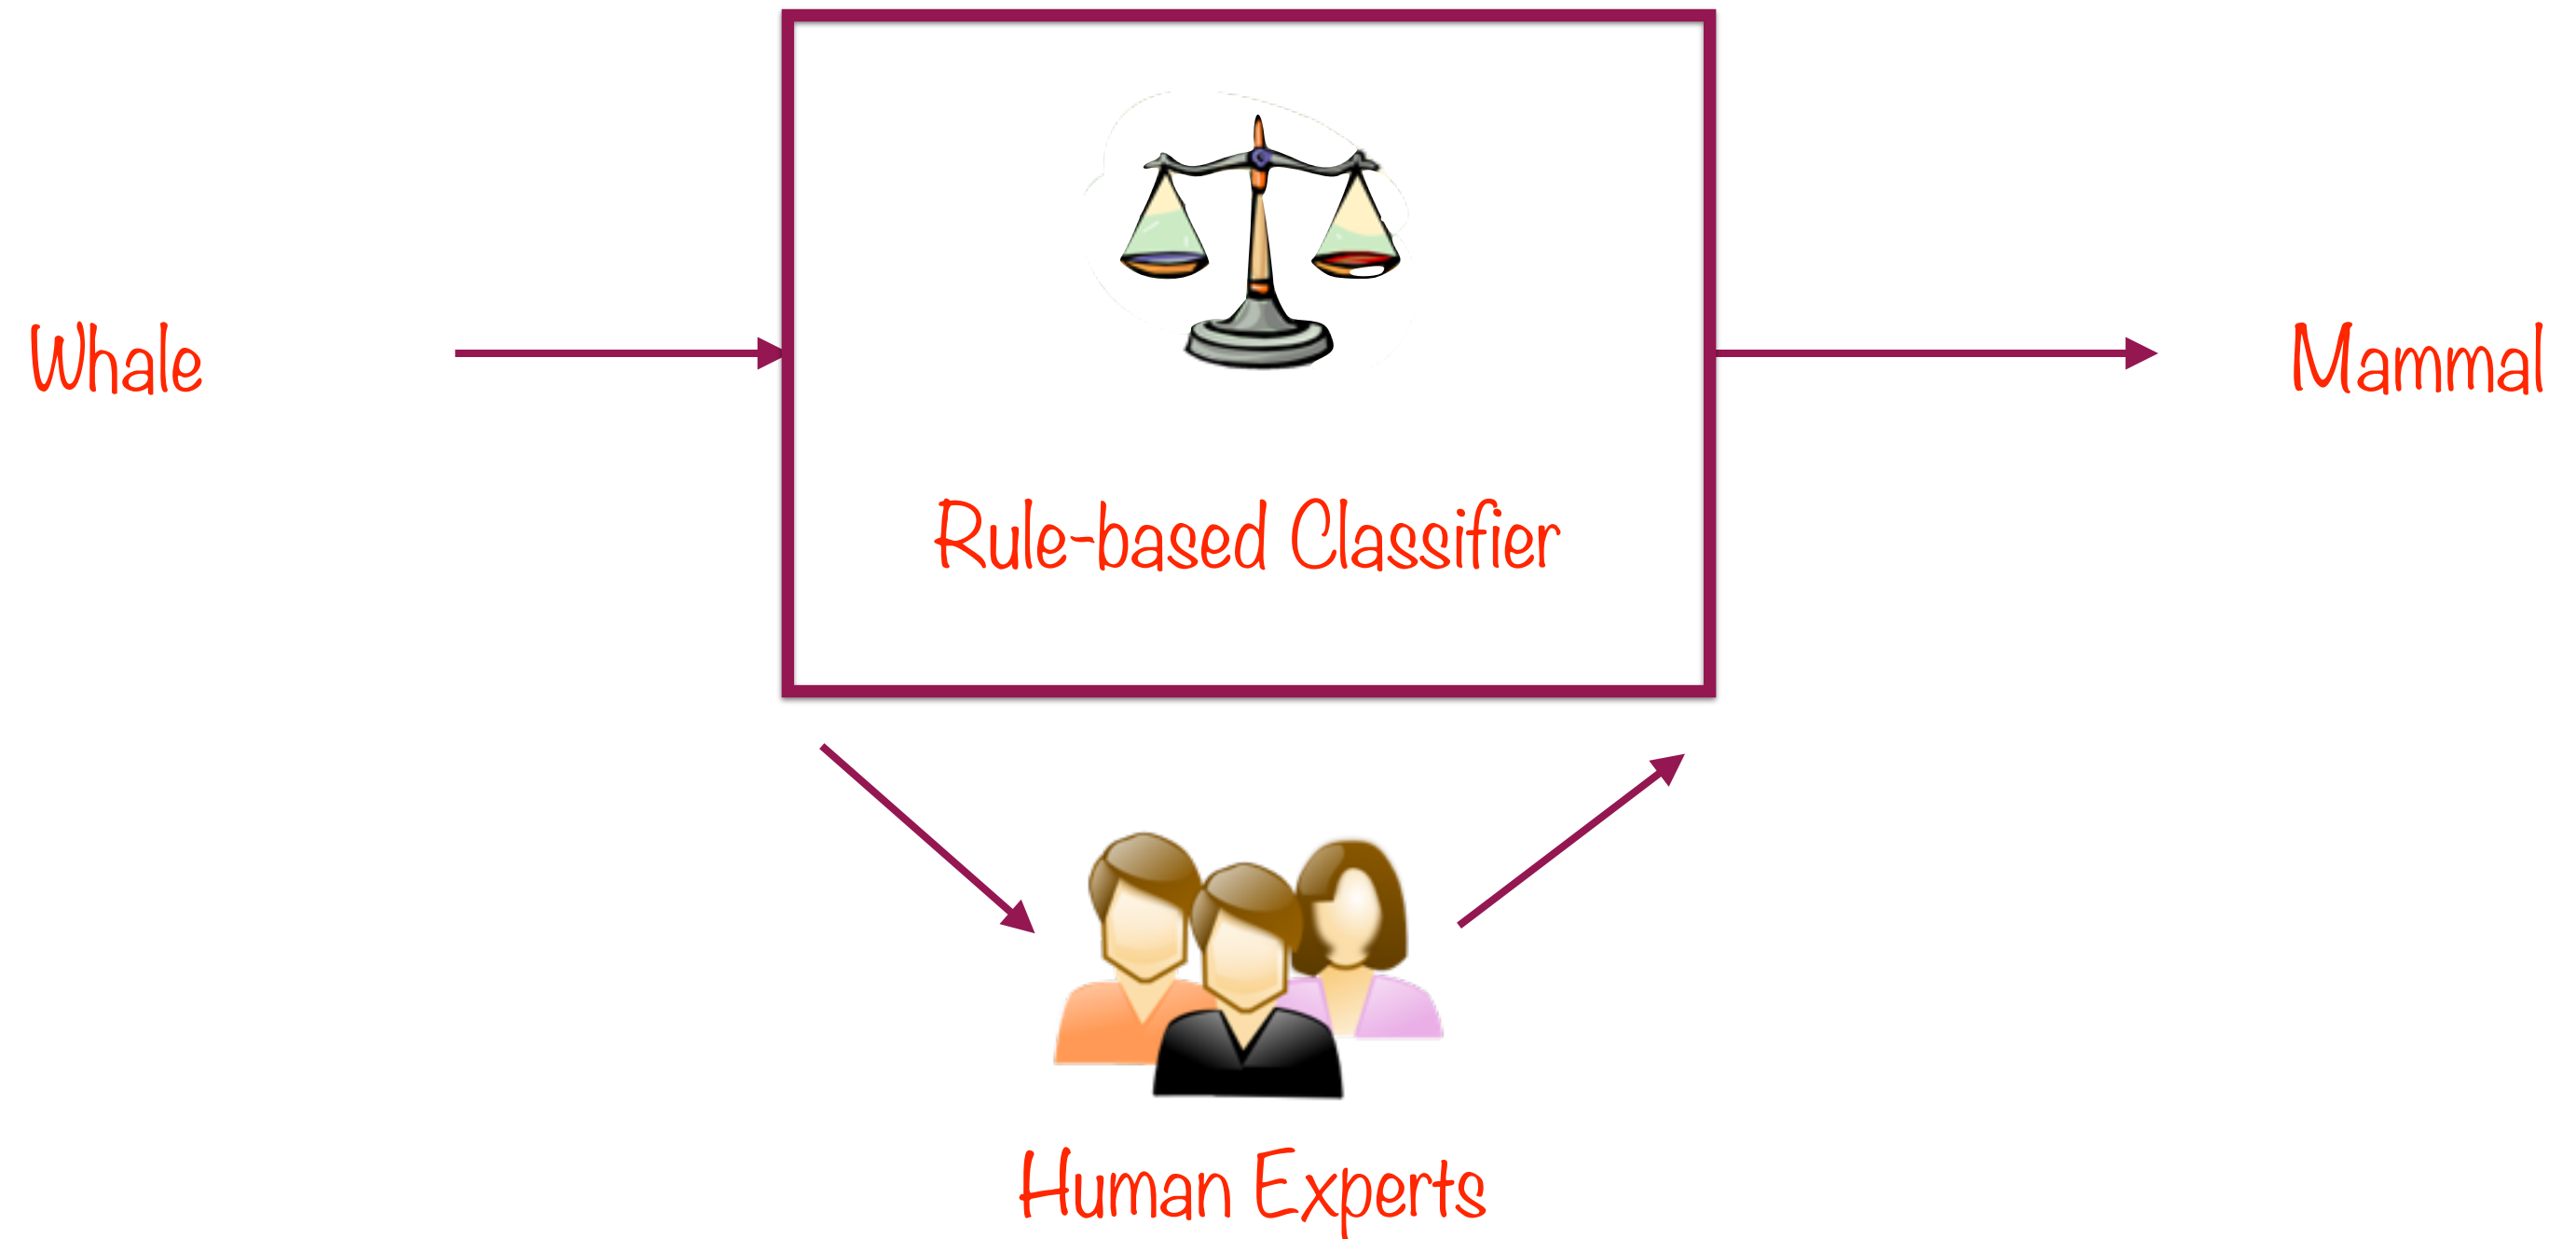
Members of the infraorder Cetacea



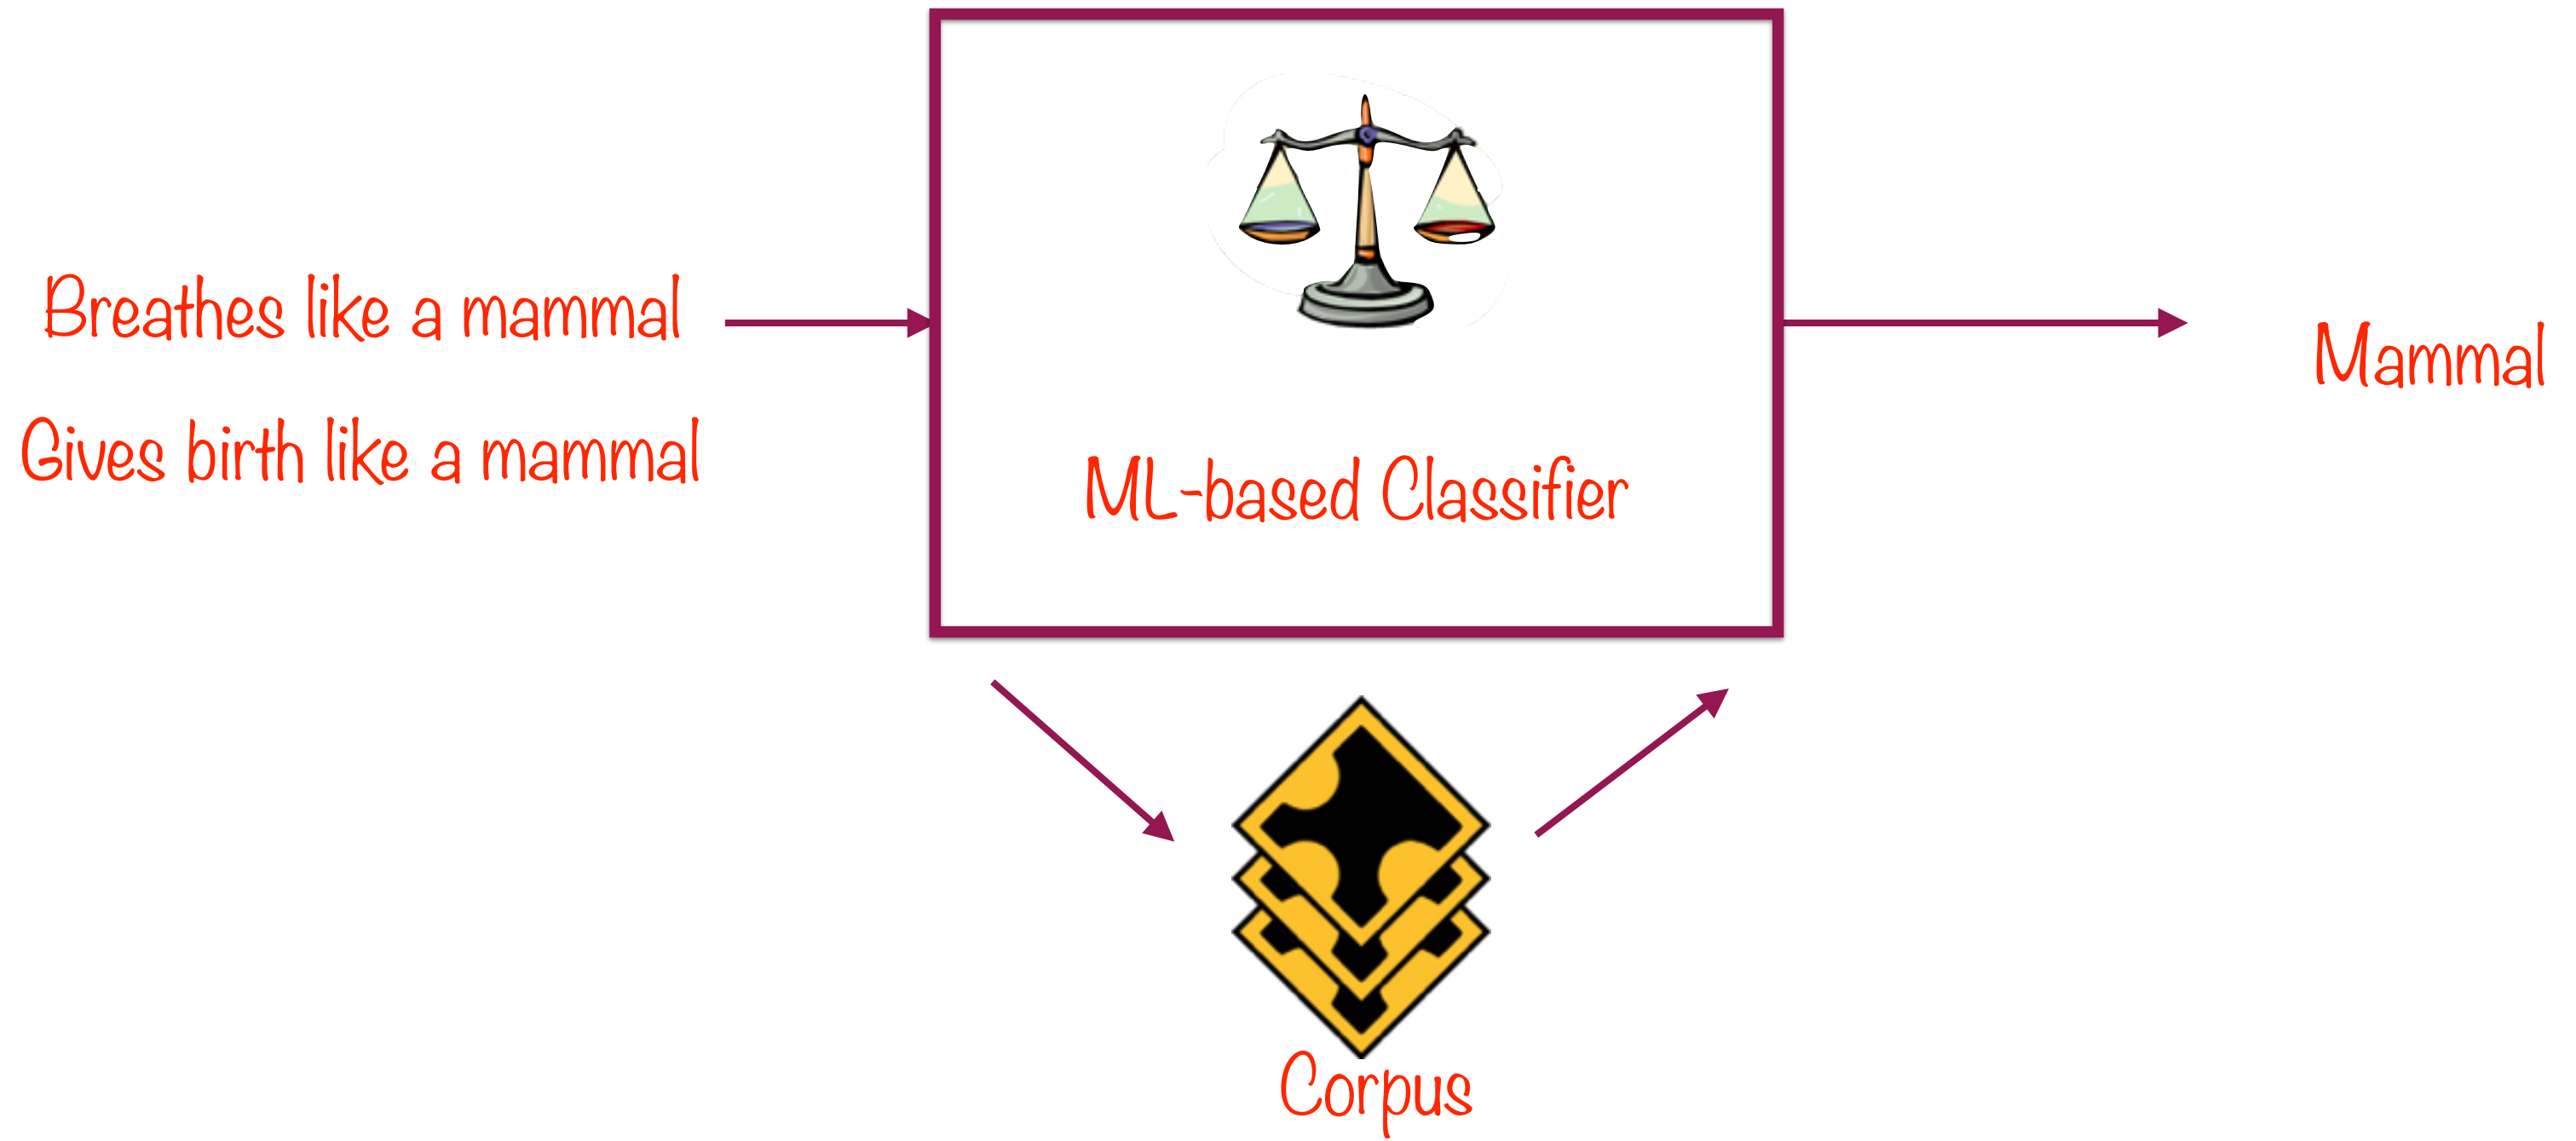
Fish

Look like fish, swim like fish, move with fish

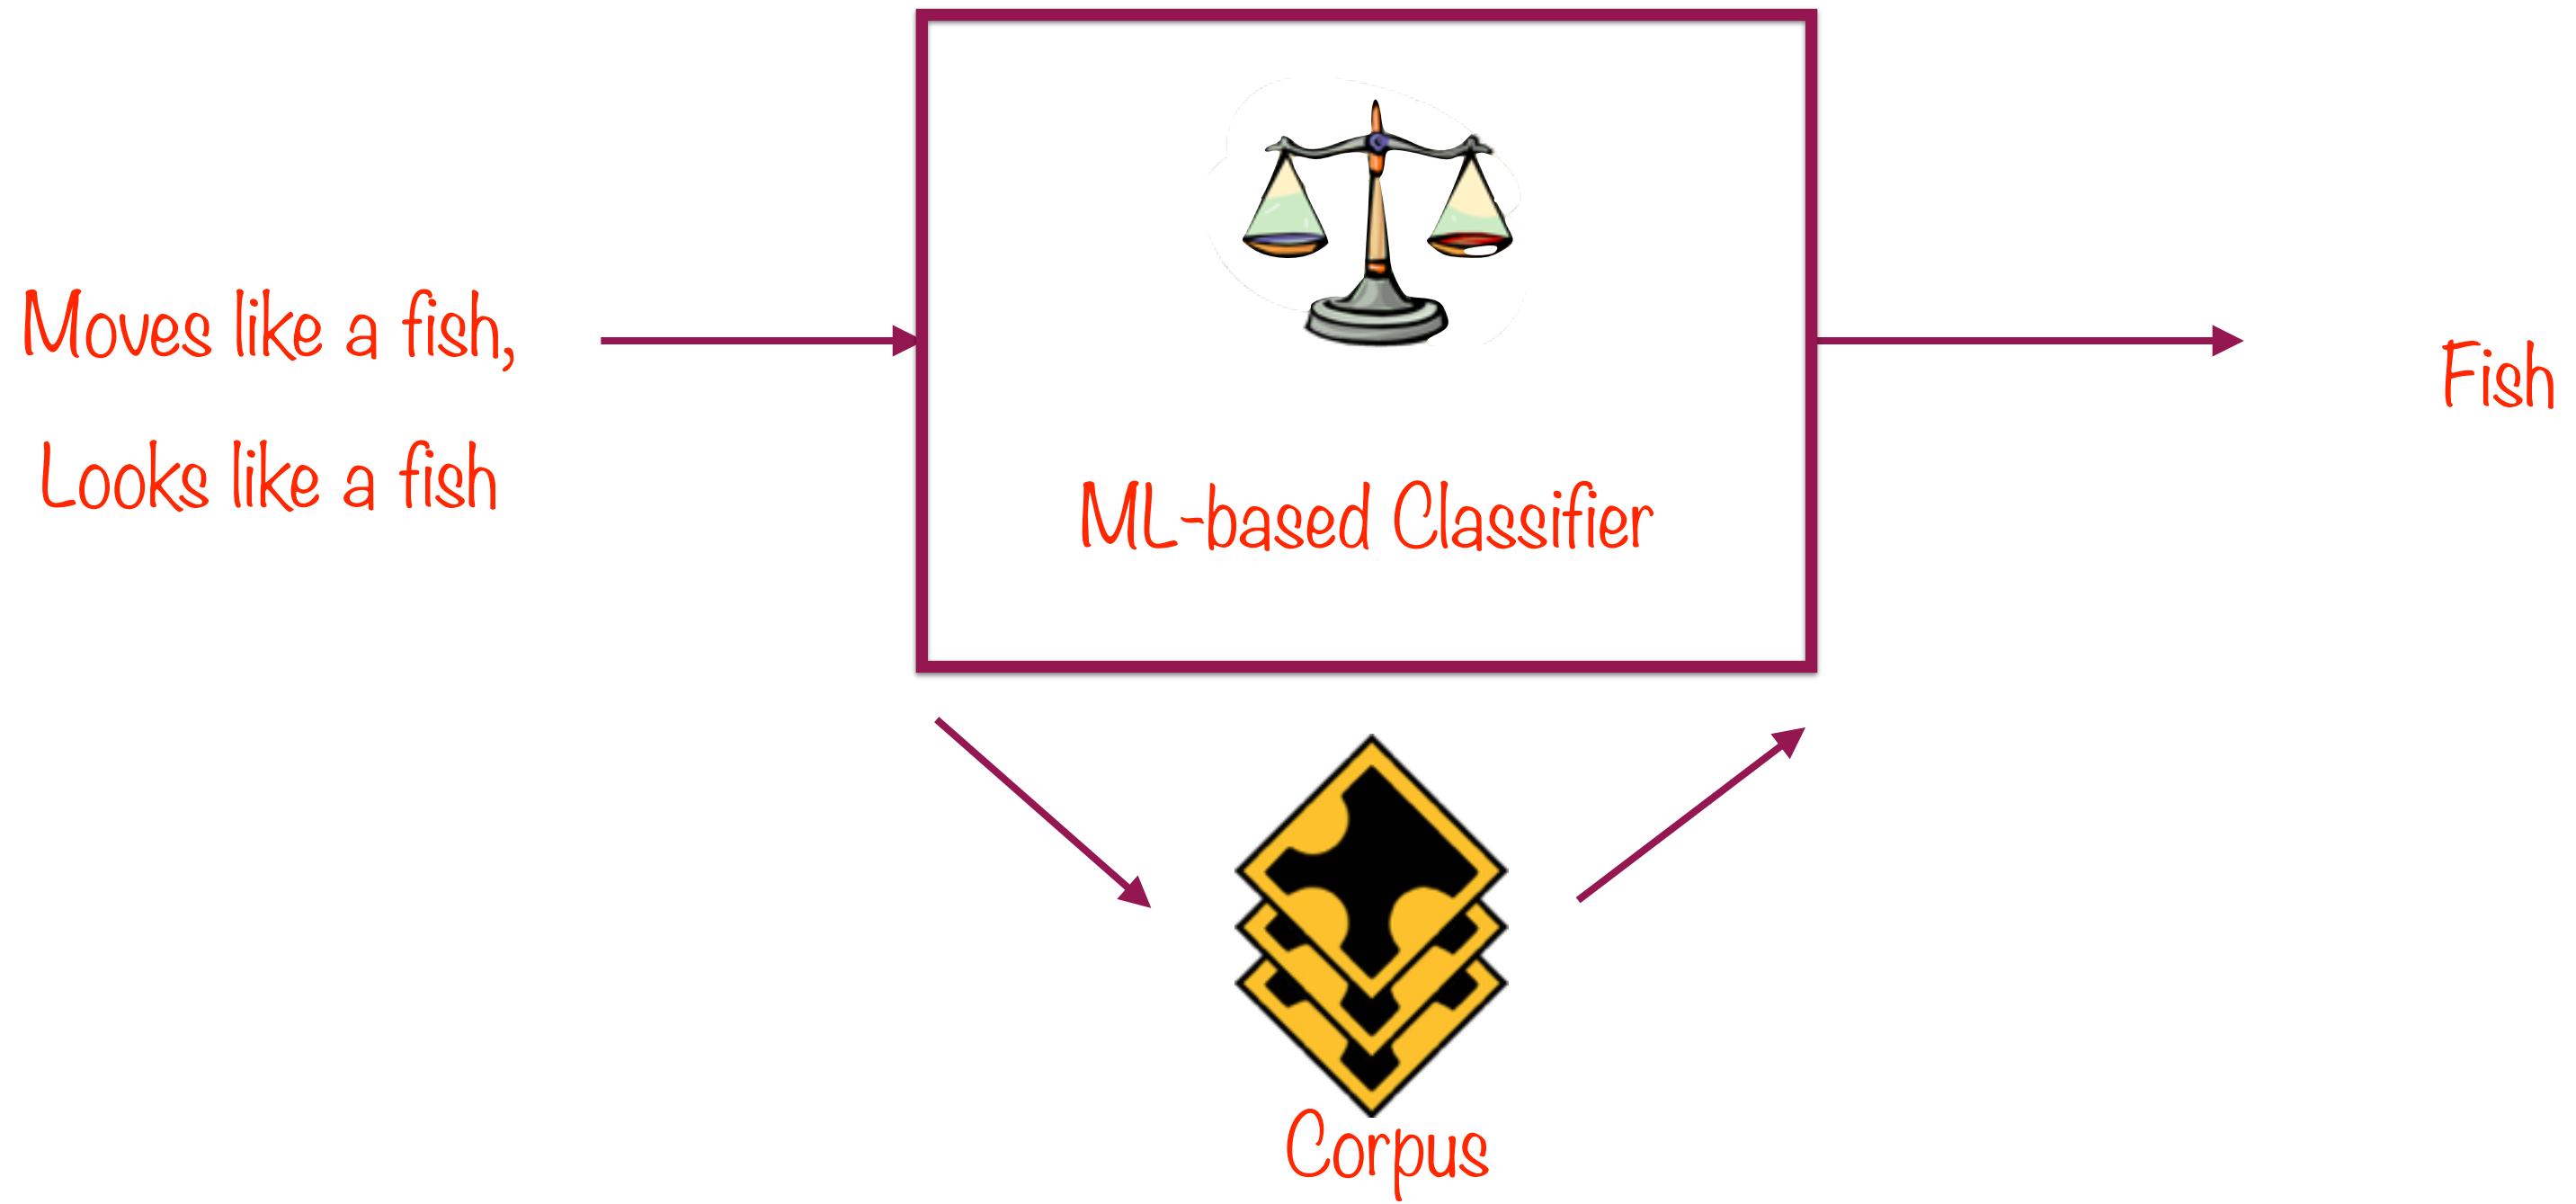
Rule-based Binary Classifier





ML-based Binary Classifier




“Traditional” ML-based Binary Classifier



ML-based Binary Classifier



Corpus



Classification Algorithm



ML-based Classifier

ML-based Binary Classifier

ML-based

Dynamic

Experts optional

Corpus required

Training step

Rule-based

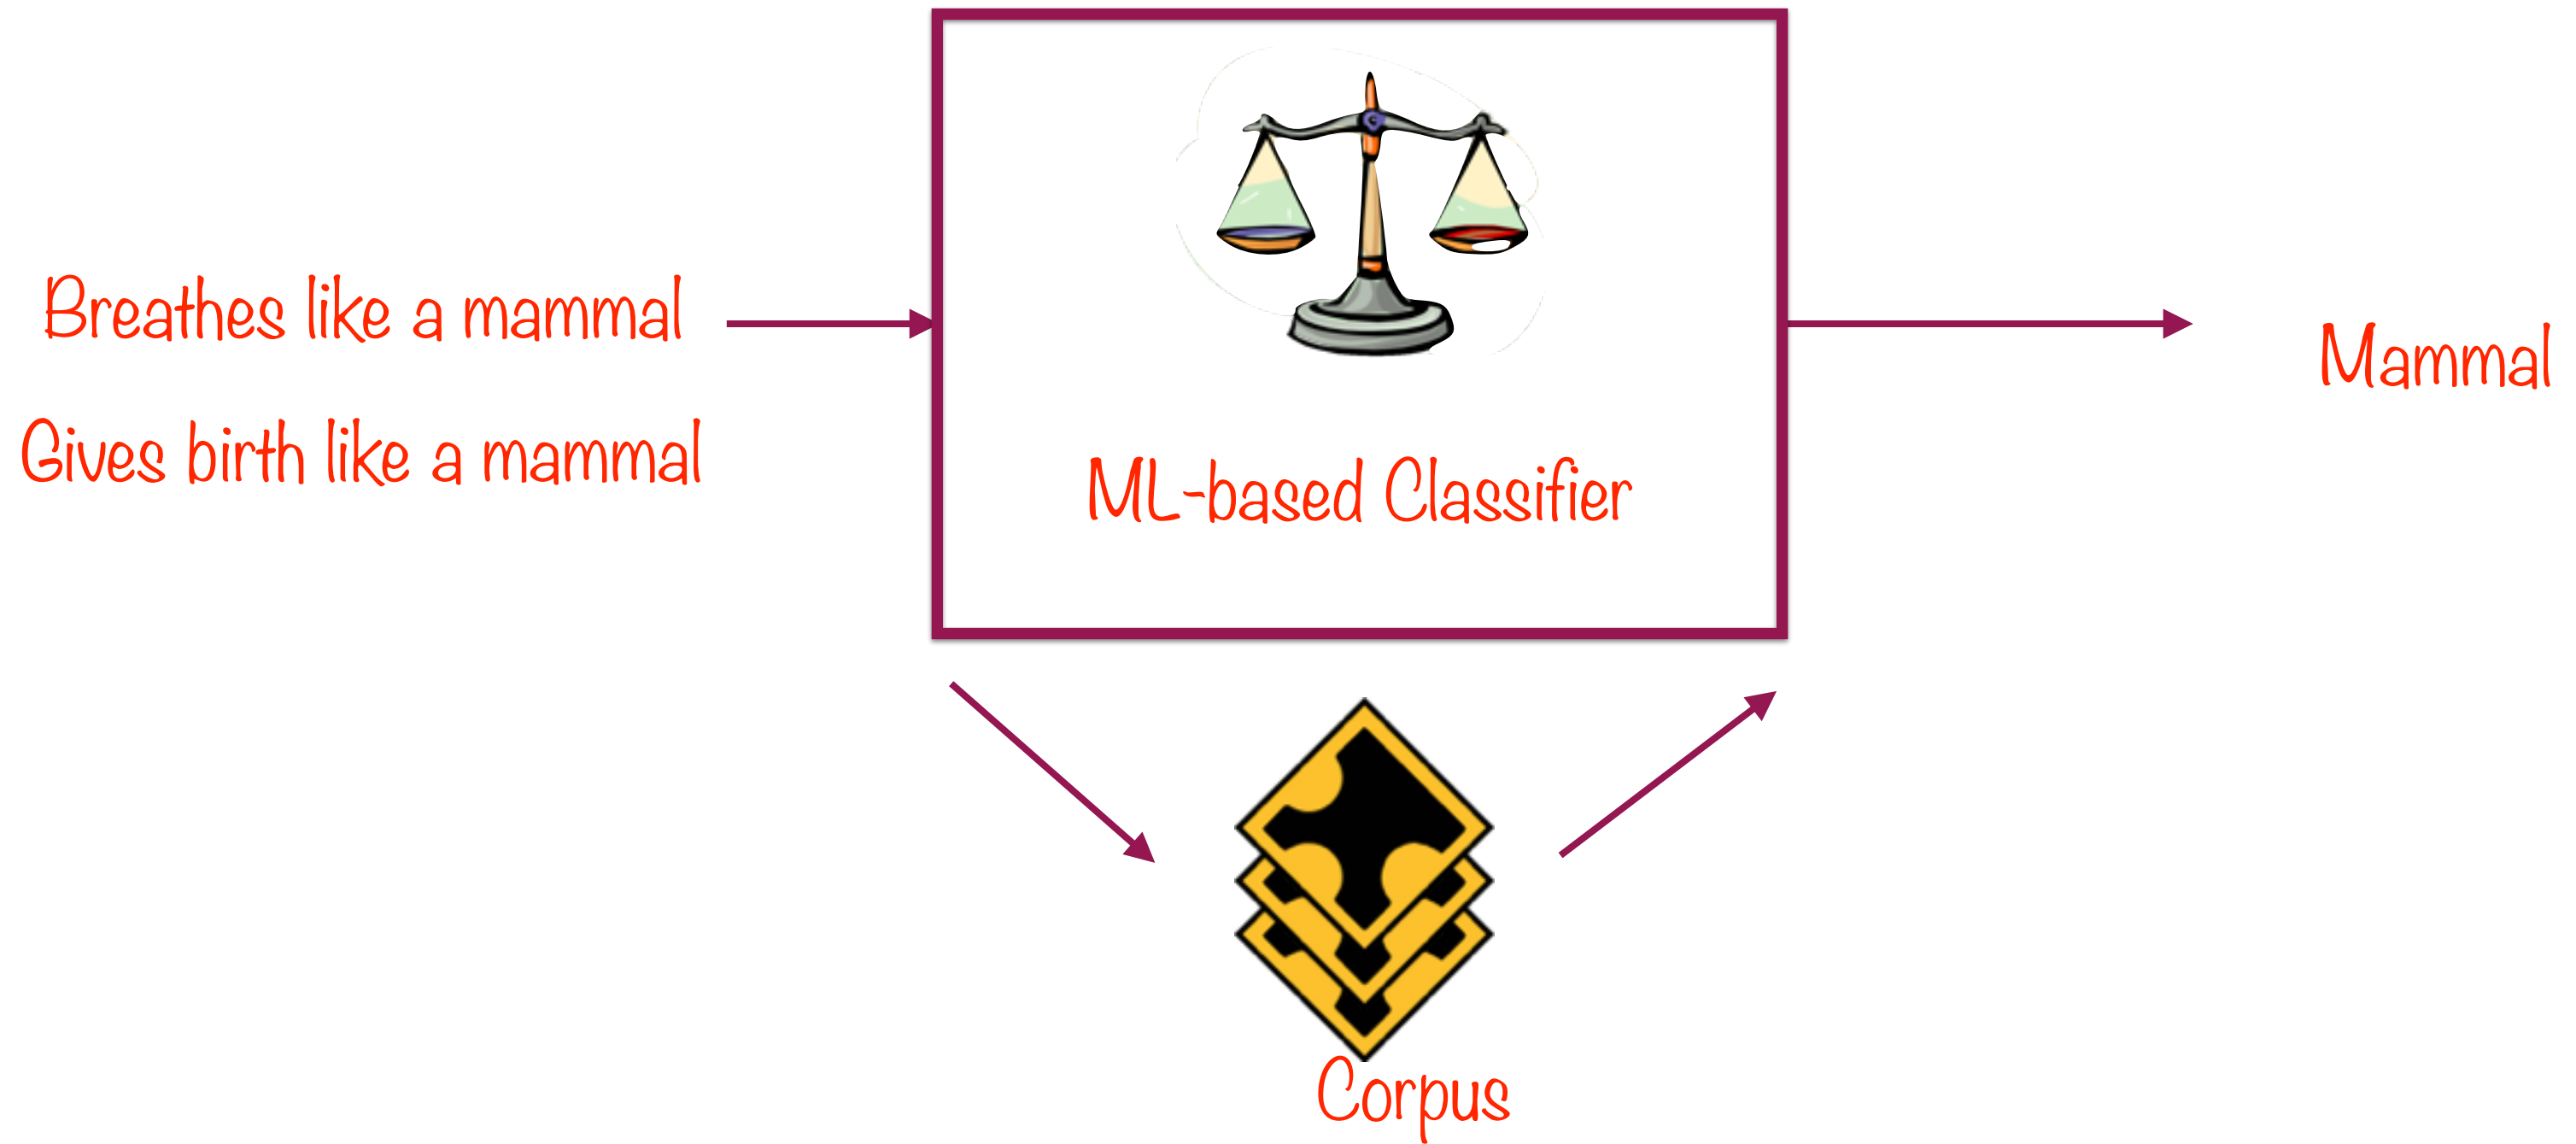
Static

Experts required

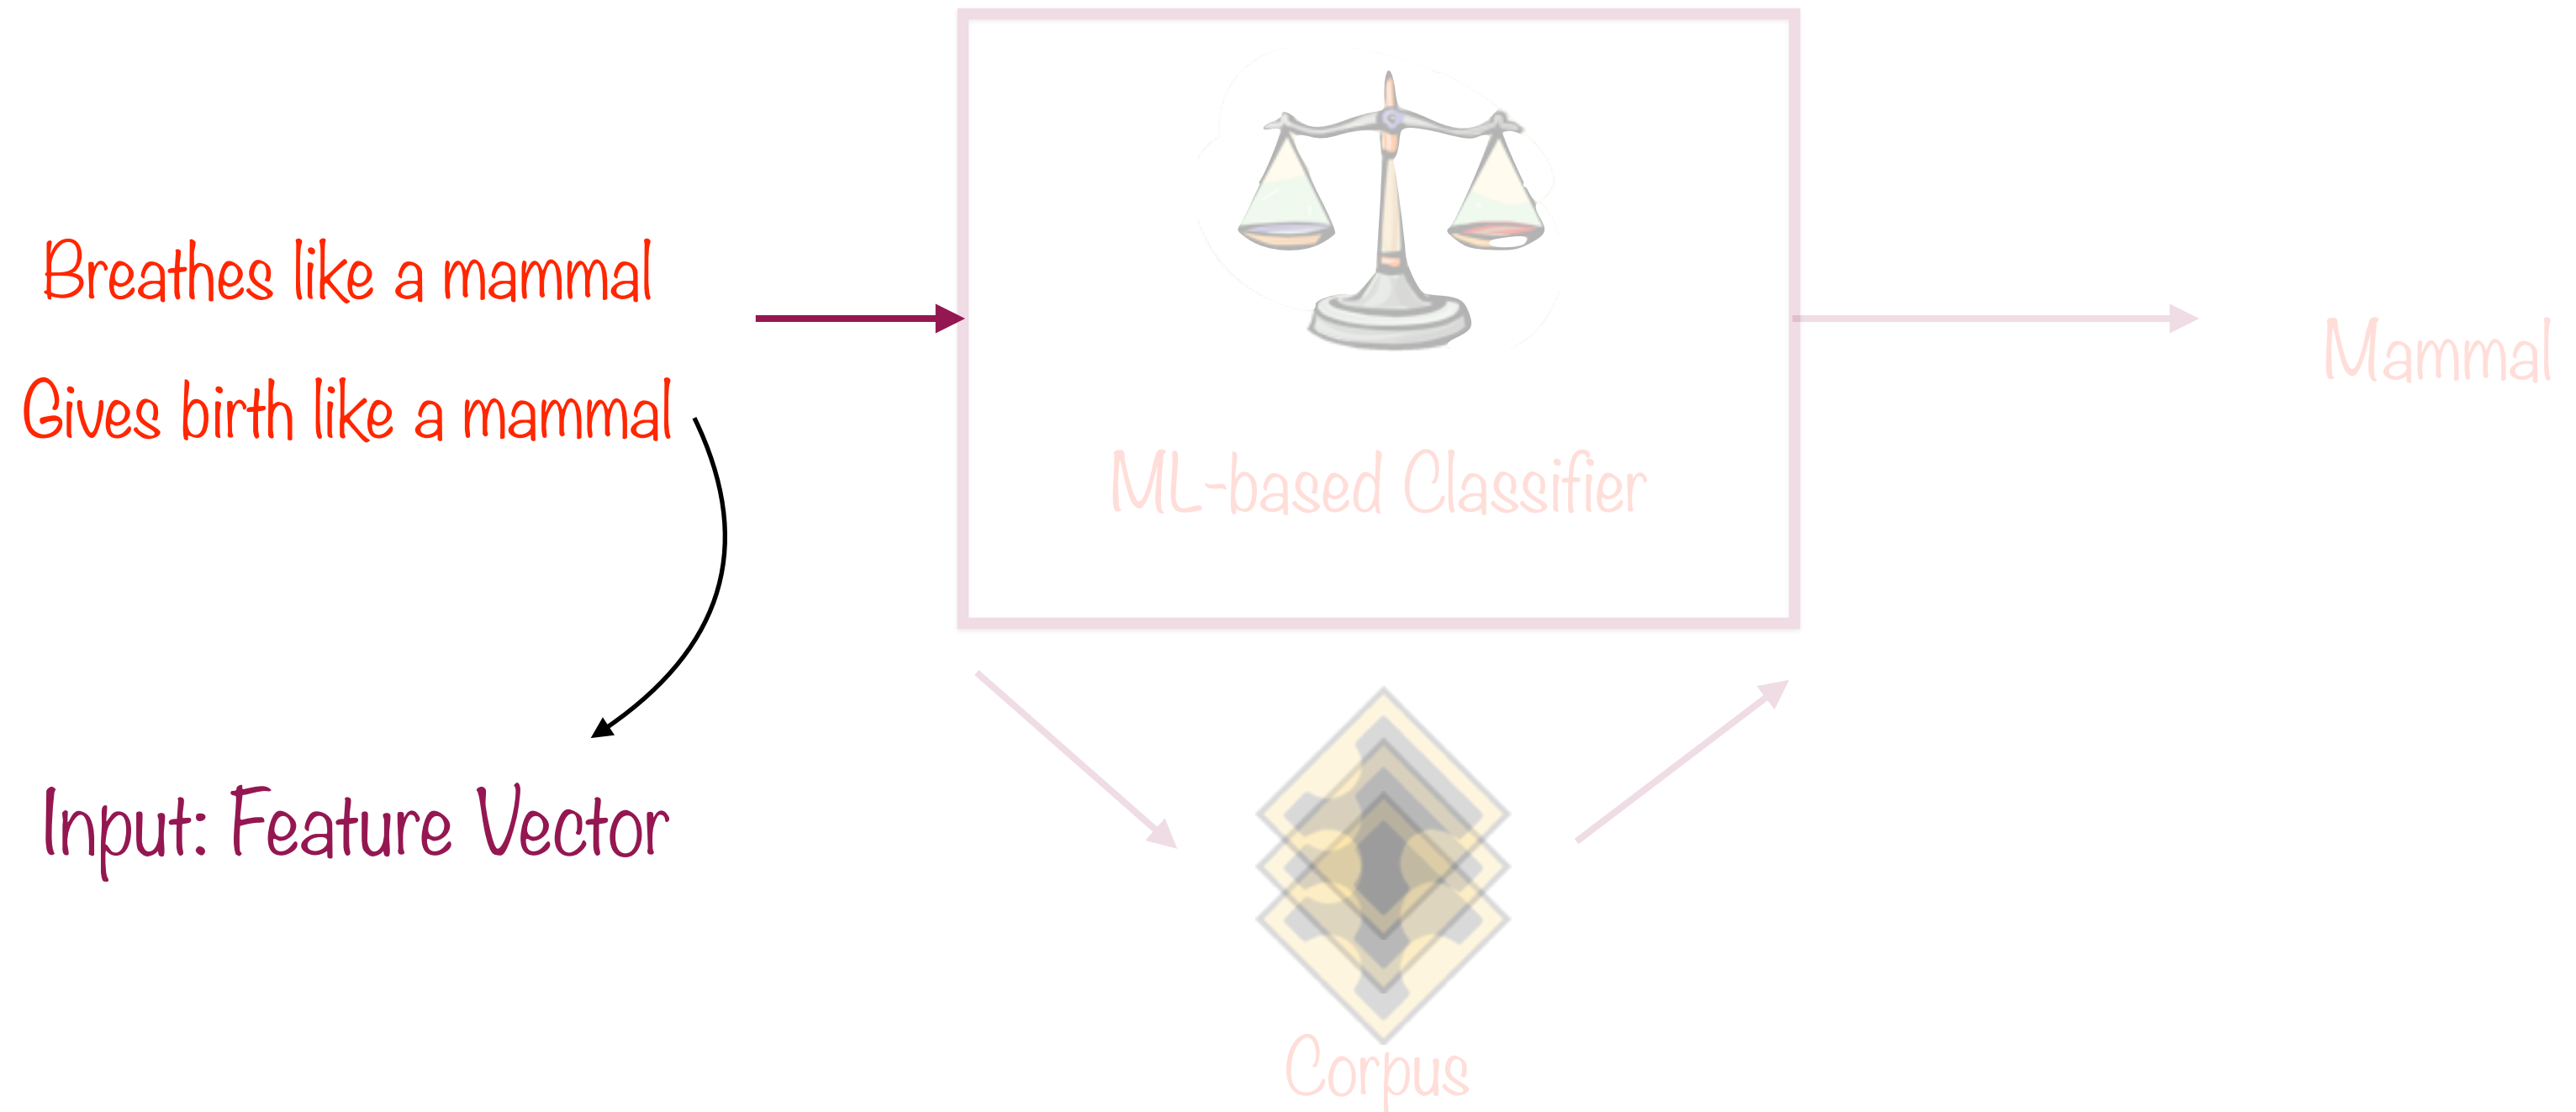
Corpus optional

No training step

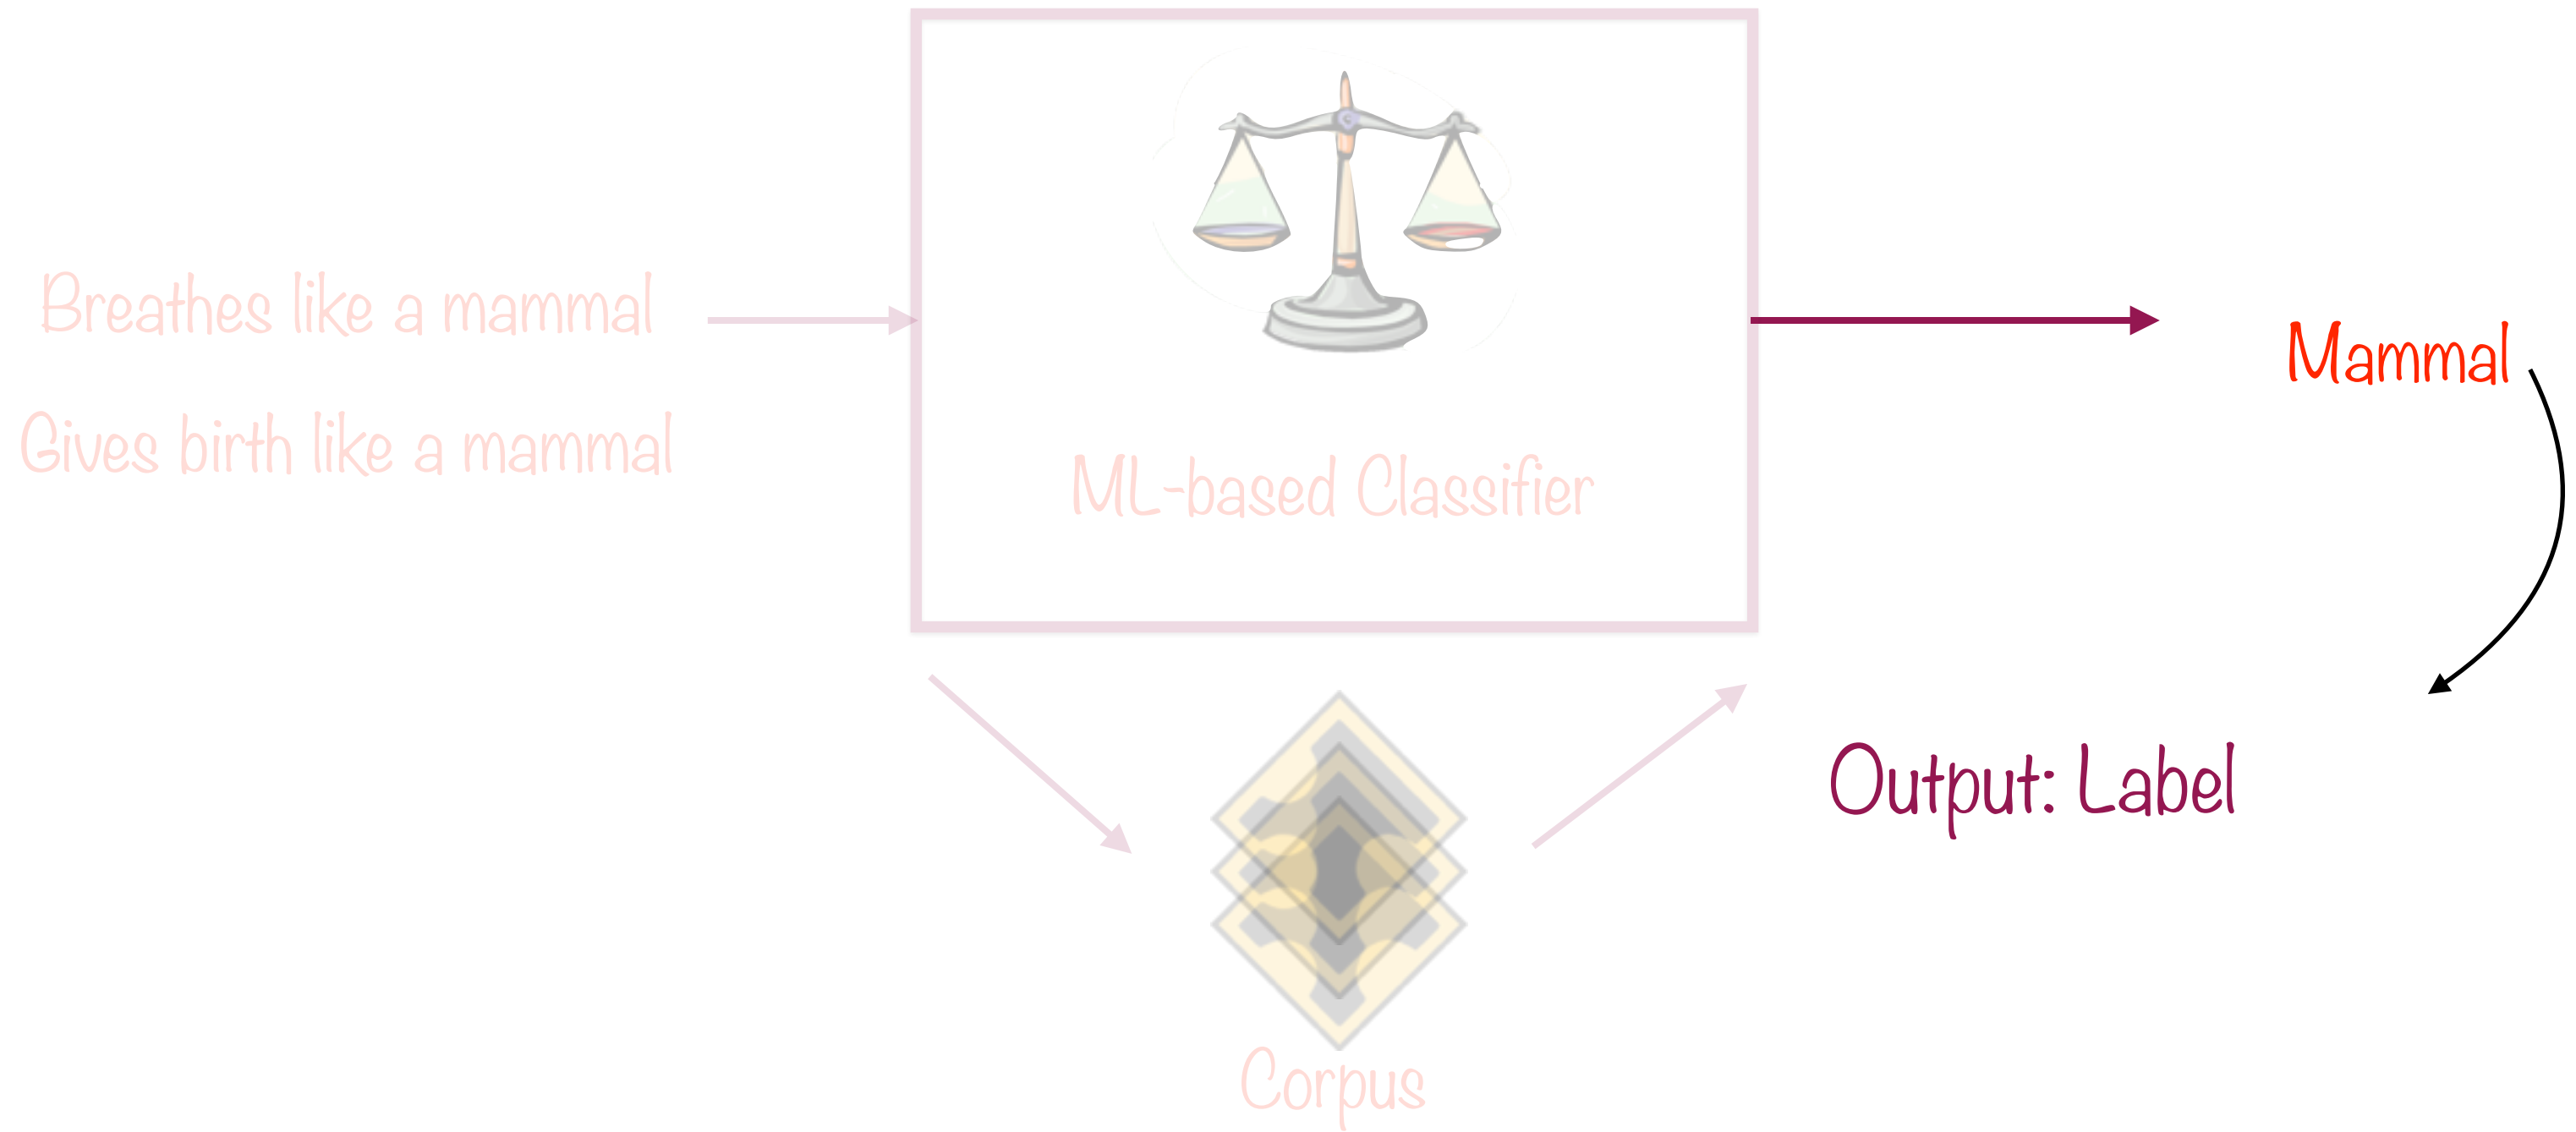
“Traditional” ML-based Binary Classifier



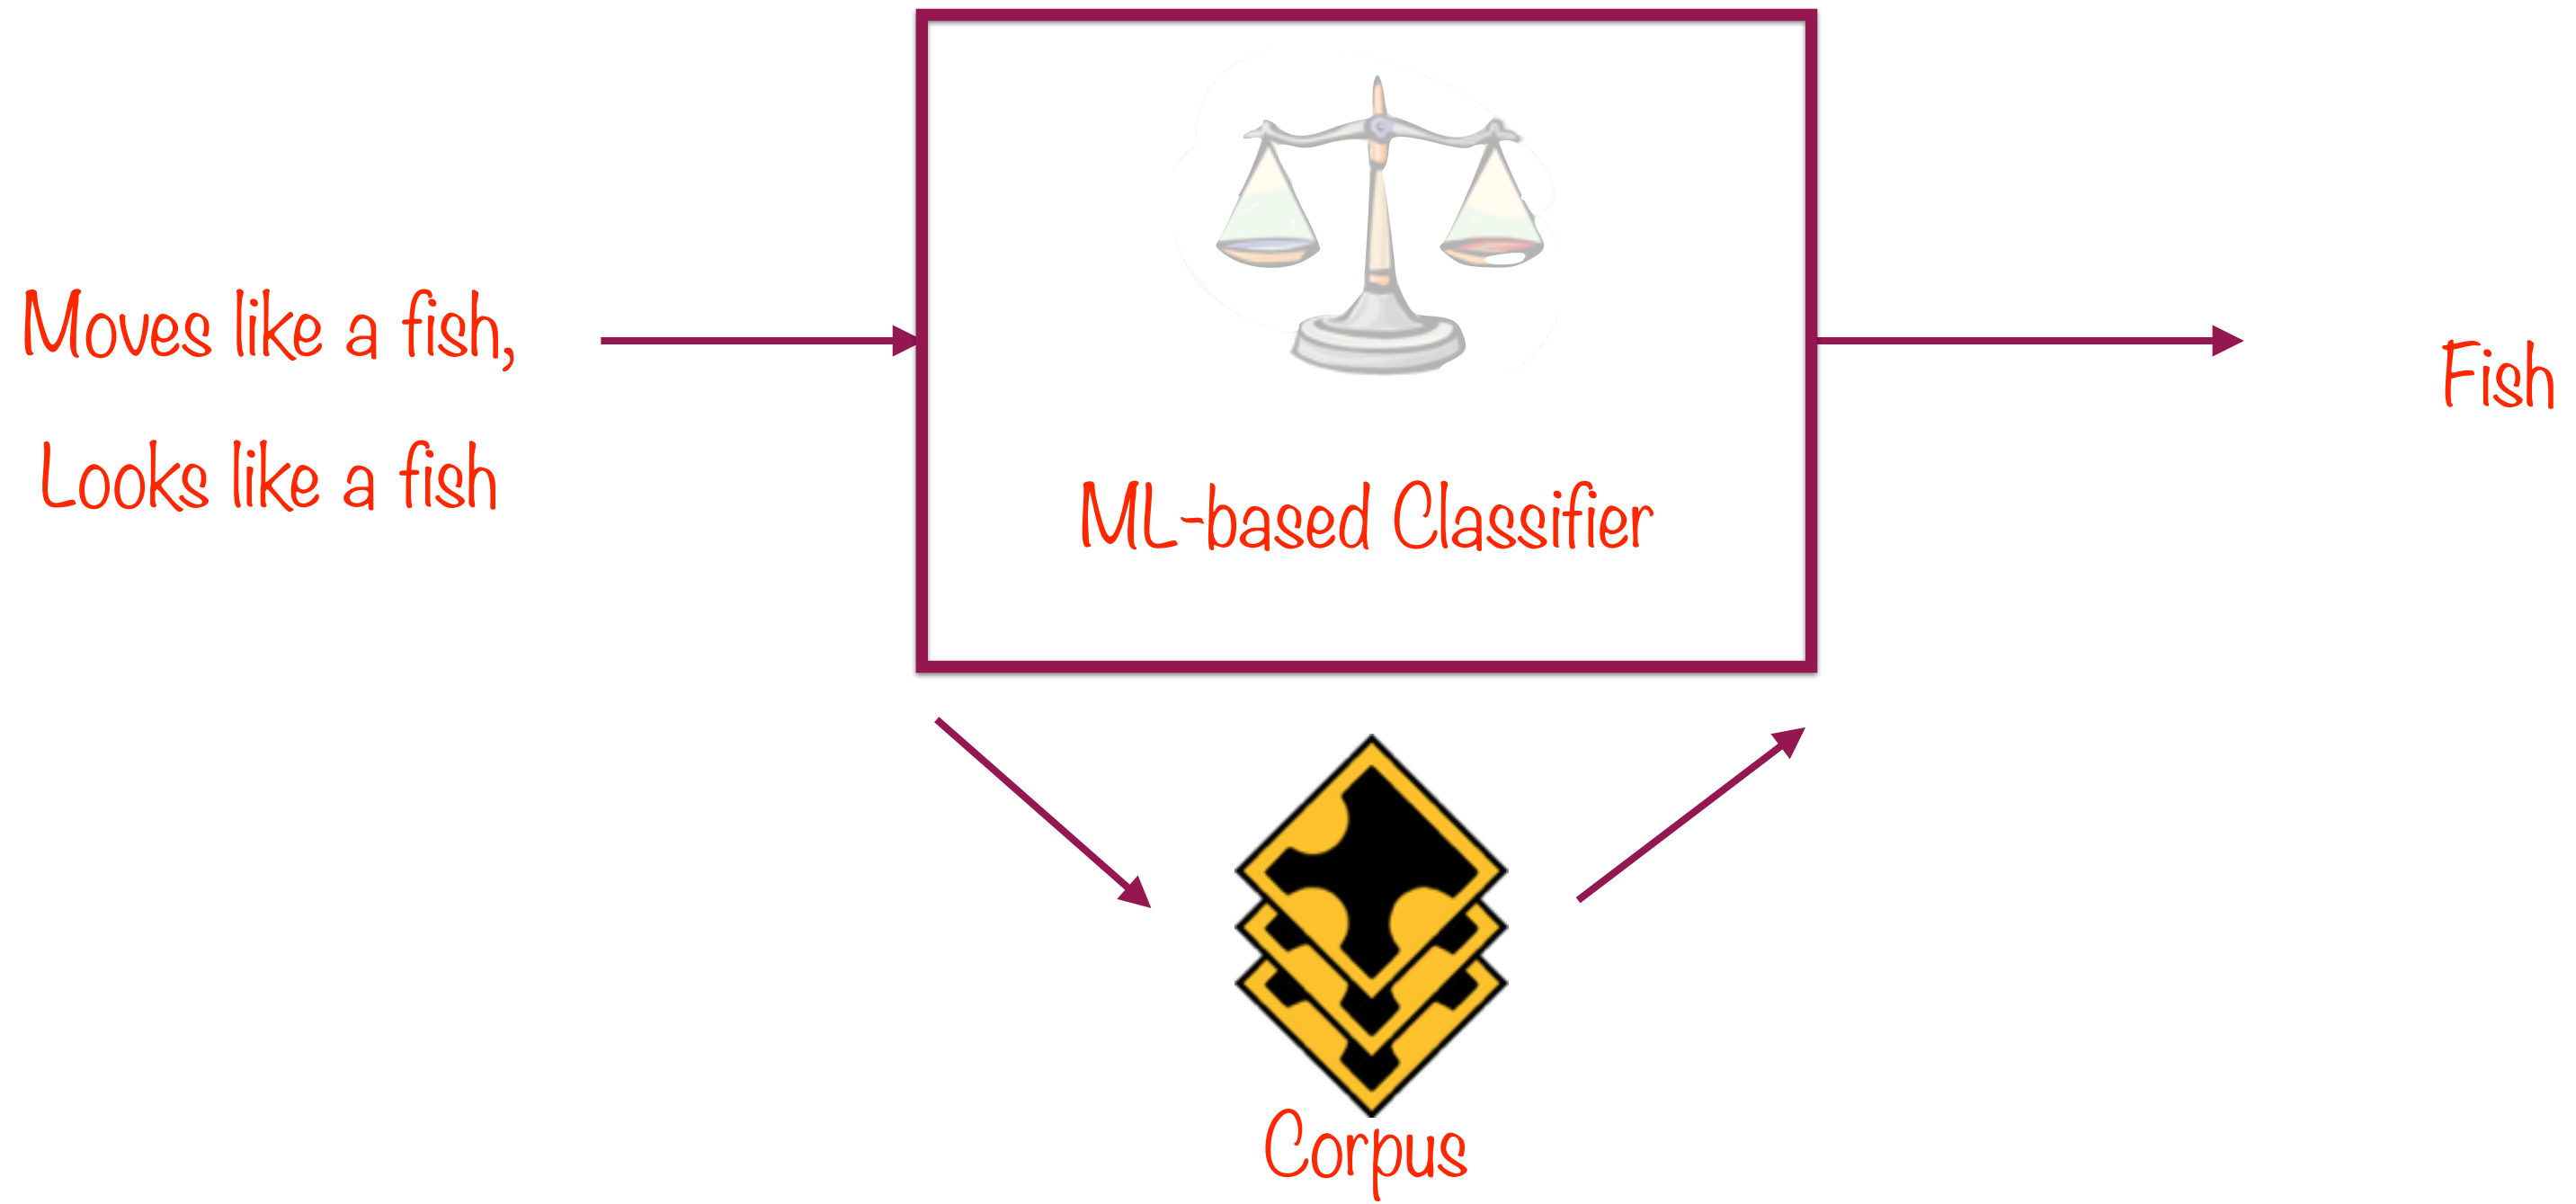
“Traditional” ML-based Binary Classifier



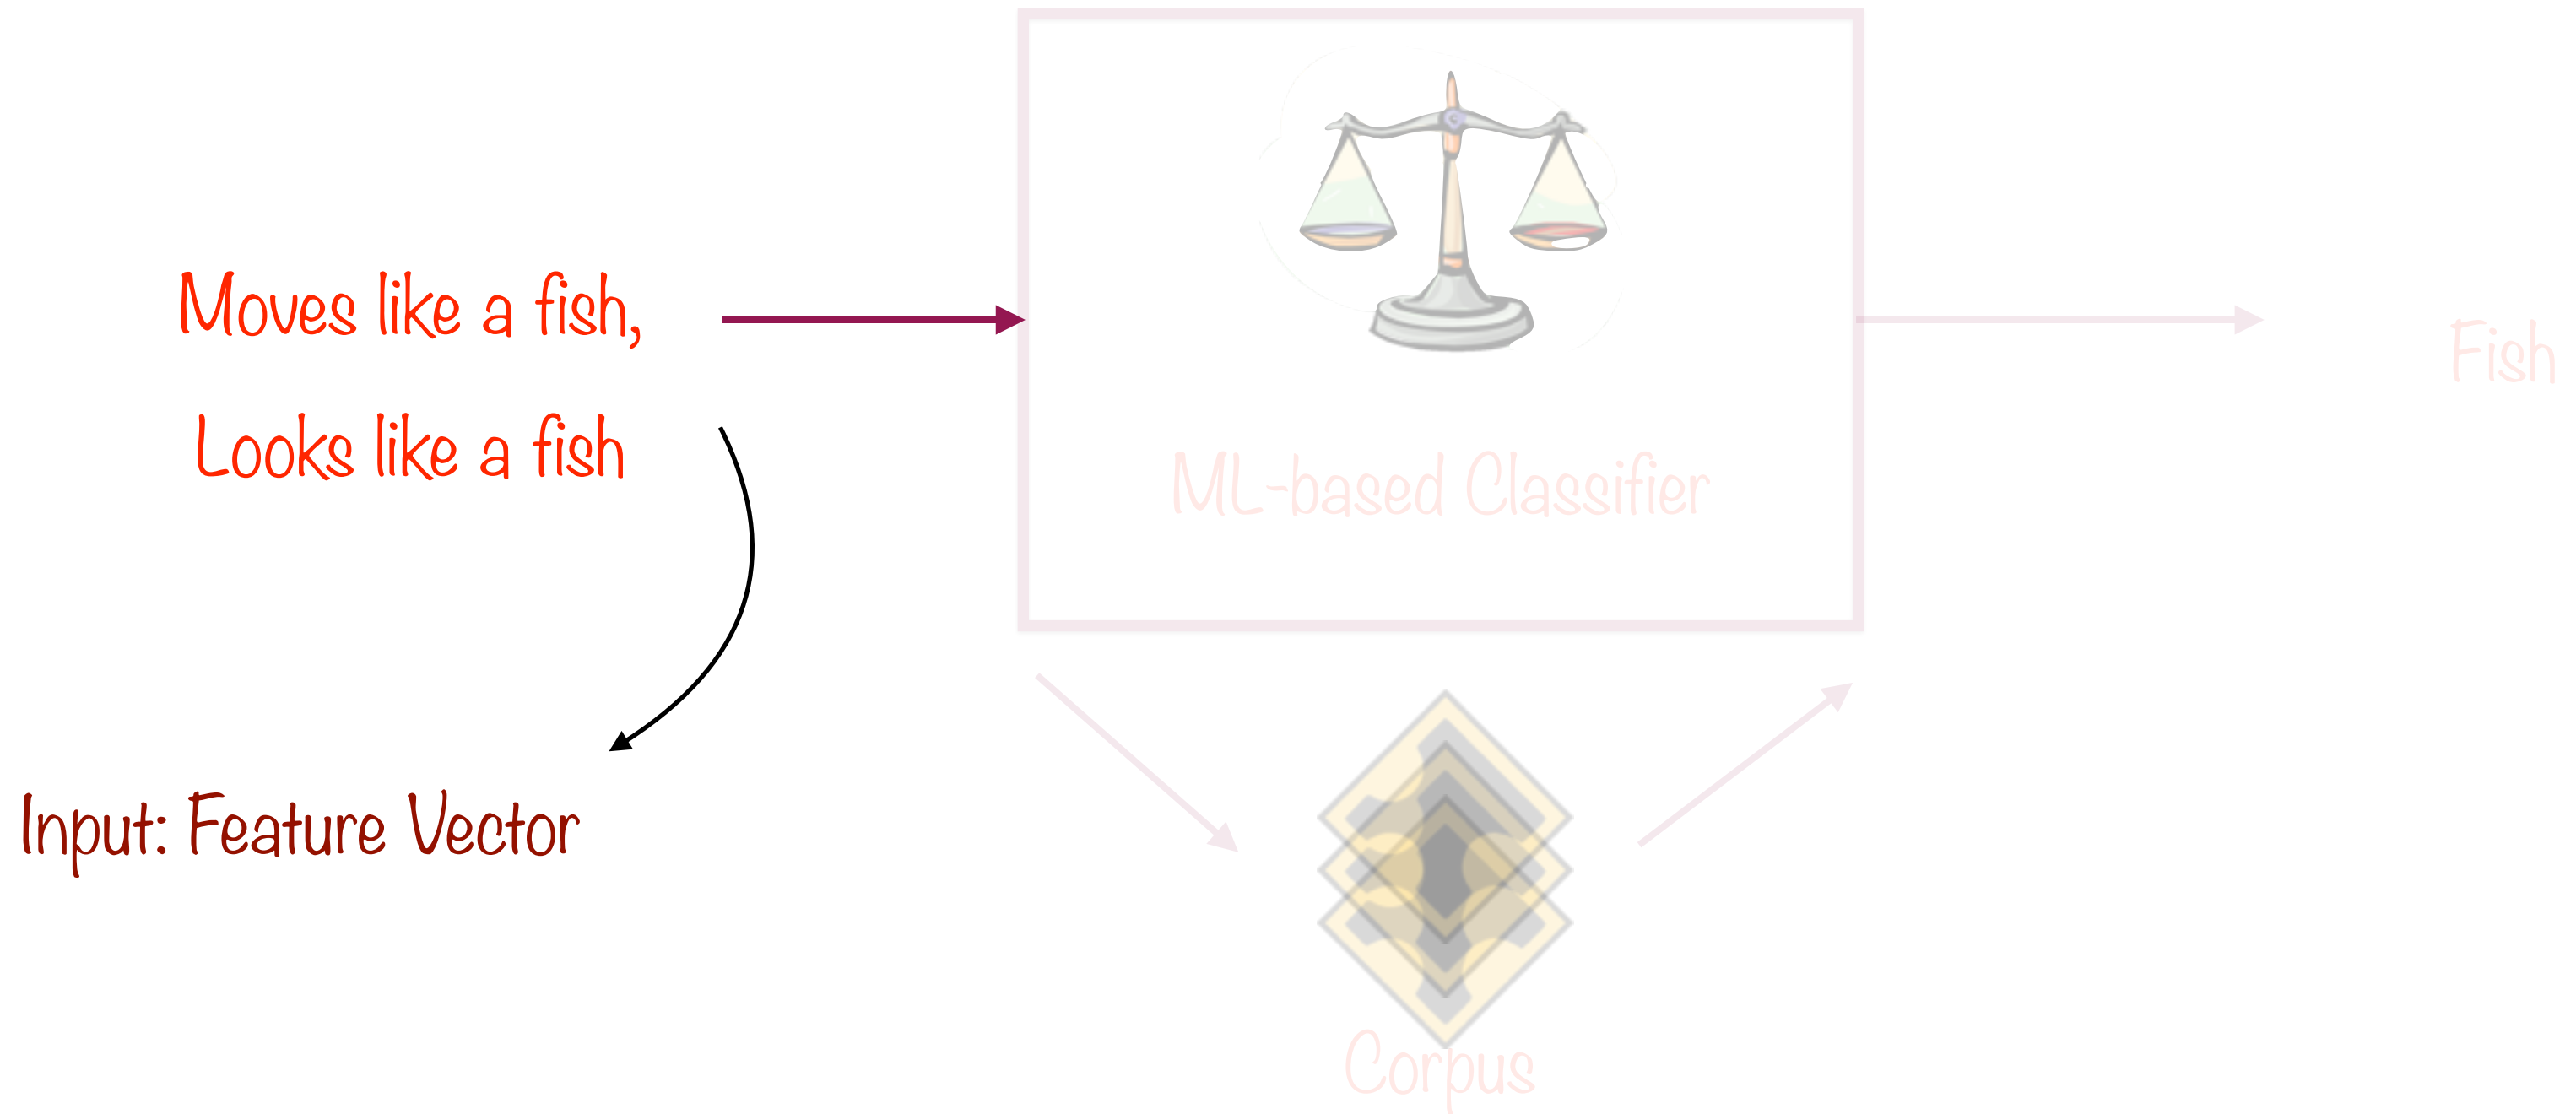
“Traditional” ML-based Binary Classifier



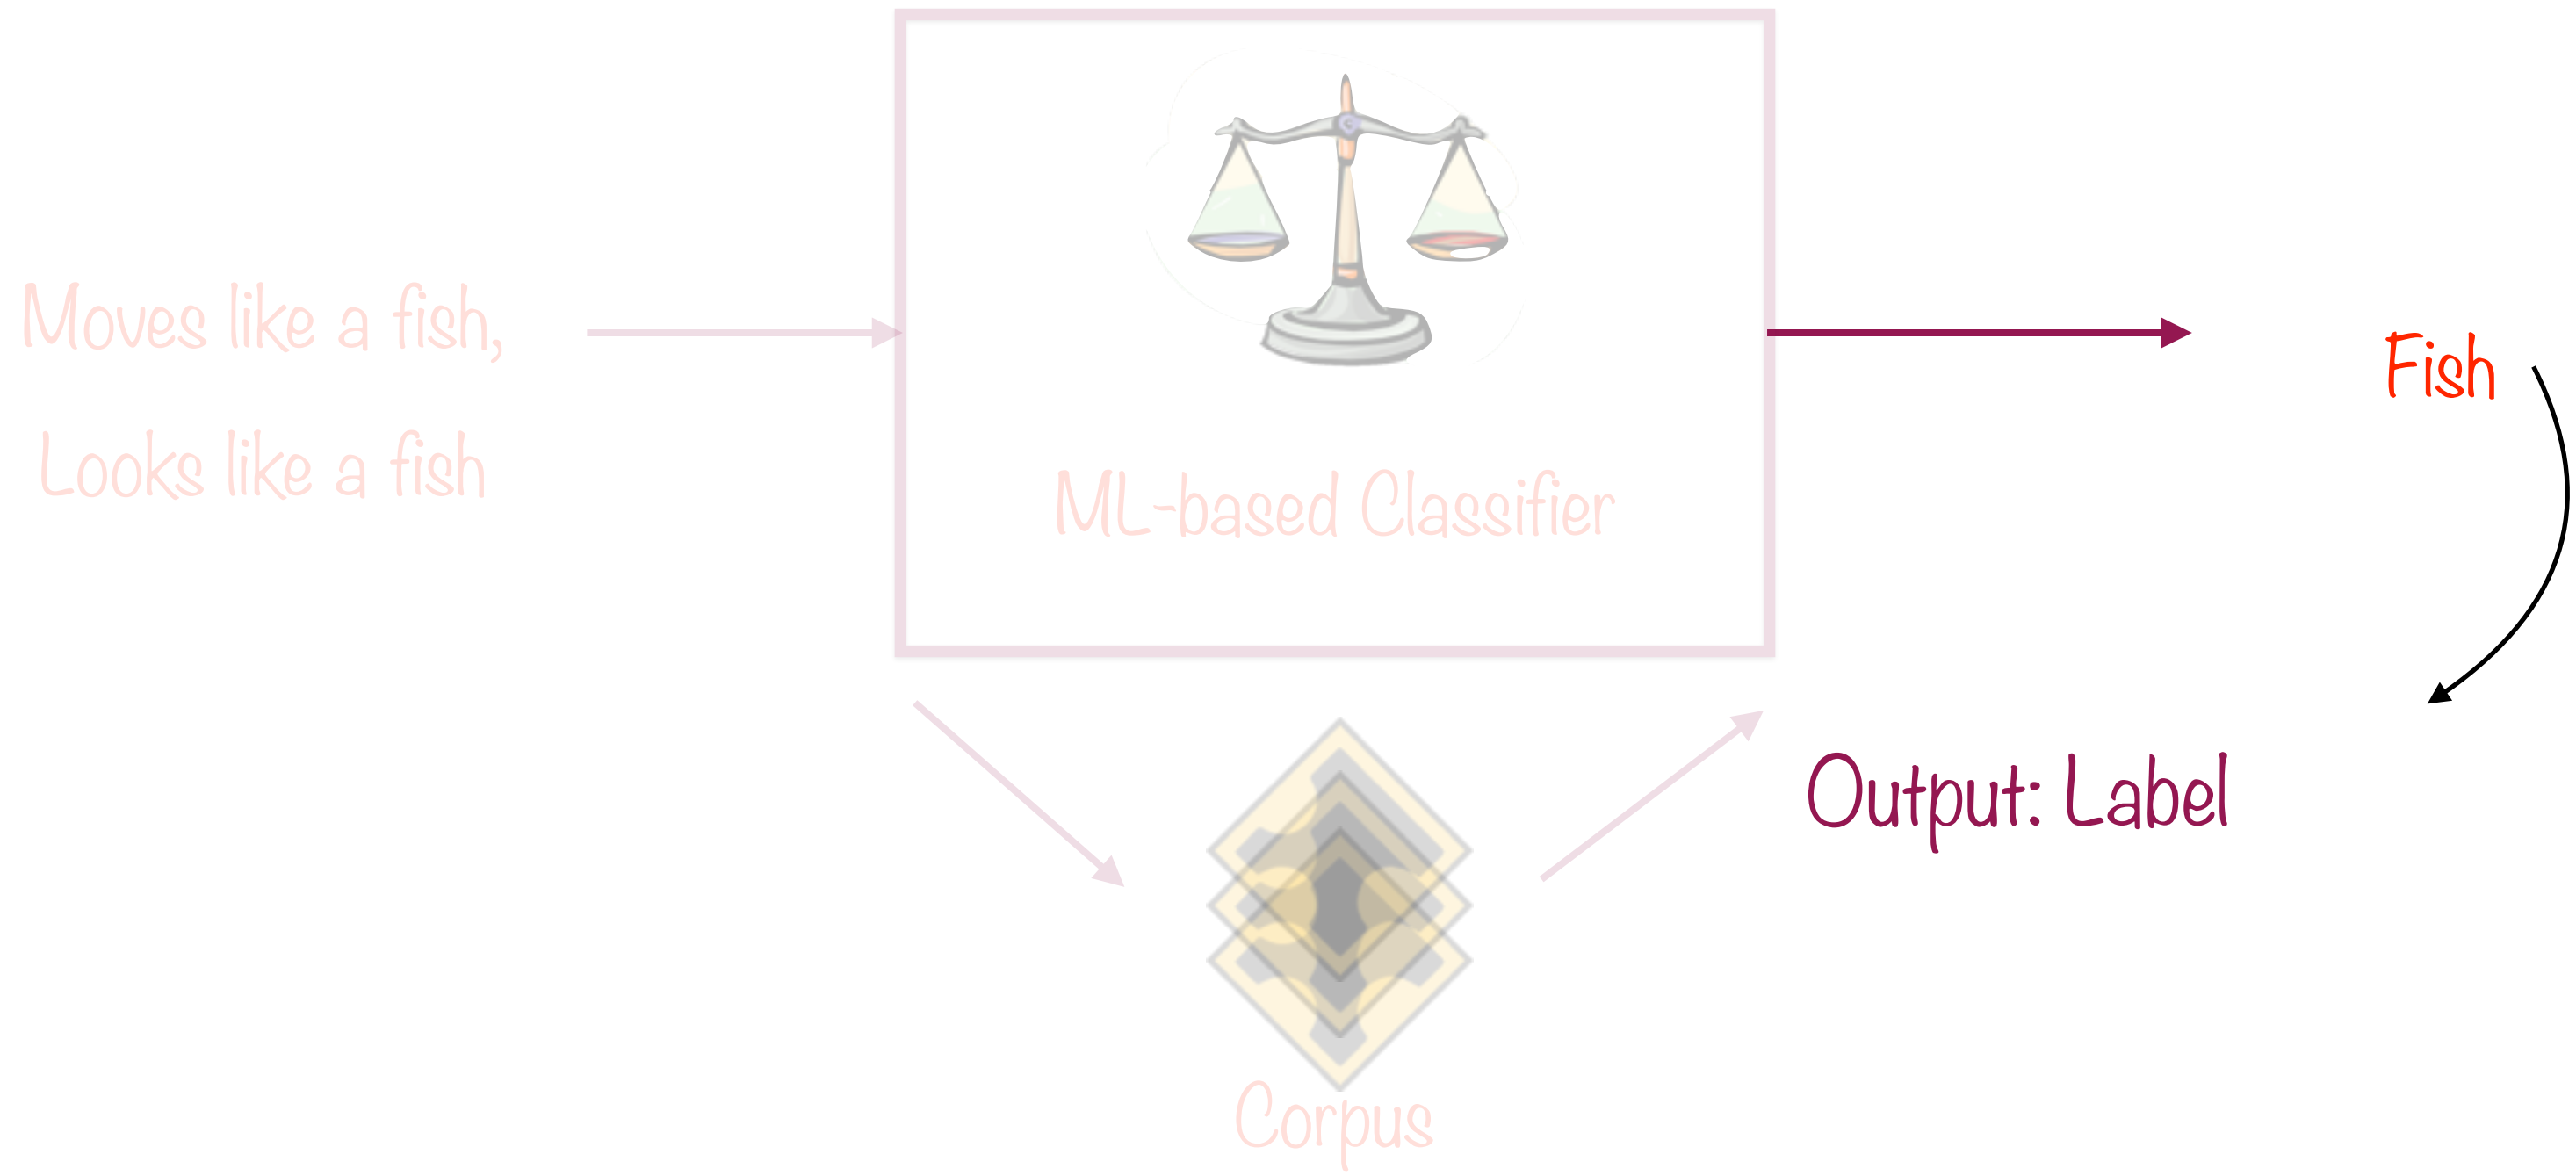
“Traditional” ML-based Binary Classifier



“Traditional” ML-based Binary Classifier



“Traditional” ML-based Binary Classifier



Feature Vectors

The attributes that the ML algorithm focuses on are called features

Each data point is a list - or vector - of such features

Thus, the input into an ML algorithm is a feature vector

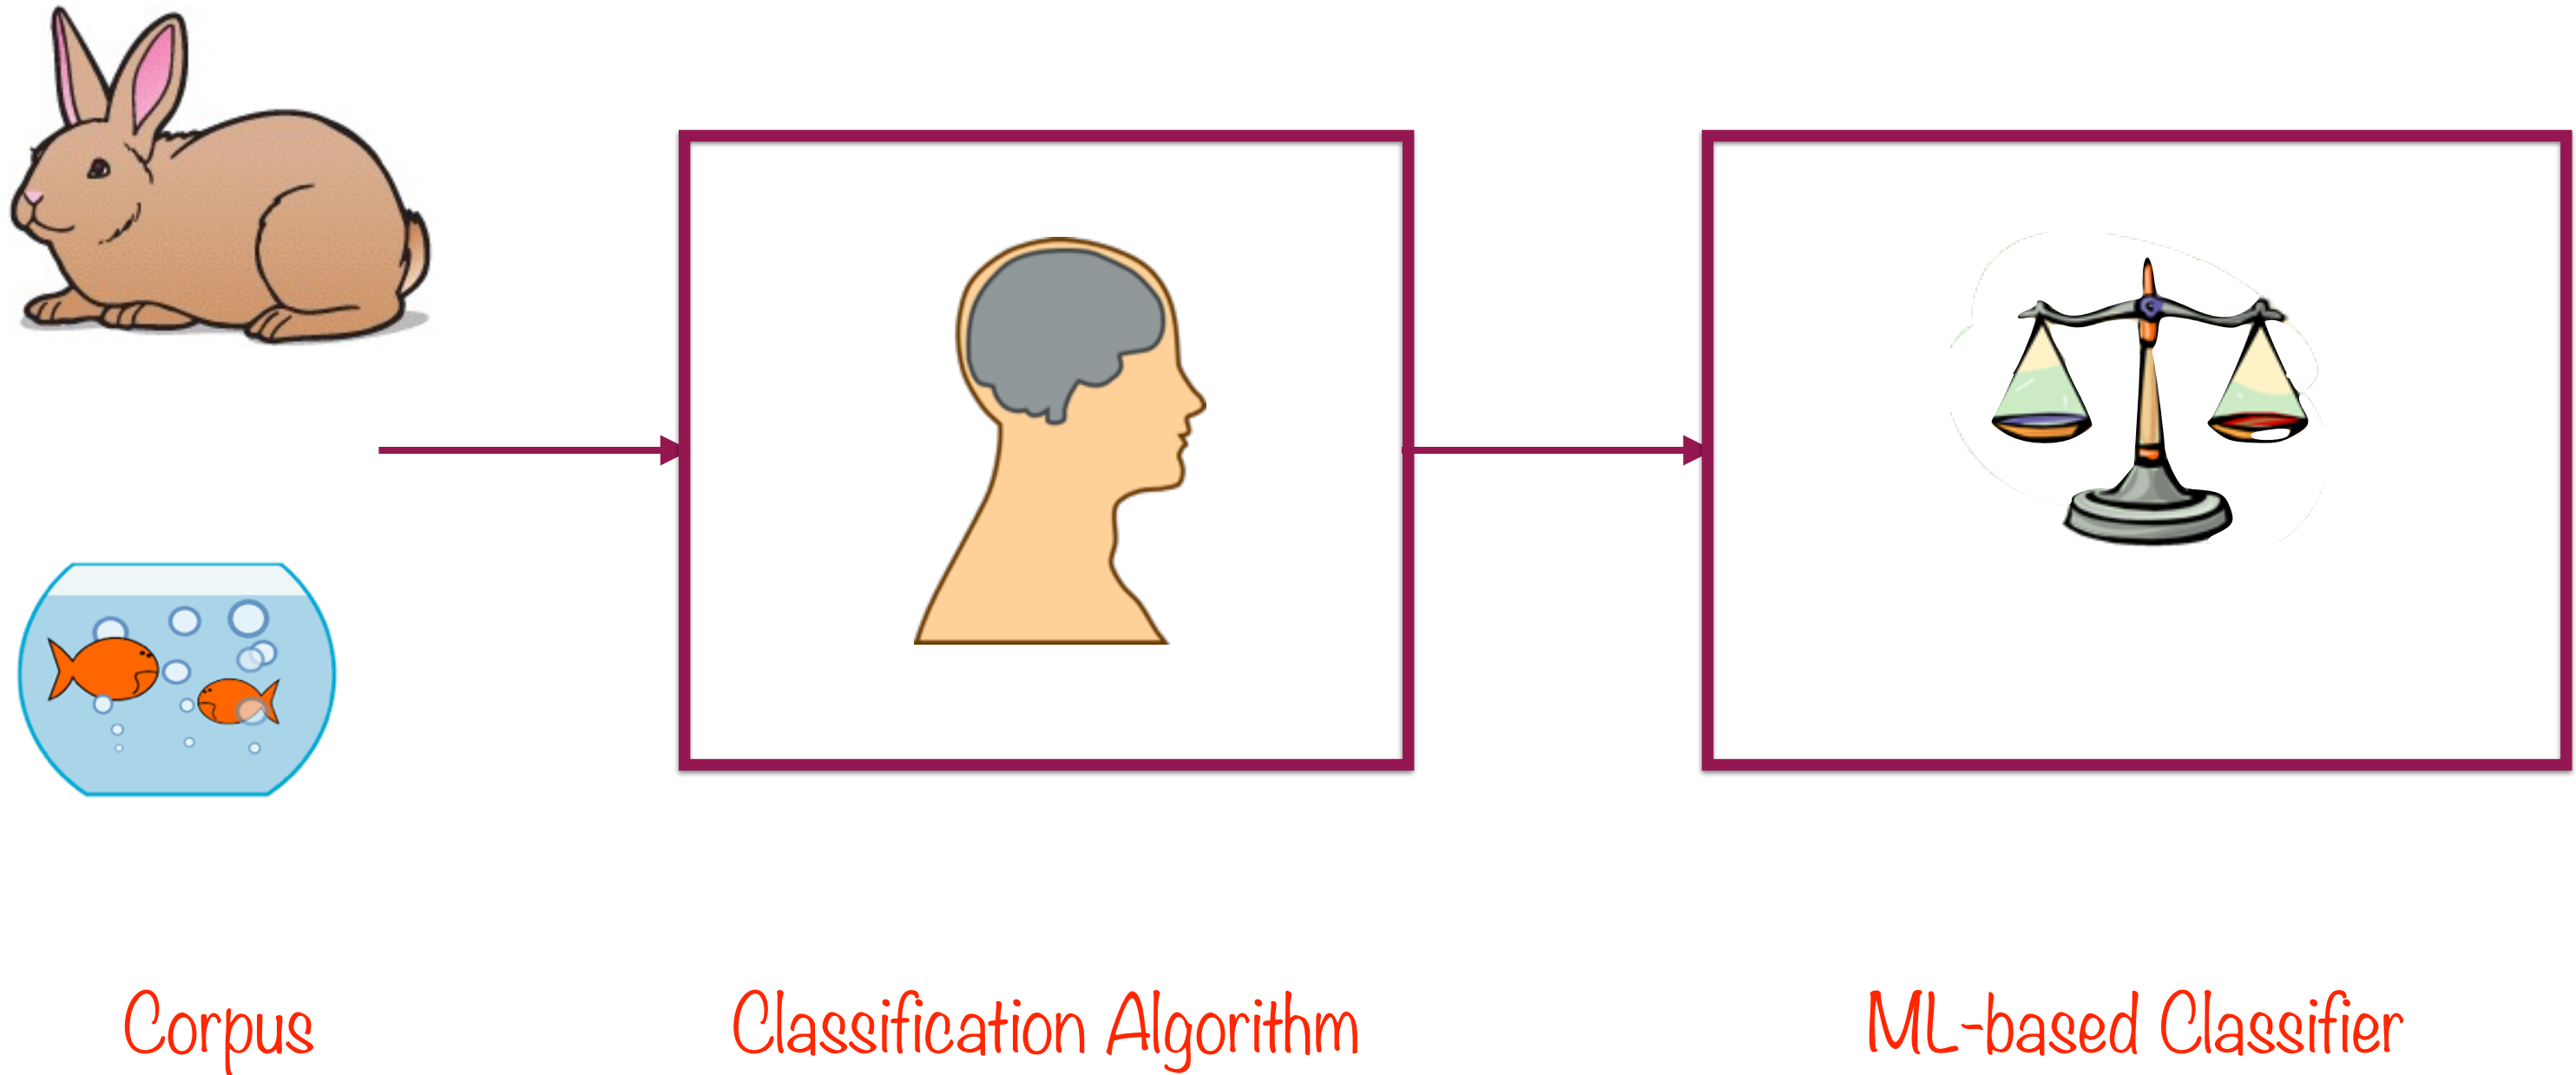
“Traditional” ML-based systems still rely on experts to decide what features to pay attention to

“Representation” ML-based systems figure out by themselves what features to pay attention to

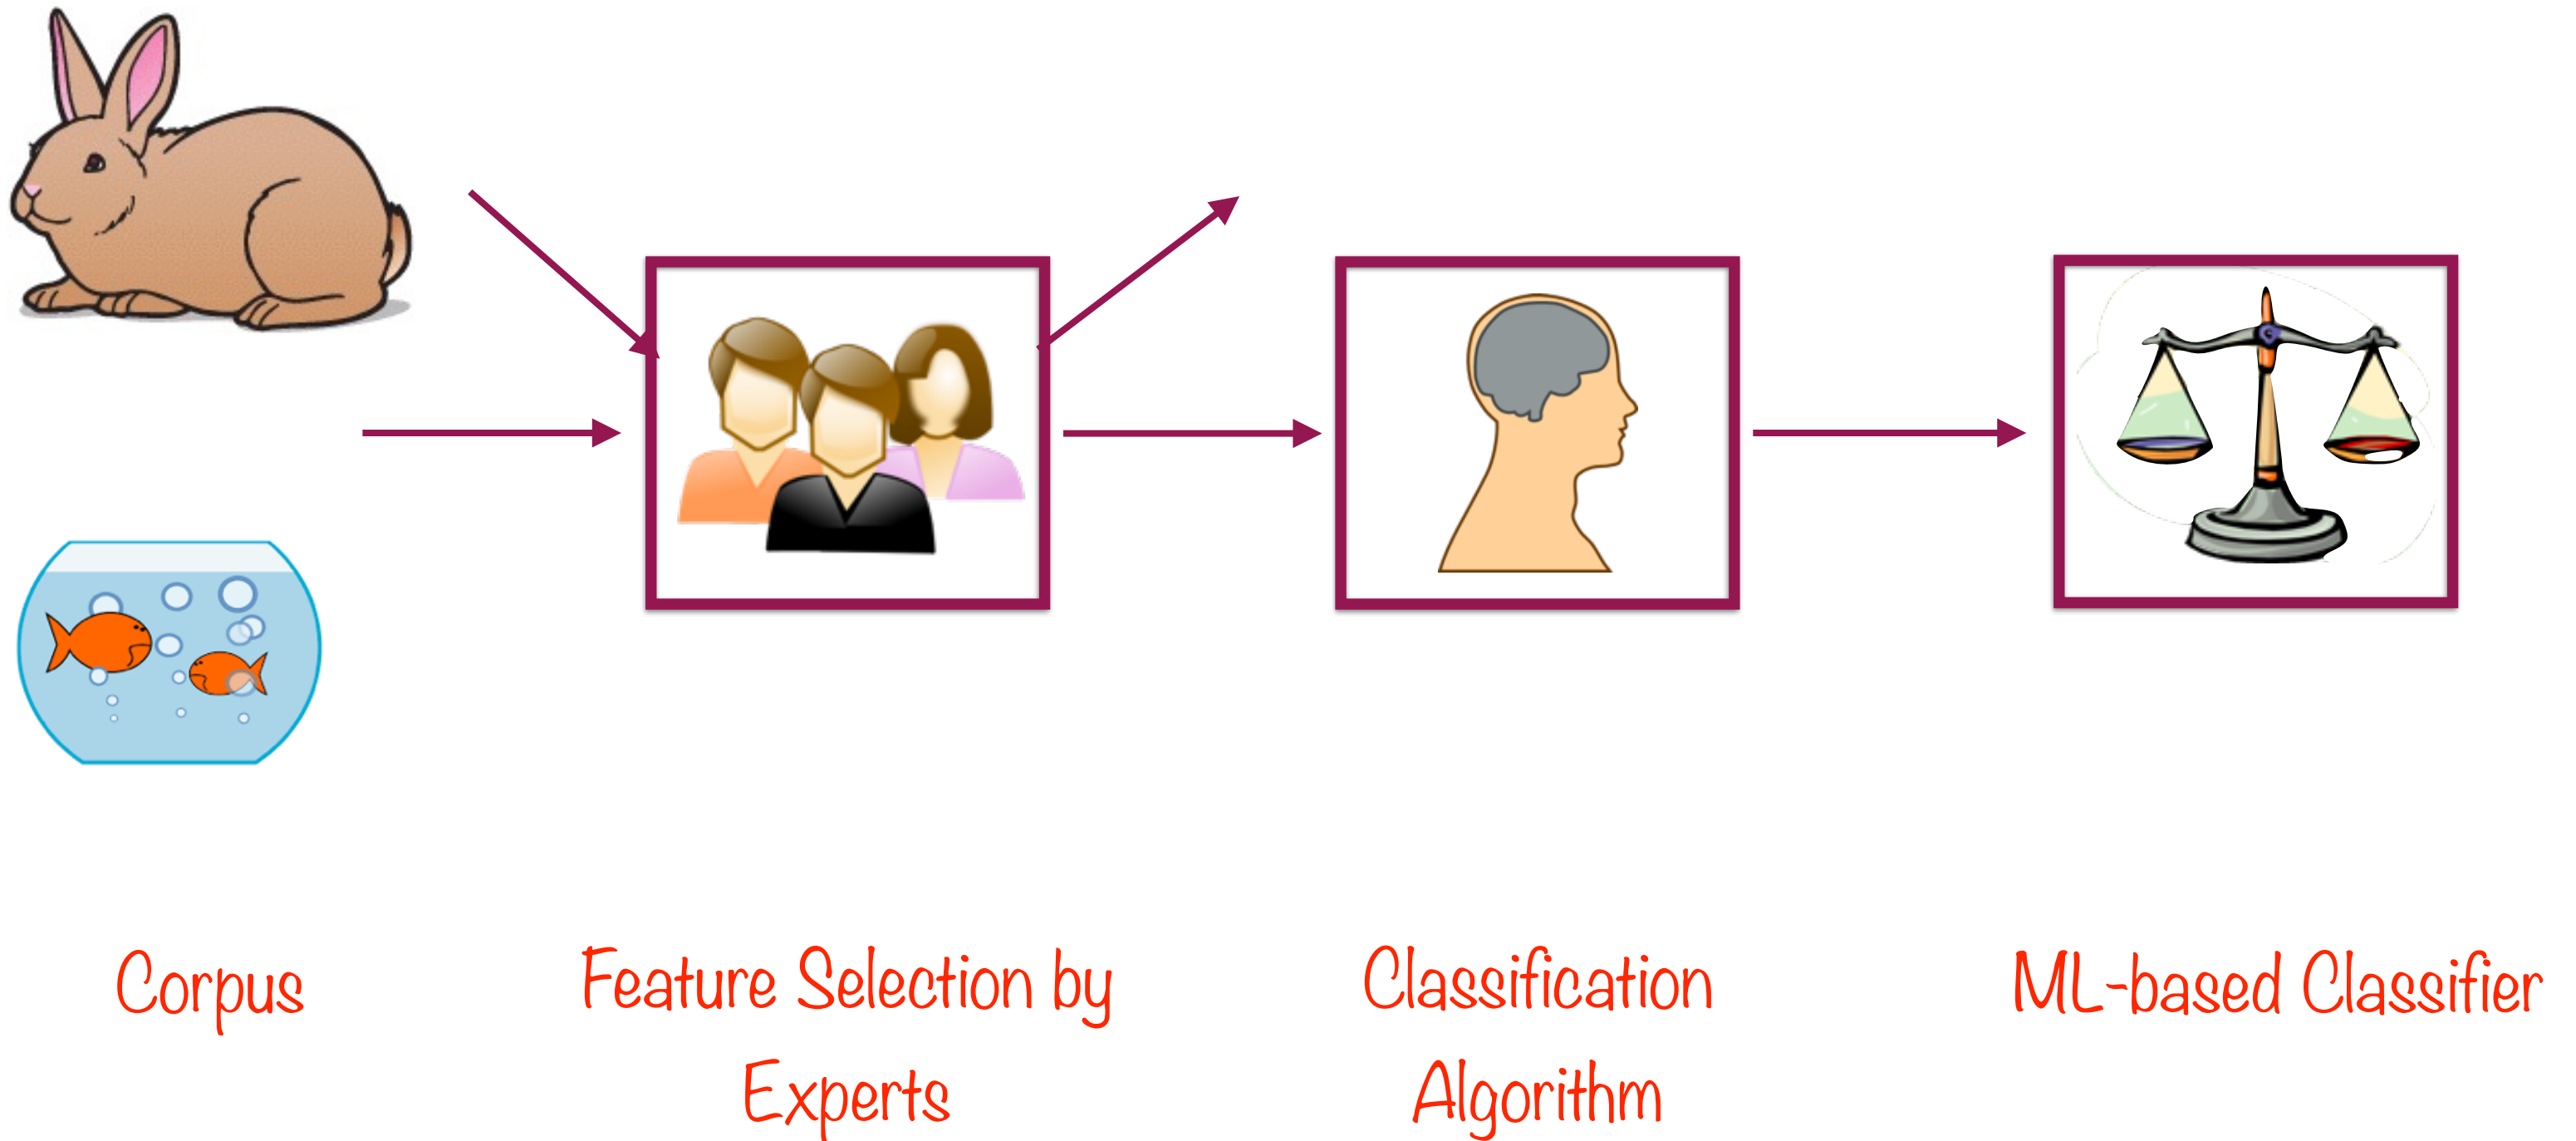
Understanding Deep Learning

“Representation” ML-based systems figure out by themselves what features to pay attention to


“Traditional” ML-based Binary Classifier



“Traditional” ML-based Binary Classifier



“Traditional” ML-based Binary Classifier




Corpus

Feature Selection by
Experts

Classification
Algorithm

ML-based Classifier

“Traditional” ML-based Binary Classifier



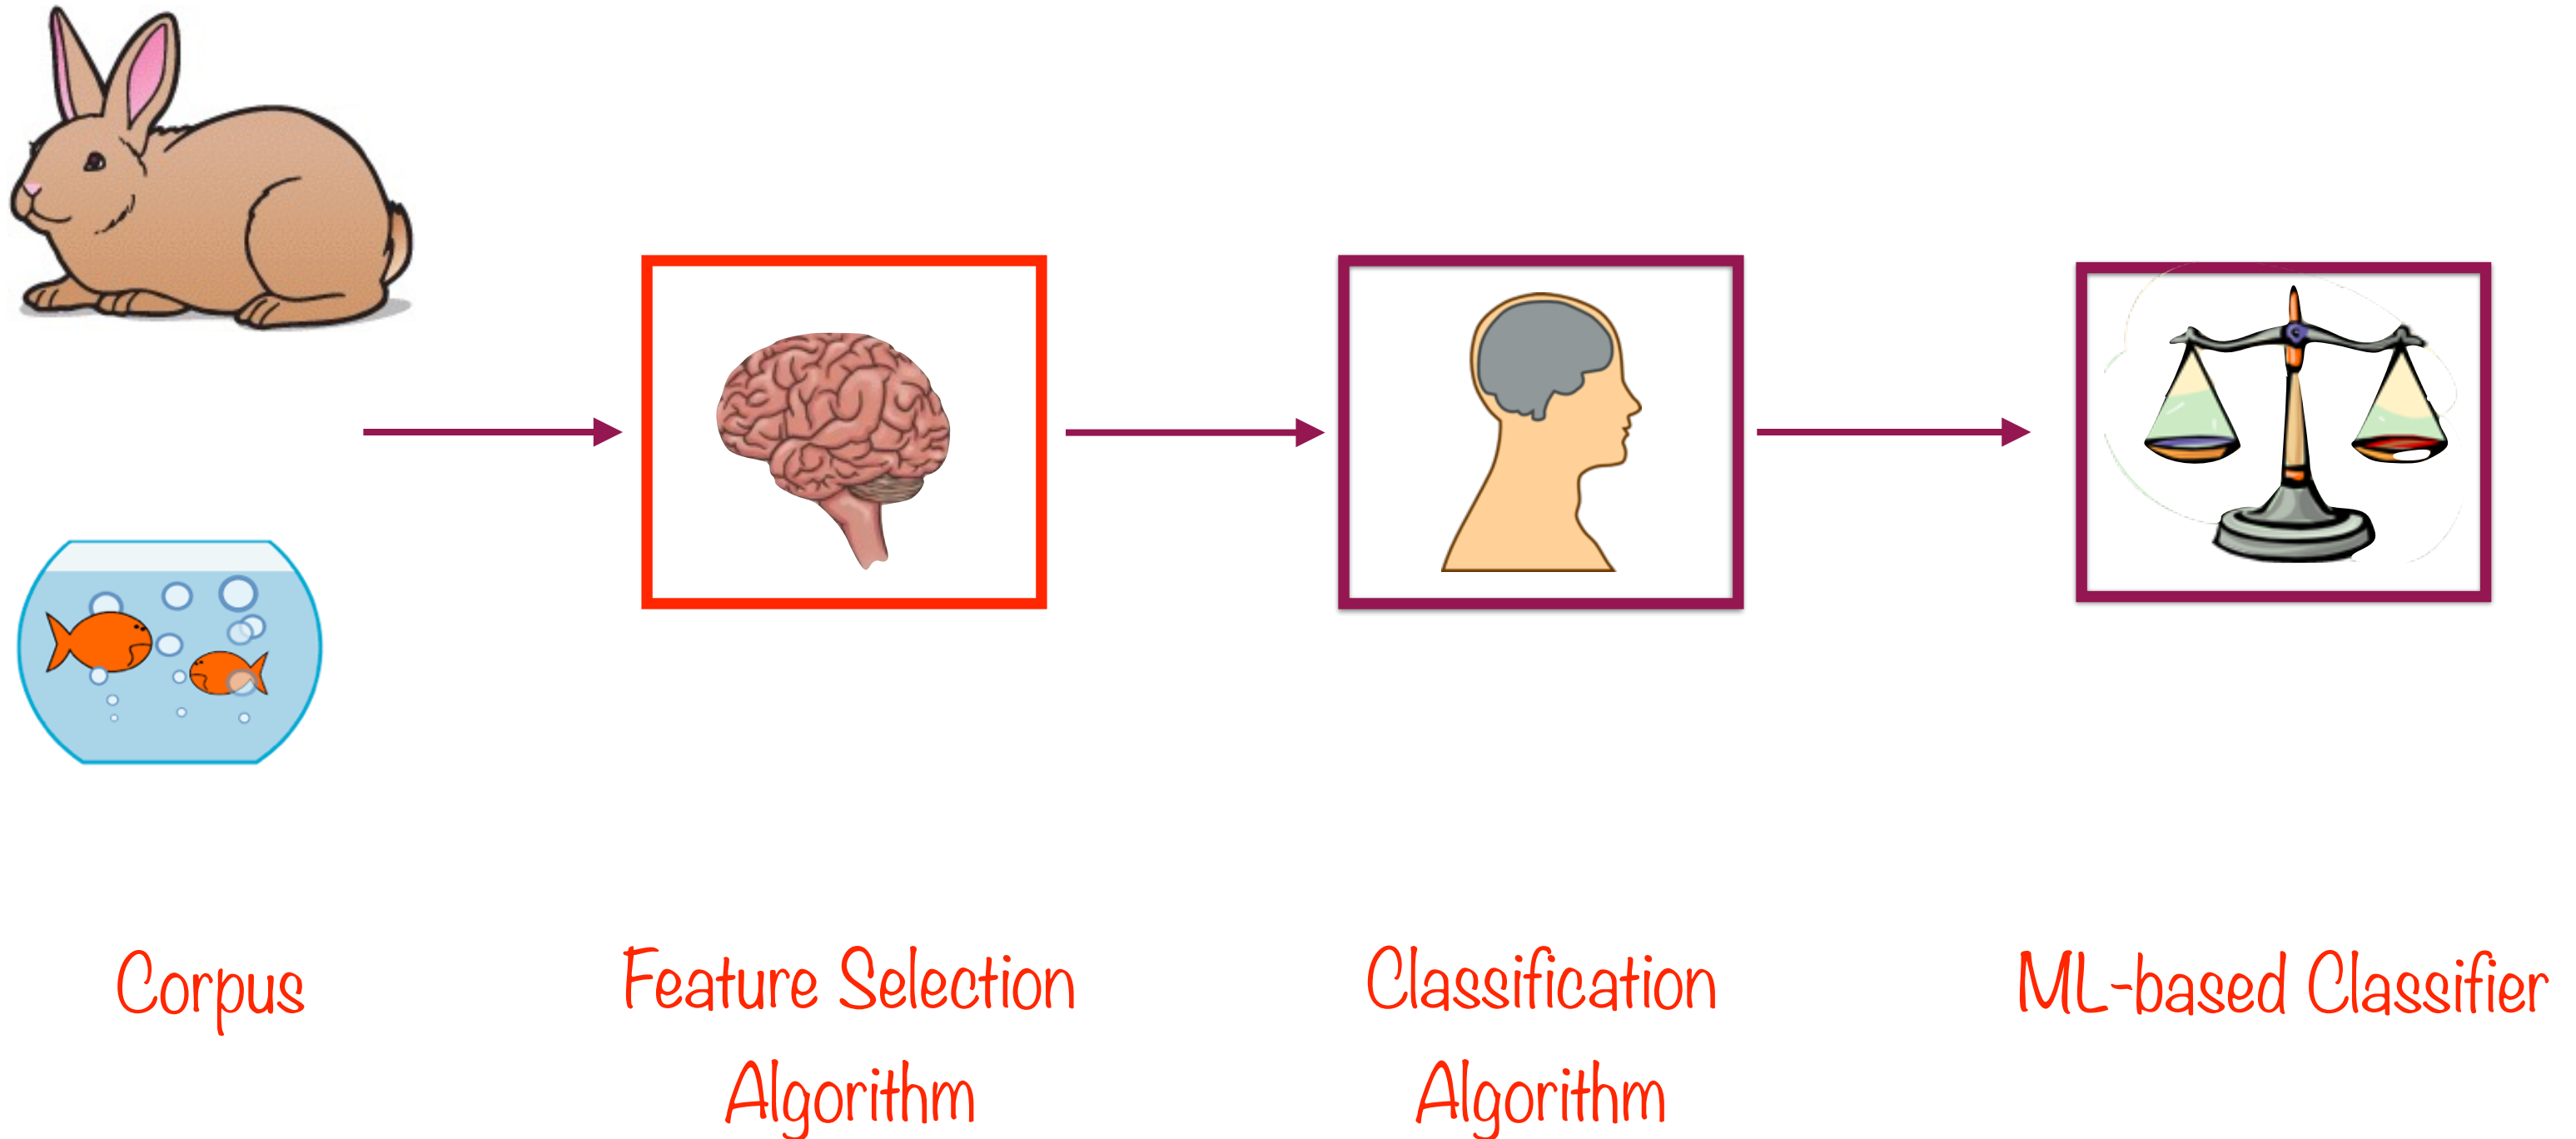
Corpus

Feature Selection by
Experts

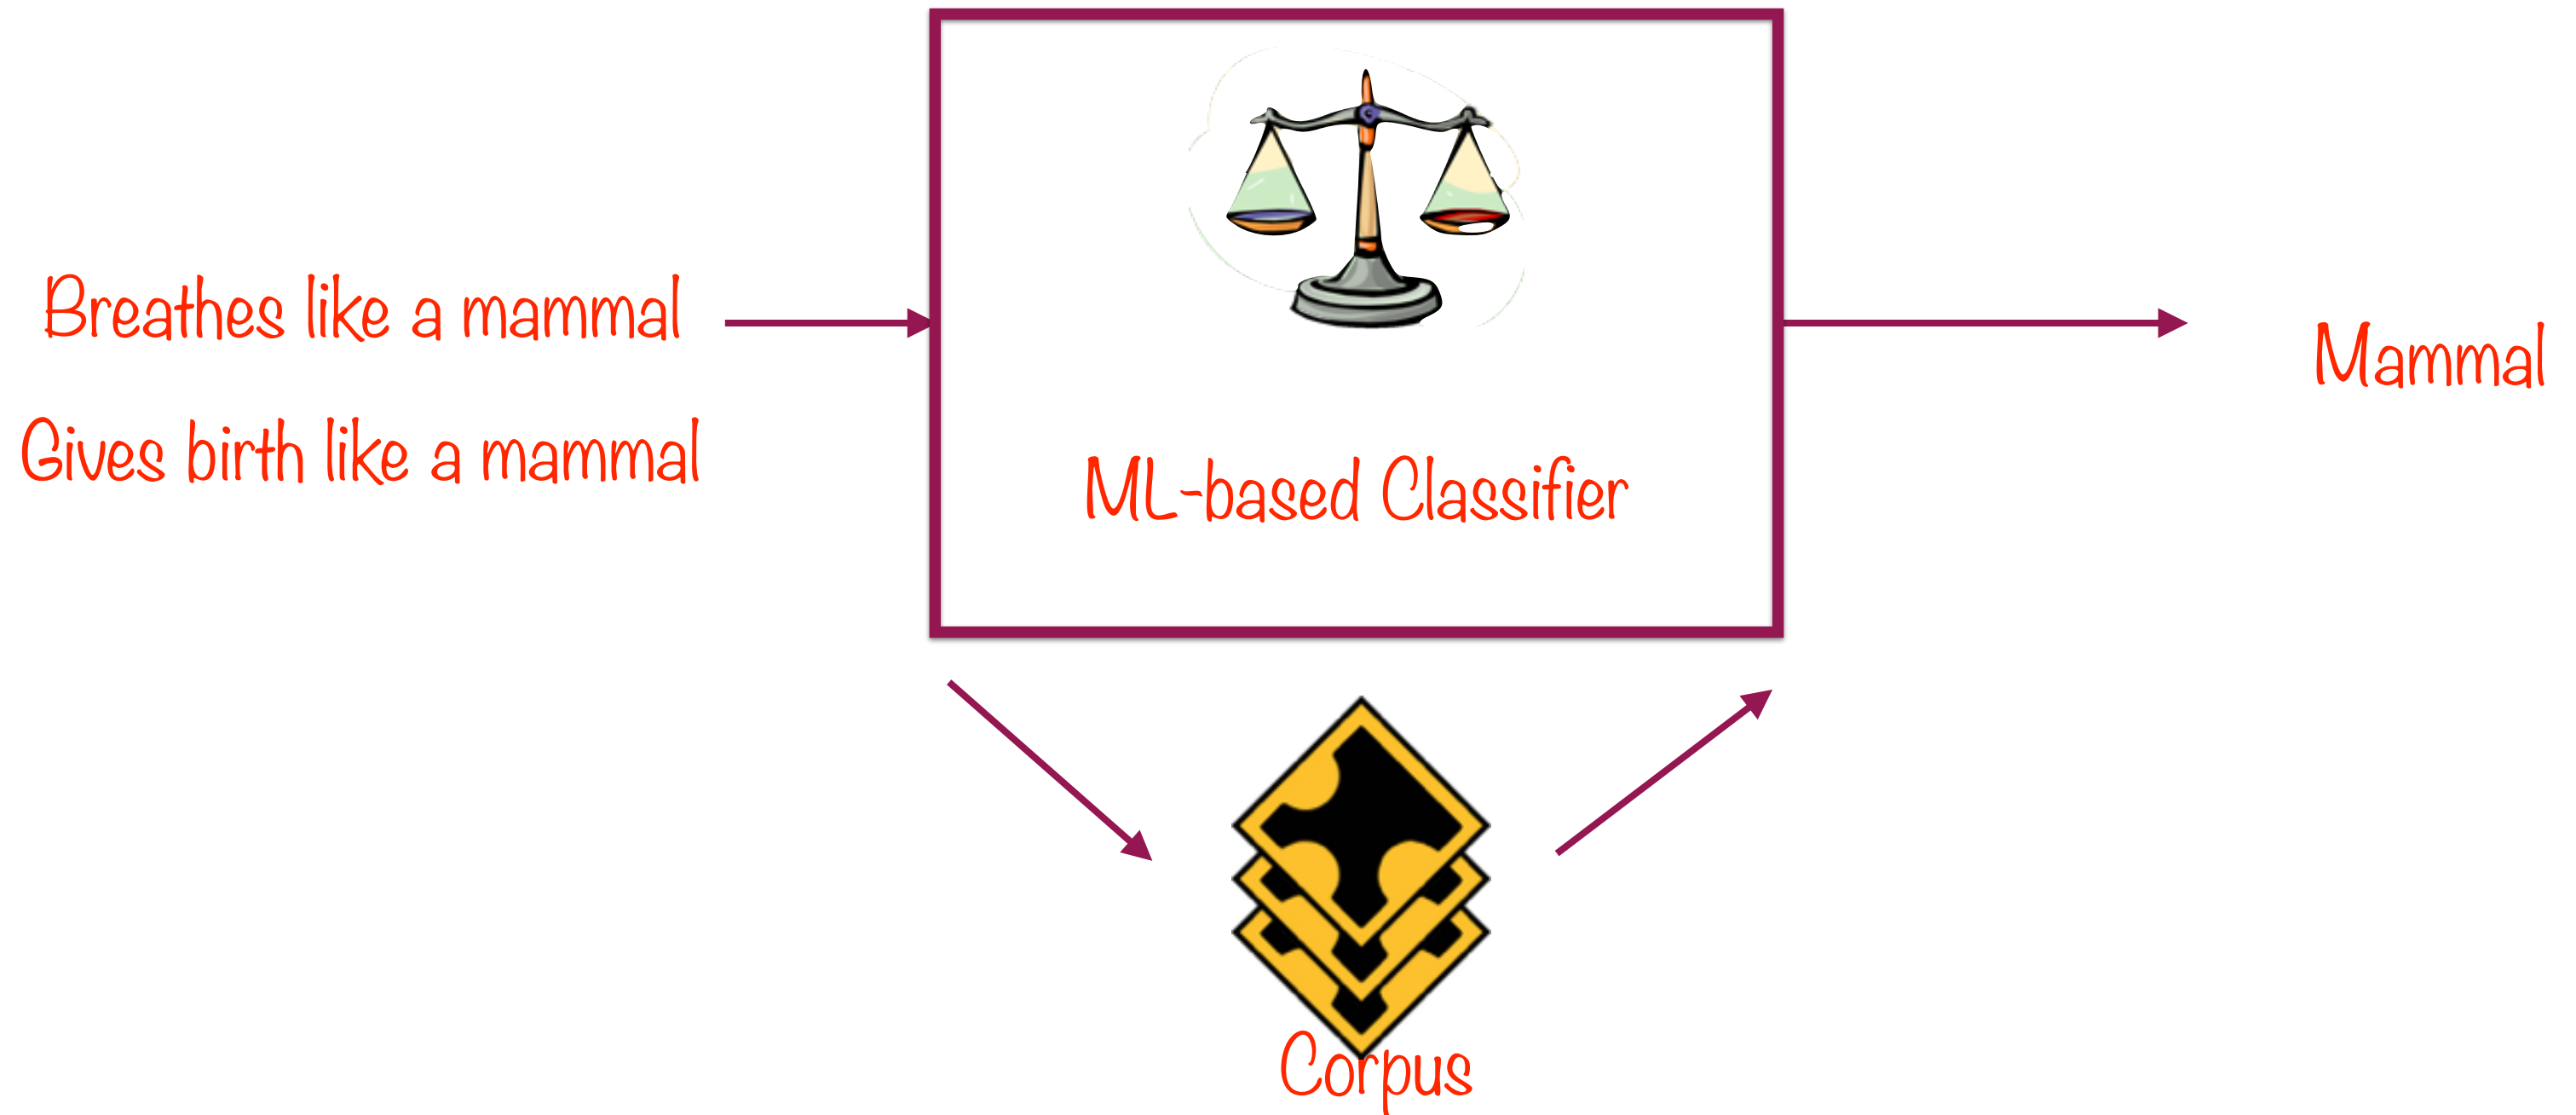
Classification
Algorithm

ML-based Classifier

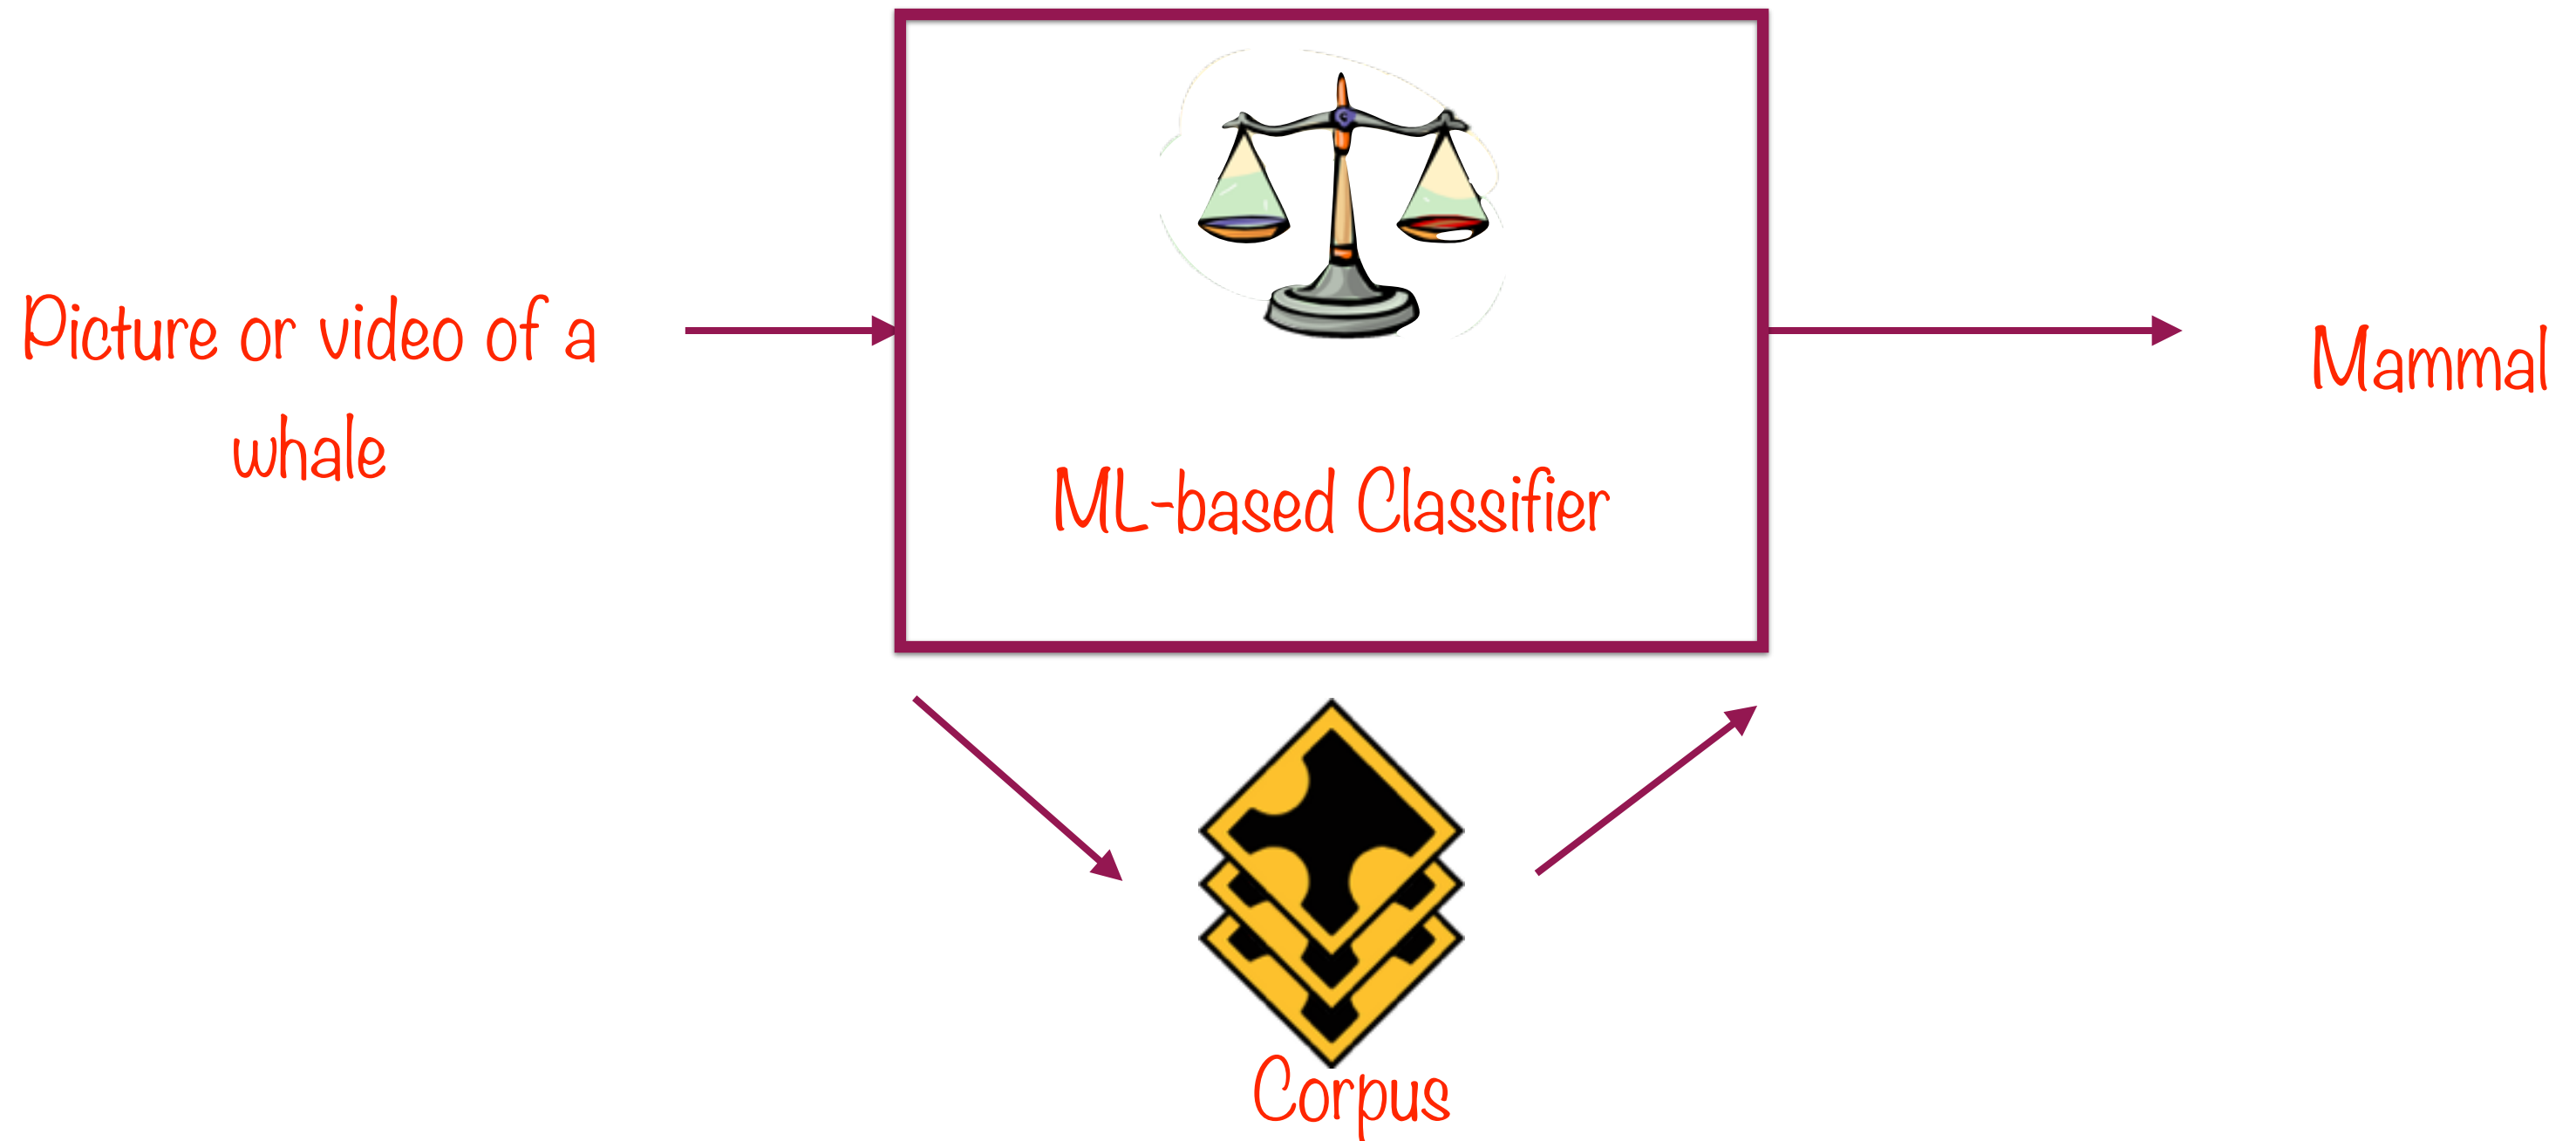
“Representation” ML-based Binary Classifier



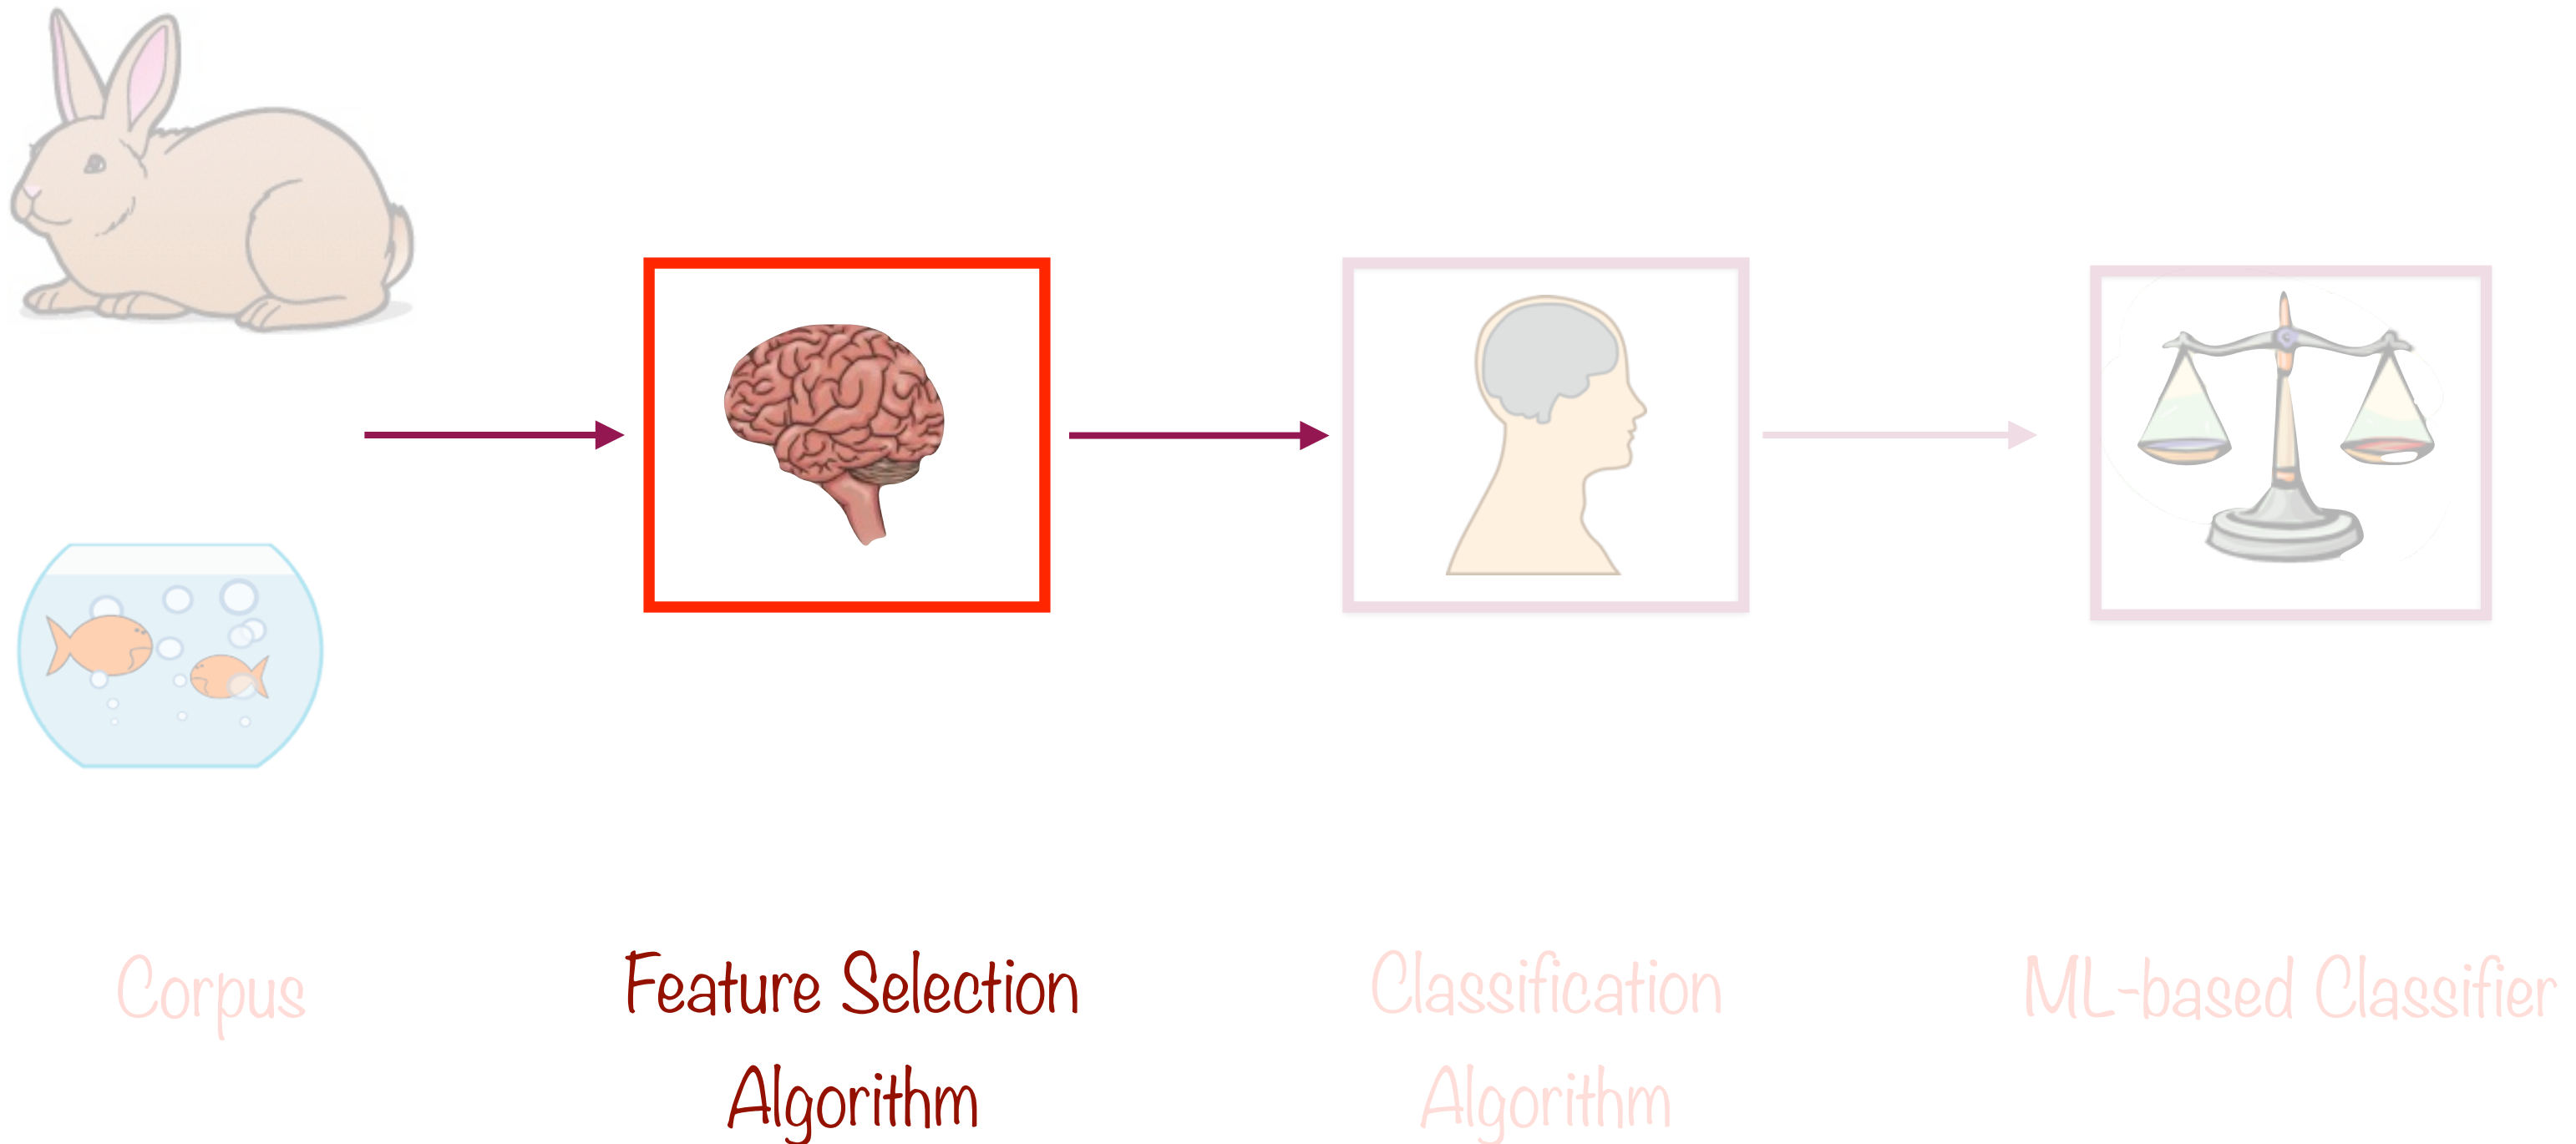
“Traditional” ML-based Binary Classifier



“Representation” ML-based Binary Classifier



“Representation” ML-based Binary Classifier



**“Deep Learning” systems are one type of
representation systems**

Deep Learning and Neural Networks

Deep Learning and Neural Networks

Deep Learning

Algorithms that learn what features matter

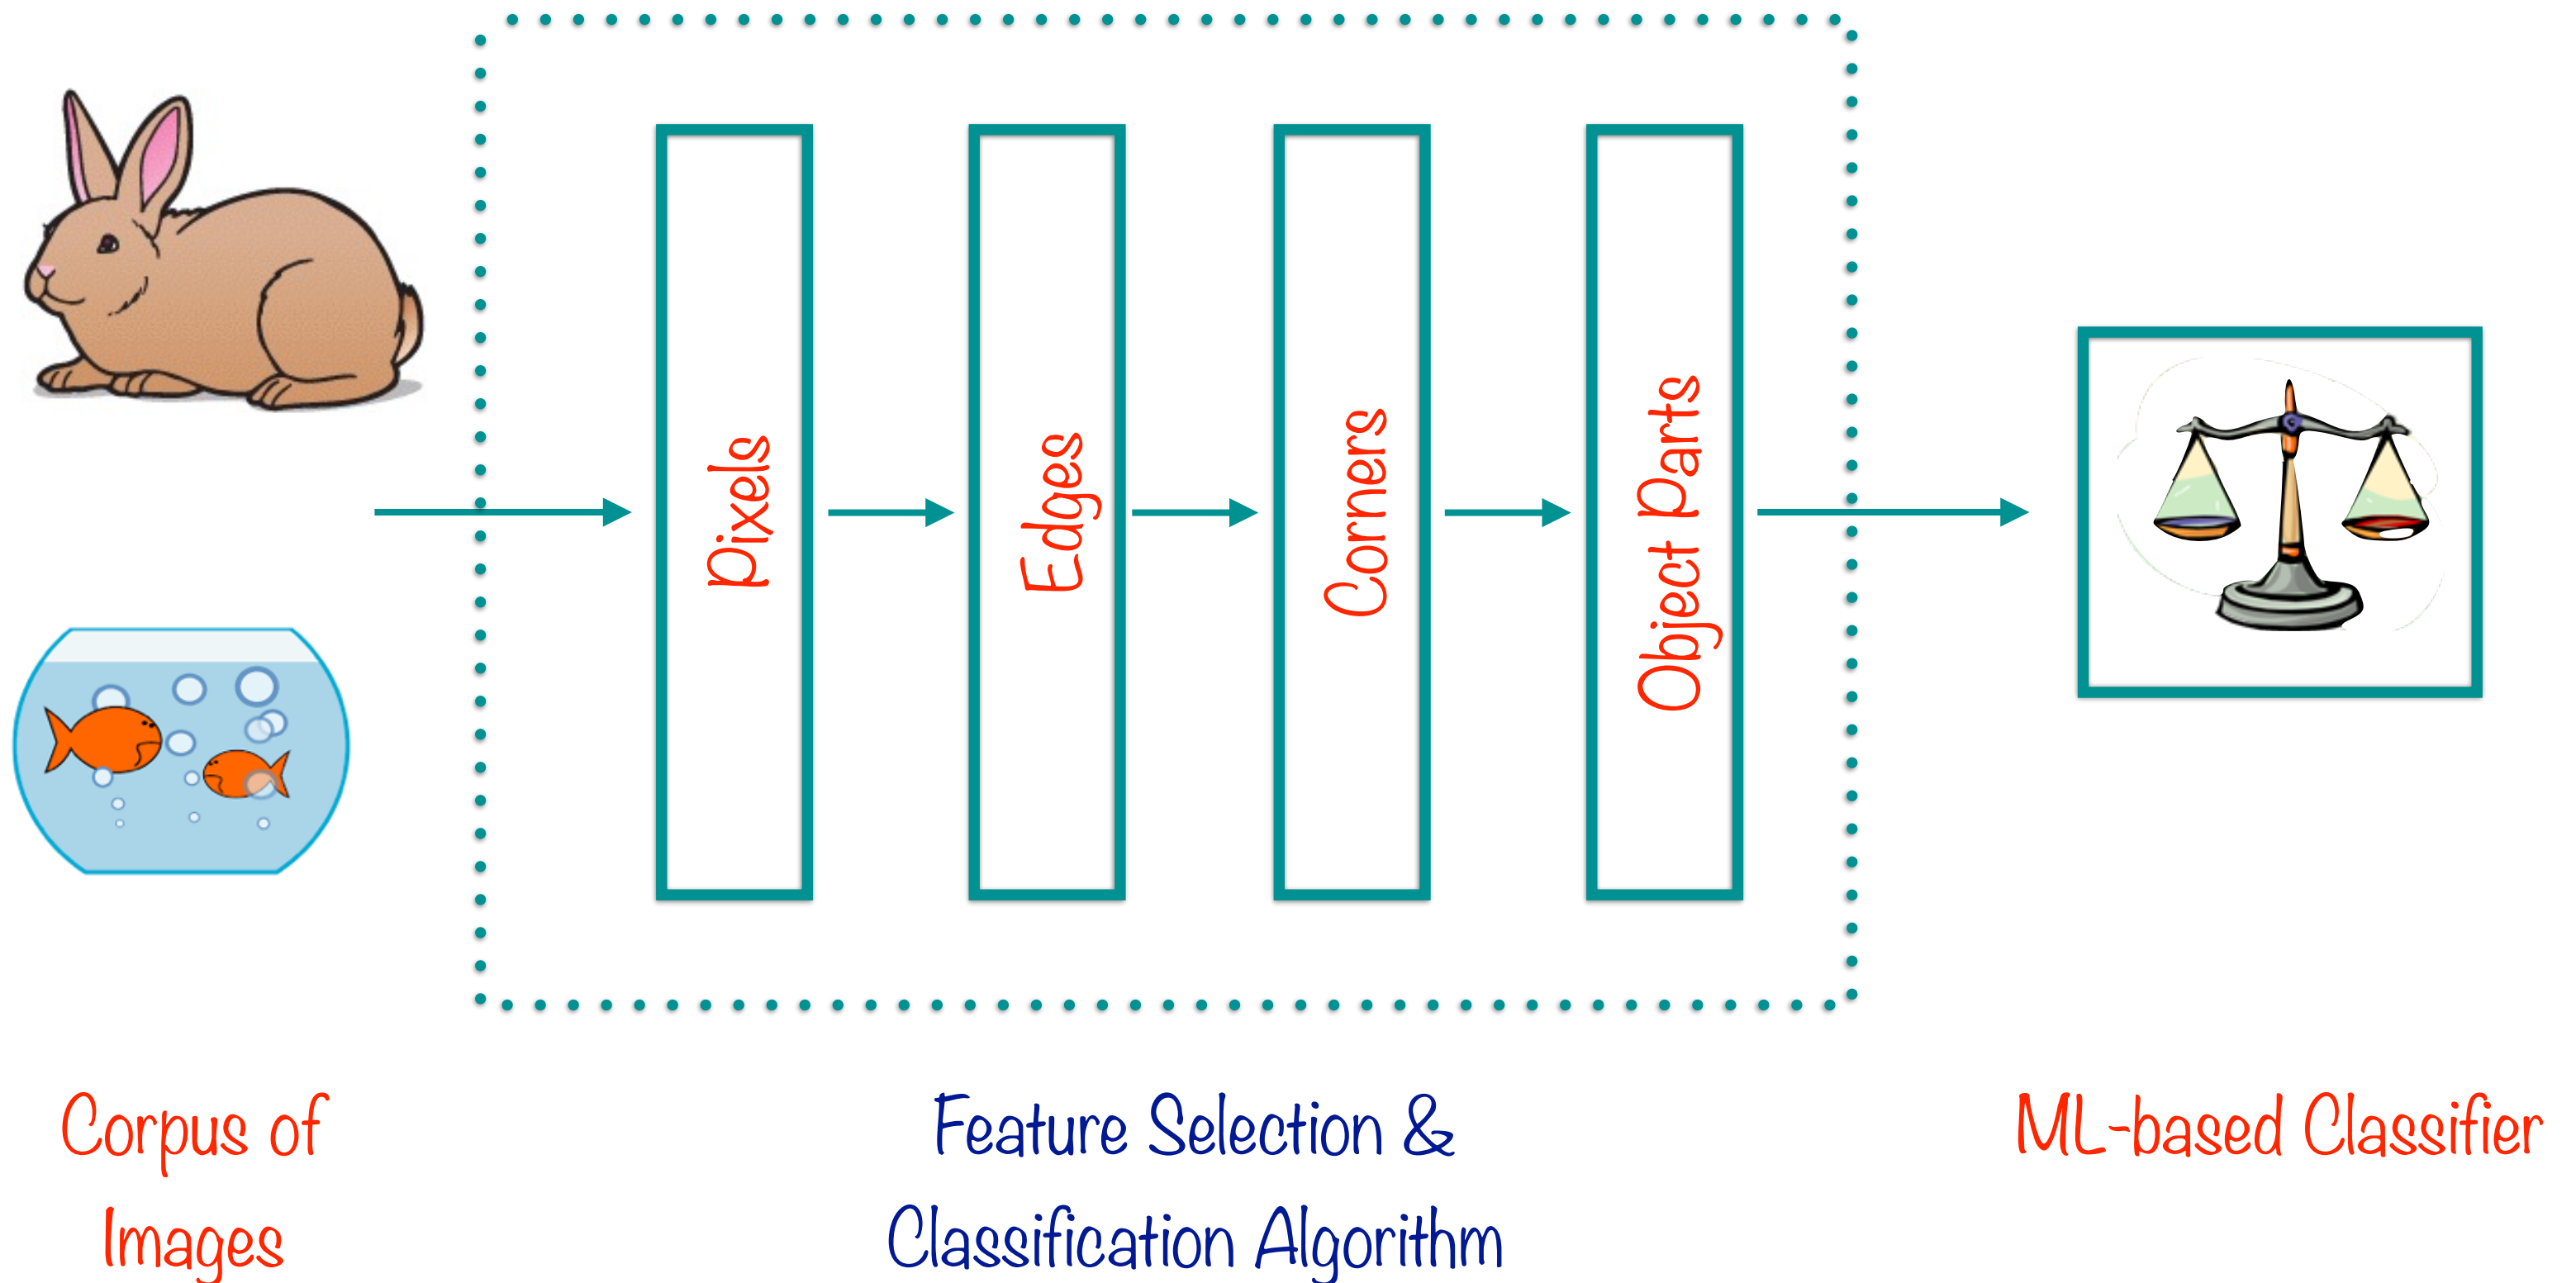
Neural Networks

The most common class of deep learning algorithms


Neurons

Simple building blocks that actually “learn”

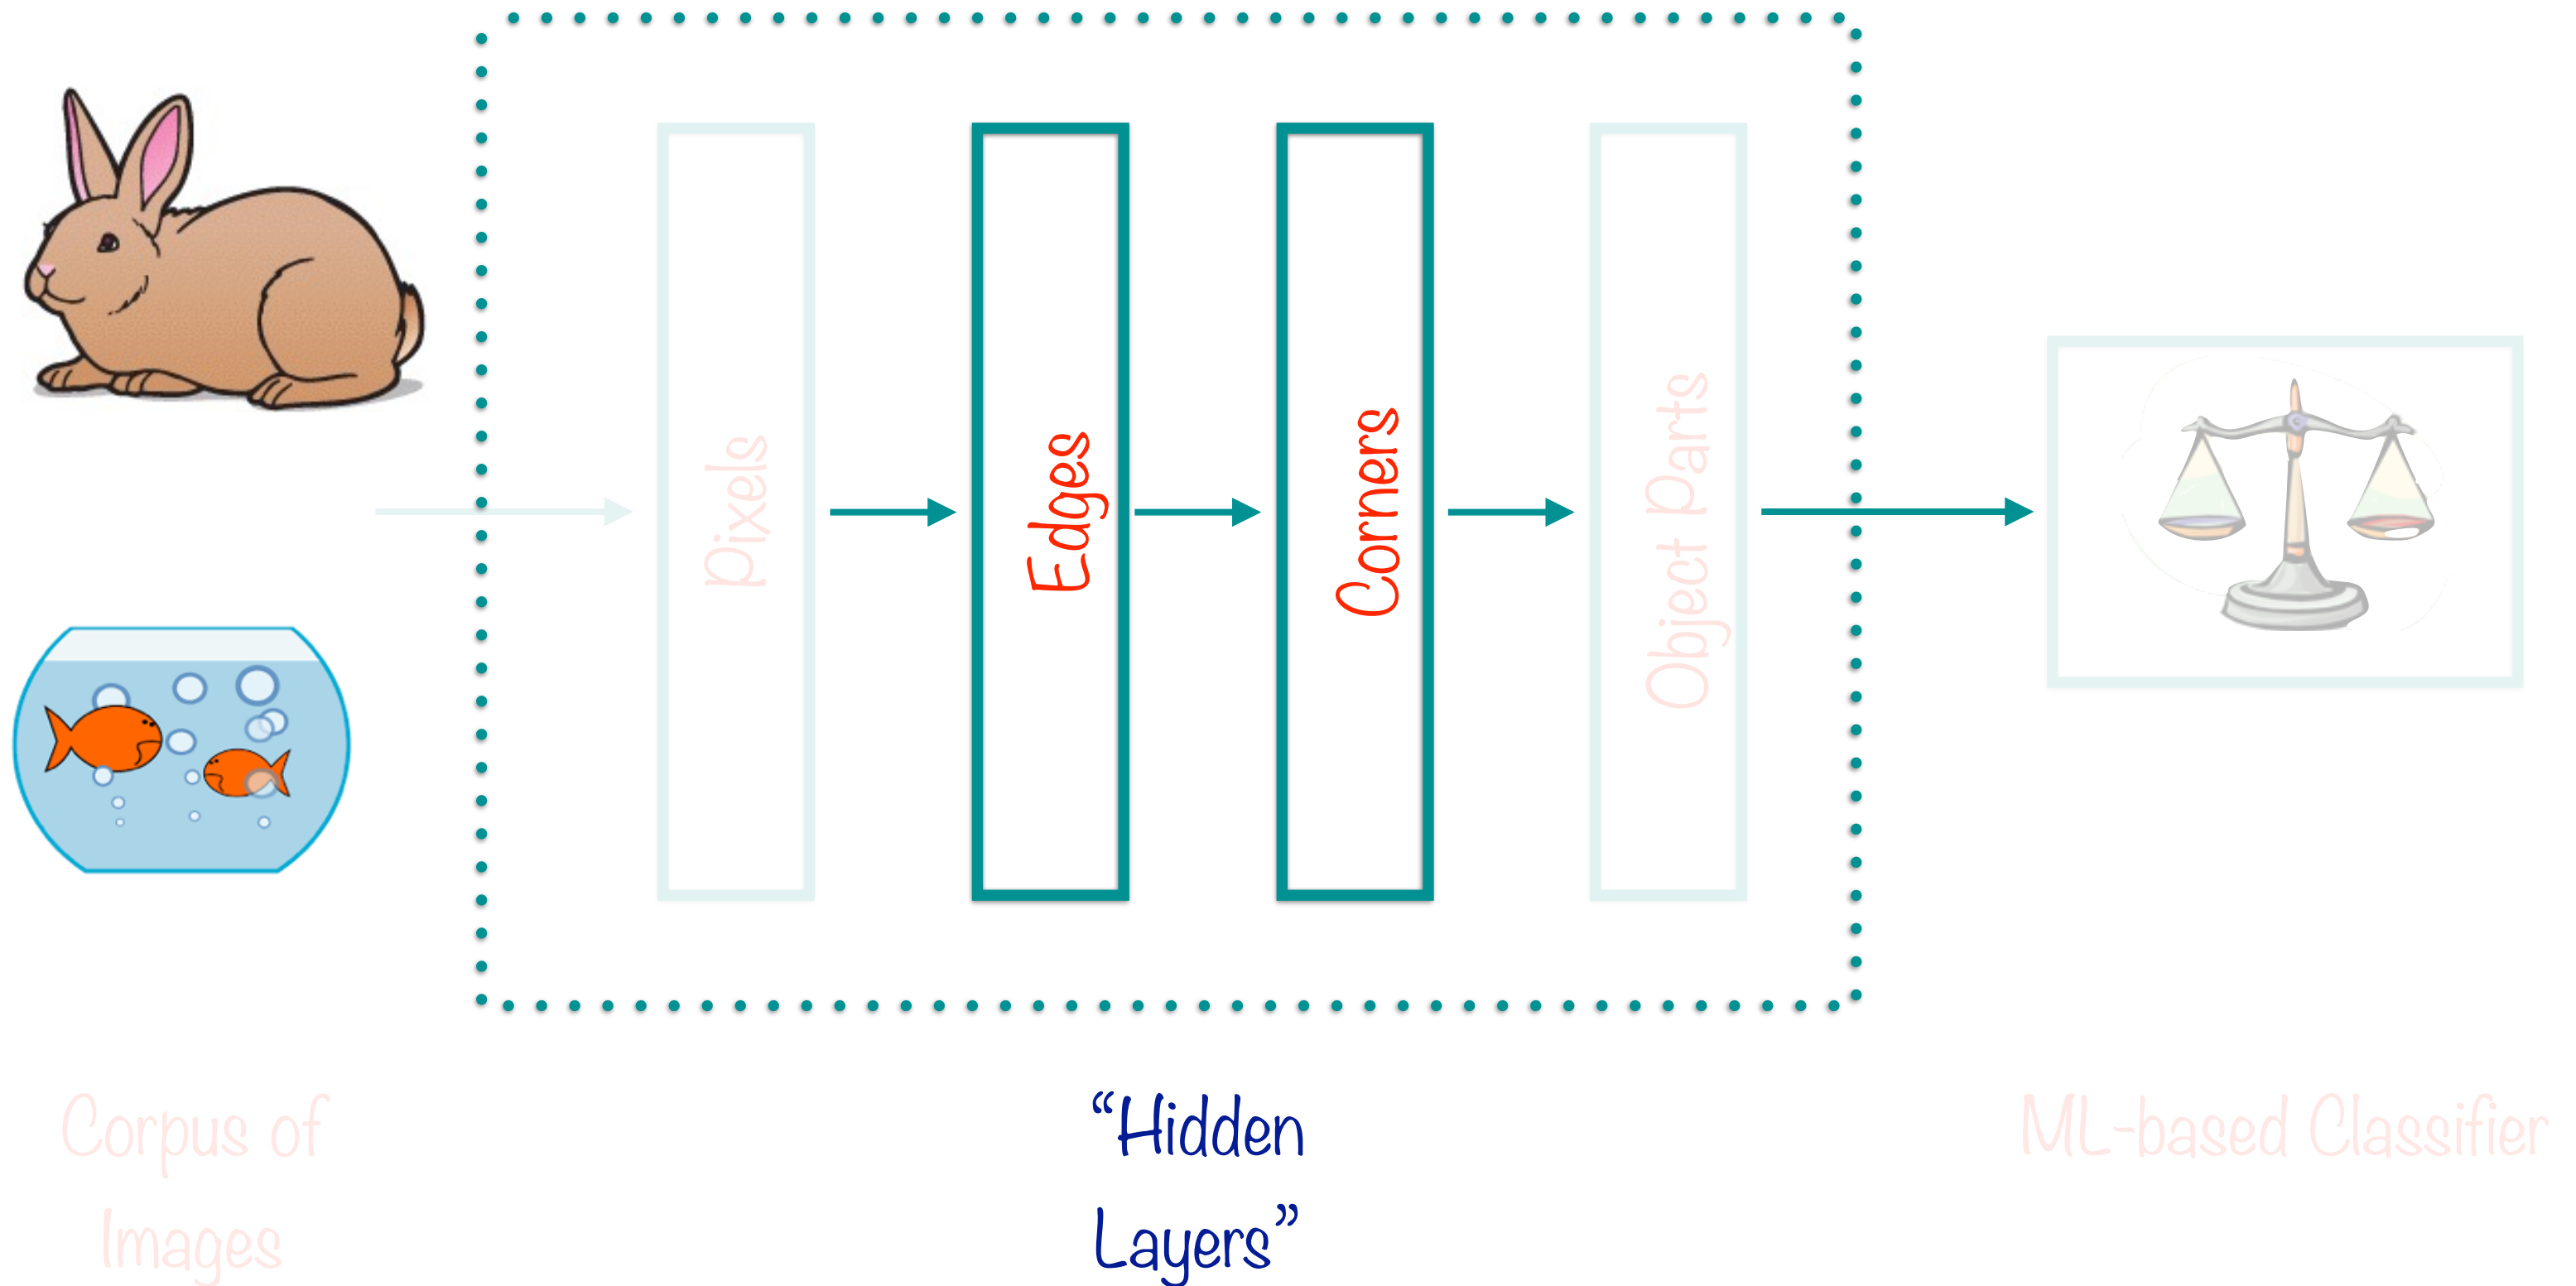
"Deep Learning"-based Binary Classifier



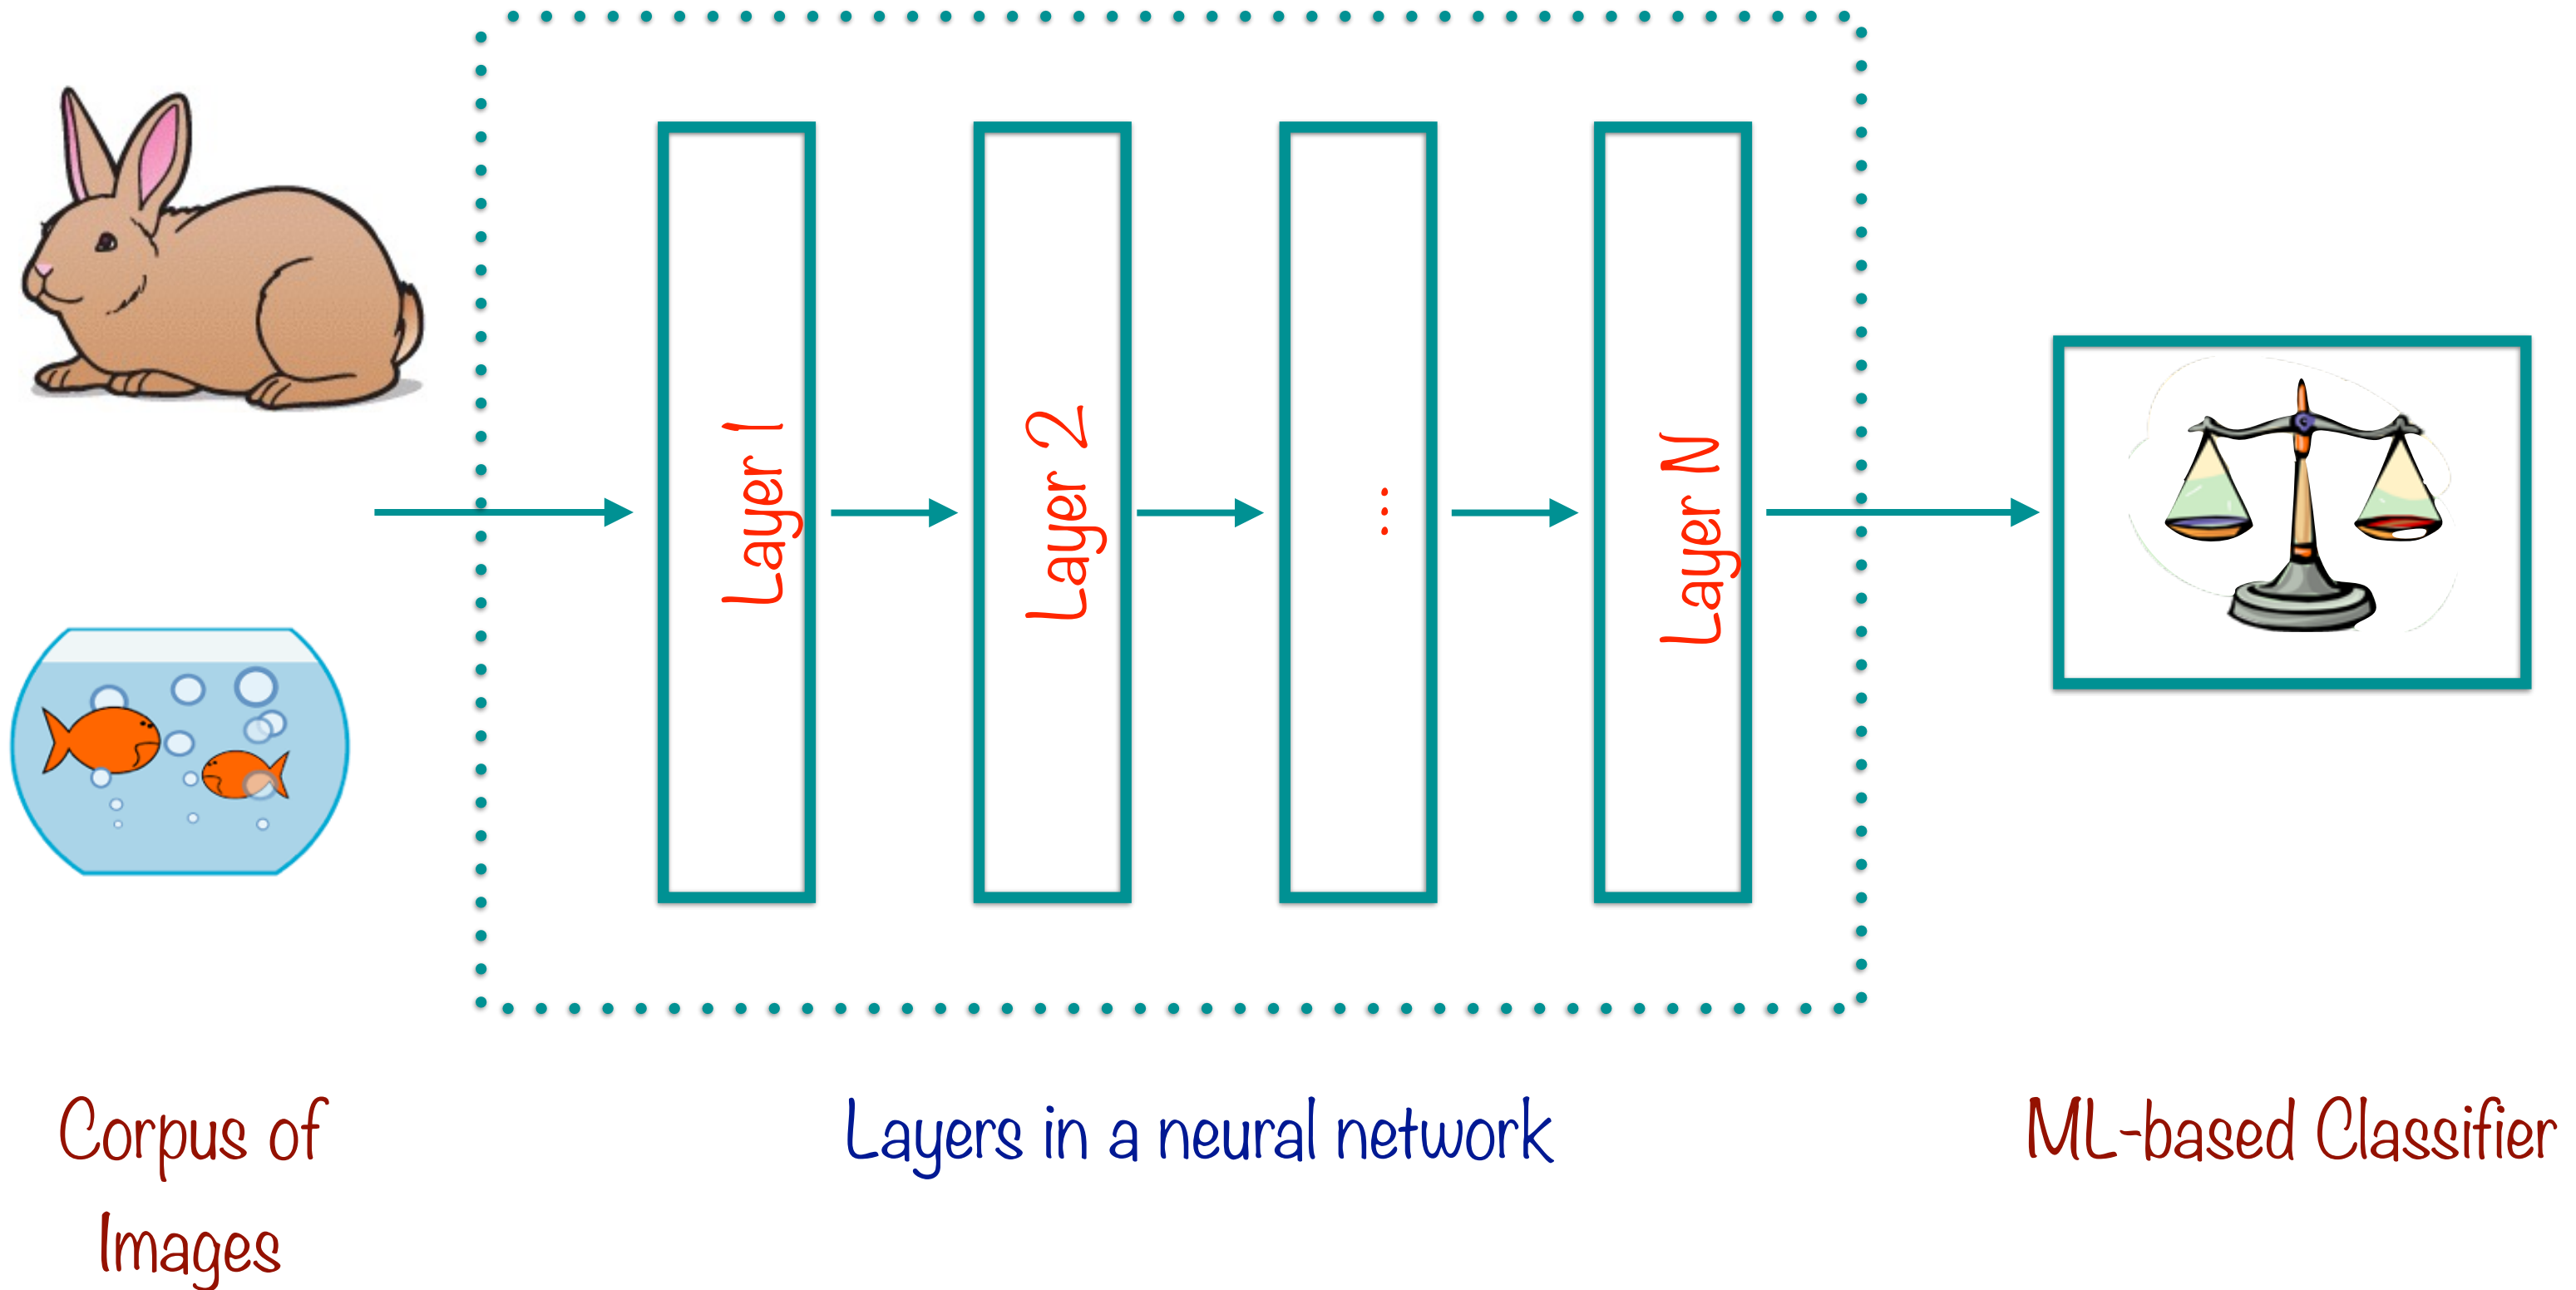
“Deep Learning”-based Binary Classifier



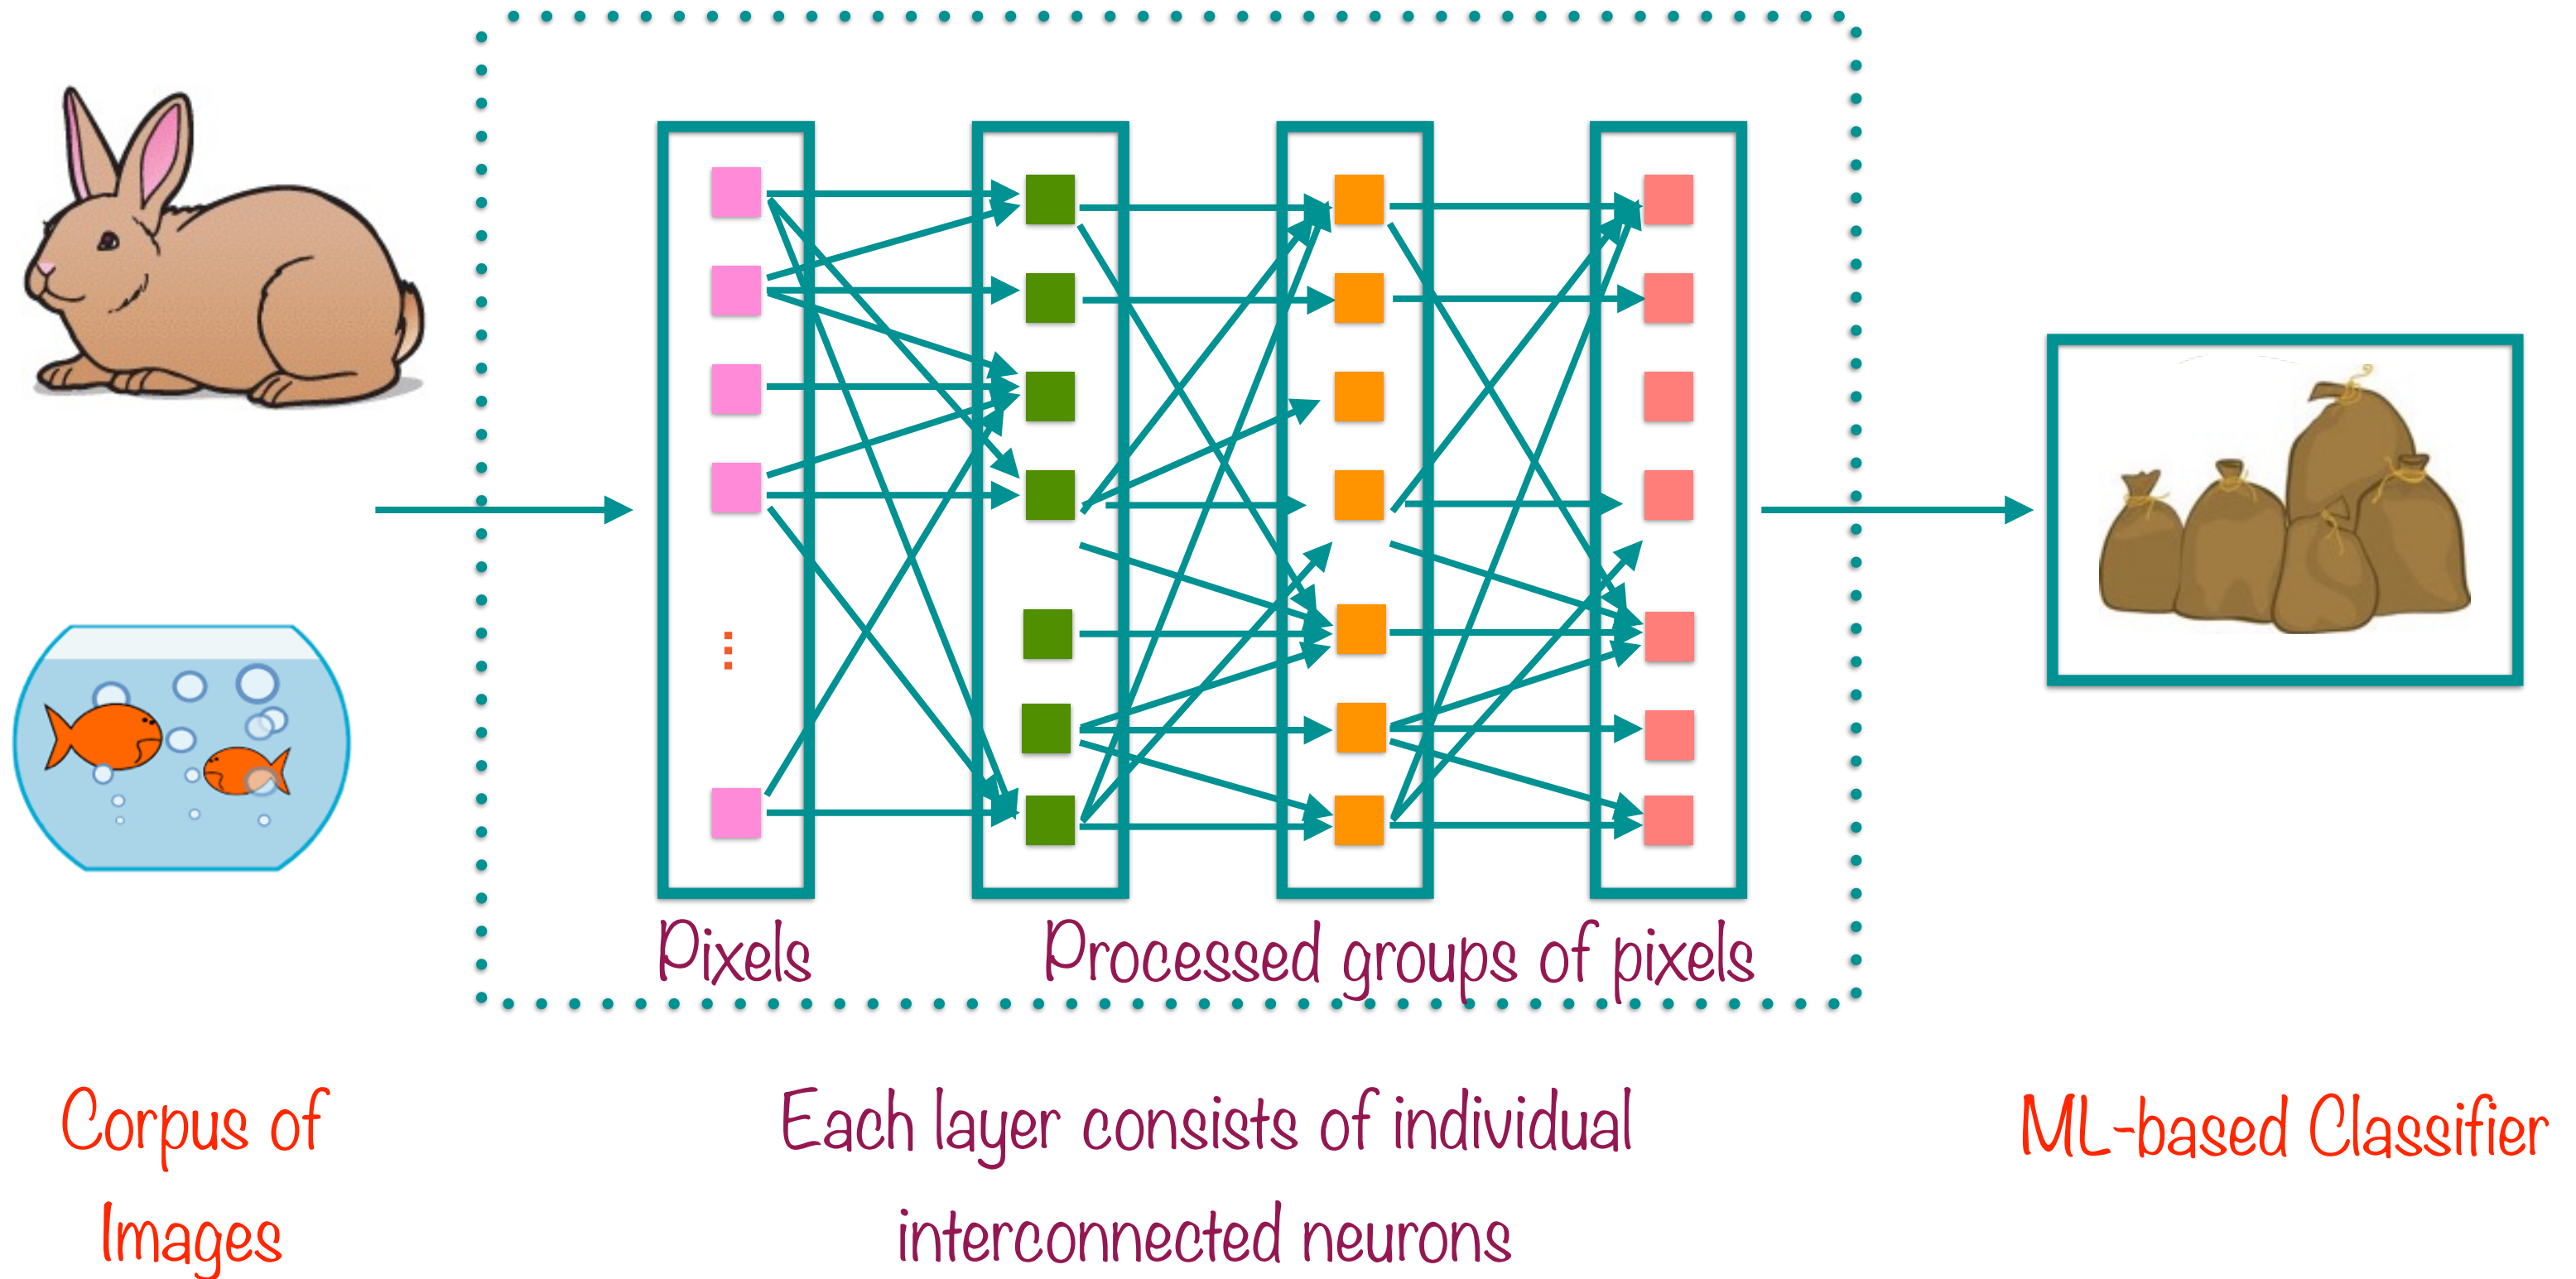
“Deep Learning”-based Binary Classifier



Neural Networks Introduced



Neural Networks Introduced



Neural networks help find unknown patterns in massive data sets

TensorFlow for Machine Learning

TensorFlow

TensorFlow™ is an open source software library for numerical computation using data flow graphs.

<https://www.tensorflow.org/>



TensorFlow

TensorFlow™ is an open source software library for numerical computation using data flow graphs.


<https://www.tensorflow.org/>



TensorFlow

TensorFlow™ is an open source software library for numerical computation using data flow graphs.

<https://www.tensorflow.org/>




TensorFlow

TensorFlow™ is an open source software library for numerical computation using data flow graphs.

<https://www.tensorflow.org/>

Advantages of TensorFlow



Distributed

Runs on a cluster or machines
or multiple CPUs/GPUs on the
same machine



Suite of software

TensorFlow, TensorBoard,
TensorFlow Serving

TensorFlow



Uses



Strengths



Challenges



Uses

Research and development of new ML algorithms

Taking models from training to production

Large scale distributed models

Models for mobile and embedded systems



Strengths

Easy to use, stable Python API

Runs on large as well small systems

Efficient and performant

Great support from Google

Additional tools like TensorBoard and TensorFlow serving

Challenges



Distributed support still has a ways to go


Libraries still being developed

Writing custom code is not straightforward

TensorFlow is on its way to becoming the
default library for machine learning


The TensorFlow World

Everything Is a Graph




A network

Everything Is a Graph




Computations

Everything Is a Graph




Data
Tensors

Tensors Flow Through the Graph



...and get transformed along the way

Tensors Flow Through the Graph



TensorFlow

Demo

Download and install TensorFlow on your local machine

Validate that the TensorFlow libraries work and can be referenced

Summary

Learnt the basics of machine learning, deep learning and neural networks

Understood the strengths and challenges of using TensorFlow for ML

Understood the modeling of problem as a computational graph

Got TensorFlow up and running on your local machine

Introducing Computation Graphs

Overview


Model nodes, edges and dependencies in a computation graph

Understand the basic parts of a program in TensorFlow

Run TensorFlow programs and visualize results using TensorBoard


The TensorFlow World

Everything Is a Graph




A network

Everything Is a Graph




Computations

Everything Is a Graph




Tensor
Data

Tensors Flow Through the Graph




...and gets transformed along the way

Tensors Flow Through the Graph




TensorFlow

Tensors Flow Through the Graph




TensorFlow

Tensors Flow Through the Graph




$$Y = (\text{round}(A) + \text{floor}(B)) * \text{round}(A) + \text{abs}(\text{round}(A) + \text{floor}(B))$$


Directed-Acyclic Graph



Directed-Acyclic Graph




Directed-Acyclic Graph




Edges point forward towards a result i.e. directed

Dependencies



One node can send its output to multiple nodes

Dependencies




Or receive inputs from multiple nodes


Dependencies

Direct dependency

Indirect dependency




Directed-Acyclic Graph



There are no cycles in the graph i.e. *acyclic*

Directed-Acyclic Graph




A graph with cycles will never finish computation

**Problems in TensorFlow are represented as
a directed-acyclic graph**

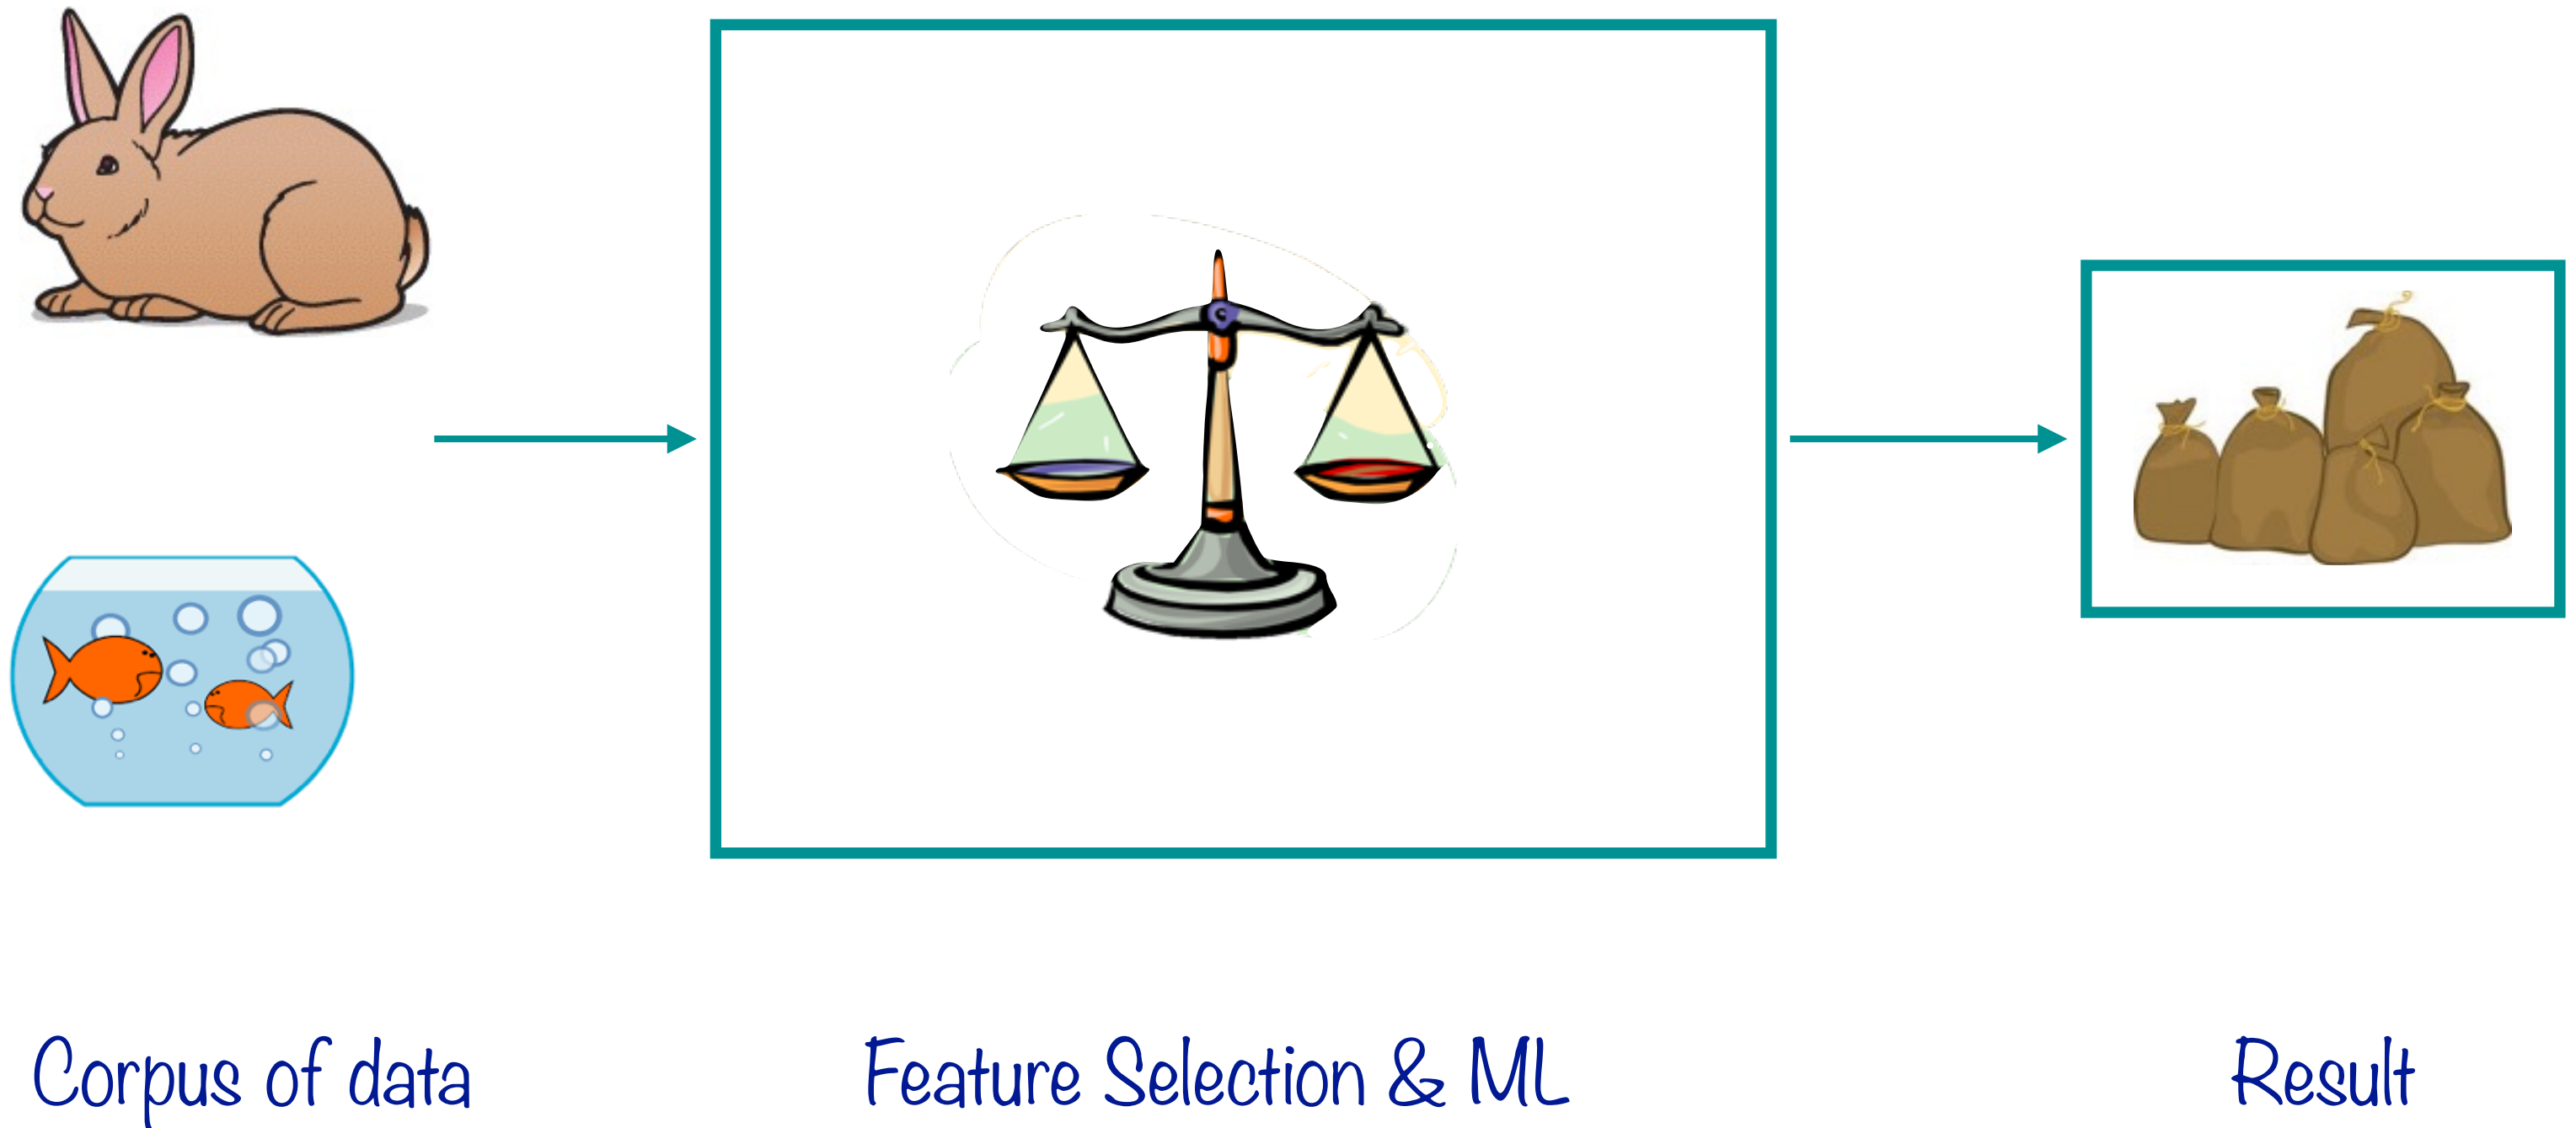
Cyclical Dependencies in Machine Learning

Directed-Acyclic Graph

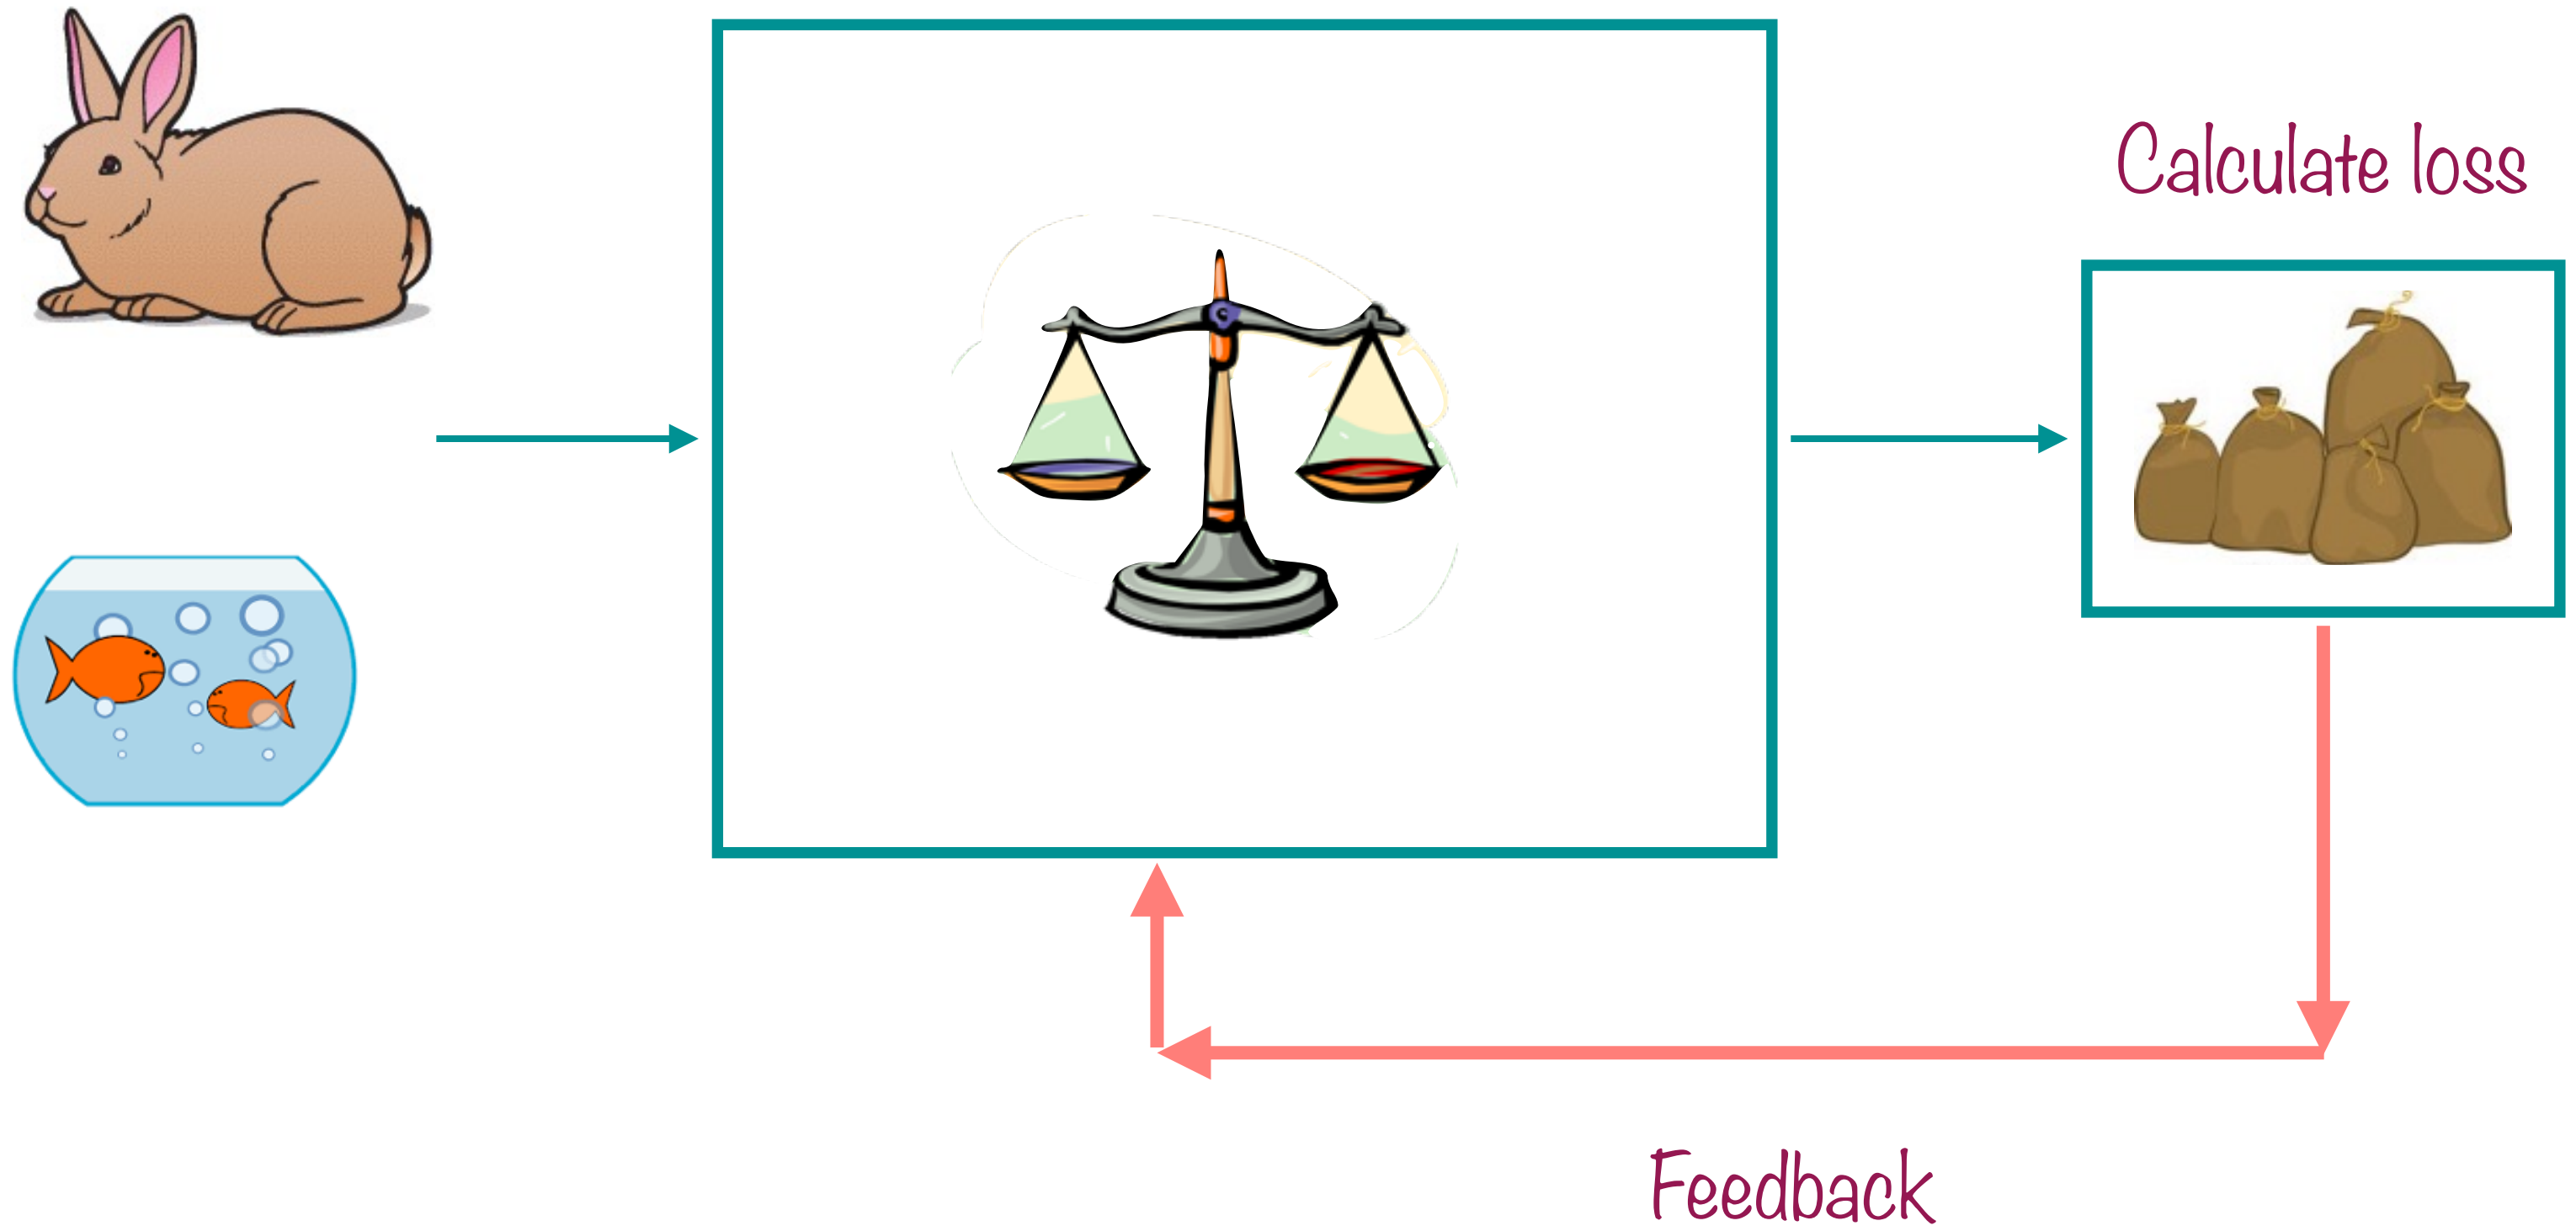


A graph with cycles will never finish computation


The Process of Machine Learning




The Process of Machine Learning




The Process of Machine Learning



Modelling Cyclical Dependencies




Modelling Cyclical Dependencies




Unrolling the graph

Modelling Cyclical Dependencies



Unrolling the graph

Modelling Cyclical Dependencies



How much you unroll depends on the number of iterations you want to run

Unroll graphs to model cyclic dependencies

Building and Running Graphs

2 Steps in a TensorFlow Program



Building a Graph

Specify the operations and the
data



Running a Graph

Execute the graph to get the final
result

Demo

Building and running a graph in TensorFlow

Exploring the graph using TensorBoard

2 Steps in a TensorFlow Program



Building a Graph

Specify the operations and the
data



Running a Graph

Execute the graph to get the final result

2 Steps in a TensorFlow Program



Building a Graph


Specify the operations and the data



Running a Graph

Execute the graph to get the
final result

Visualizing a Graph




TensorBoard


Visualize how data flows and what
computations operate on it

Modeling Computations as Graphs


Computation Graphs




Computation Graphs




Computation Graphs




Computation Graphs




Computation Graphs



Computation Graphs




Computation Graphs




**TensorFlow calculates only that portion of
the graph which is required**

Computation Graphs



Running Graphs on a Distributed System



**Multiple portions of the graph can be run
in parallel across machines in the cluster**

Demo

Executing simple math commands in TensorFlow

Tensor

Tensor

The central unit of data in TensorFlow. A tensor consists of a set of primitive values shaped into an array of any number of dimensions.

<https://www.tensorflow.org/>

Tensor

The central unit of data in TensorFlow. *A tensor consists of a set of primitive values shaped into an array of any number of dimensions.*


<https://www.tensorflow.org/>

Tensor

The central unit of data in TensorFlow. A tensor consists of a set of primitive values shaped into an array of any number of dimensions.

<https://www.tensorflow.org/>

Data Is Represented as Tensors



Scalars are 0-D tensors

3, 6.7, "a"

Data Is Represented as Tensors



Vectors are 1-D tensors


○ [1, 3, 5, 7, 9] ○

Data Is Represented as Tensors


Matrices are 2-D tensors

[[1, 3, 5],
[7, 9, 11]]


Data Is Represented as Tensors



N-Dimensional matrices are N-D tensors



$[[[1, 2], [3, 4], [5, 6]],$
 $[[7, 8], [9, 10], [11, 12]]]$



Characterization of Tensors



Rank

The number of dimensions
in a tensor



Shape

The number of elements in
each dimension




Data Type

The data type of each
element in the tensor



Rank

Tensor	Rank
4	0
[1, 2, 3]	1
[[1, 2], [3, 4]]	2
[[[1], [2]], [[3], [4]]]	3



Shape

Tensor	Shape
4	<code>[]</code>
<code>[1, 2, 3]</code>	<code>[3]</code>
<code>[[1, 2, 3], [4, 5, 6]]</code>	<code>[3, 2]</code>
<code>[[[1], [2]], [[3], [4]]]</code>	<code>[2, 2, 1]</code>



Data Type

int

float

string

boolean

**Rank, shape and data types are 3
important characteristics which define a
Tensor**

Summary

Worked with the directed-acyclic graph to model problems in TensorFlow

Understood constants, operators and sessions

Understood Tensor characteristics such as rank, shape and data type

Run TensorFlow programs and visualized results using TensorBoard

Digging Deeper into Fundamentals

Overview

Work through the logic of a specific machine learning problem

Run programs on different inputs using placeholders and feed dictionaries

Use variables to hold values which the program updates

Make TensorBoards more useful using named scopes

Understanding Linear Regression

X Causes Y



Cause

Independent variable



Effect

Dependent variable

Wealth Increases Life Expectancy



Cause

Wealth of individuals



Effect

How long they expect to live

Lower Home Prices Away from the City



Cause

Distance in miles from the city
center



Effect

Price per square foot of homes

X Causes Y



Cause


Explanatory variable



Effect

Dependent variable

Linear Regression



Simple Regression


Regression Equation:

$$y = A + Bx$$

■	y_1	=	$A + Bx_1$
■	y_2	=	$A + Bx_2$
■	y_3	=	$A + Bx_3$

■	y_n	=	$A + Bx_n$

The “Best” Regression Line




Linear Regression involves finding the “best fit” line

The "Best" Regression Line



y



Line 1: $y = A_1 + B_1x$


Line 2: $y = A_2 + B_2x$

x



Let's compare two lines, Line 1 and Line 2


The "Best" Regression Line



The "Best" Regression Line



y



x increases by 1

y decreases by B_1

Line 1: $y = A_1 + B_1x$


Line 2: $y = A_2 + B_2x$

x



In the first line, if x increases by 1 unit, y decreases by B_1 units

The "Best" Regression Line




The second line has y-intercept A_2



The "Best" Regression Line



y



Line 1: $y = A_1 + B_1x$

Line 2: $y = A_2 + B_2x$

x




In the second line, if x increases by 1 unit, y decreases by B_2 units

Minimising Least Square Error



y



Line 1: $y = A_1 + B_1x$

Line 2: $y = A_2 + B_2x$

x




Drop vertical lines from each point to the lines 1 and 2

Minimising Least Square Error



y



Line 1: $y = A_1 + B_1x$

Line 2: $y = A_2 + B_2x$

x




Drop vertical lines from each point to the lines 1 and 2

Minimising Least Square Error



y



Line 1: $y = A_1 + B_1x$

Line 2: $y = A_2 + B_2x$

x




The “best fit” line is the one where the sum of the squares of the lengths of these dotted lines is minimum

Minimising Least Square Error



y



Line 1: $y = A_1 + B_1x$

Line 2: $y = A_2 + B_2x$

x




The “best fit” line is the one where the sum of the squares of the lengths of these dotted lines is minimum

Minimising Least Square Error



y



Line 1: $y = A_1 + B_1x$

Line 2: $y = A_2 + B_2x$

x




The “best fit” line is the one where the sum of the squares of the lengths of the errors is minimum

Minimising Least Square Error



y



Line 1: $y = A_1 + B_1x$

Line 2: $y = A_2 + B_2x$

x




The “best fit” line is the one where the sum of the squares of the lengths of the errors is minimum

Minimising Least Square Error



y




Regression Line: $y = A + Bx$

x



Linear Regression



Linear Regression Algorithms in Practice

Estimate initial values for
 A and B


Find the errors for the
regression line with those
values

Feed errors back and get
new values for A and B

The "Best" Regression Line



y



$$y = A_{\text{start}} + B_{\text{start}}x$$

x




Start off with some values for A and B

The "Best" Regression Line



y



$$y = A_{\text{start}} + B_{\text{start}}x$$

x




Calculate the least square error and feed that back

The "Best" Regression Line



y



$$y = A_i + B_i x$$

x




This will give us new values for A and B

The "Best" Regression Line



y



$$y = A_i + B_i x$$

x




Adjust values of A and B by feeding back the error values

The "Best" Regression Line



y




$$y = A_i + B_i x$$

x



Adjust values of A and B by feeding back the error values

The “Best” Regression Line




The “best fit” line is called the regression line


Regression is an example of a supervised learning algorithm

Placeholders

Linear Regression



Linear Regression




Linear Regression



y

Life expectancy



Regression Line: $y = A + Bx$

x



Wealth

**Machine learning algorithms can
be applied to a variety of problems**

**The model should have the ability to accept the
different X and Y values**

Placeholder

Hold the place for a Tensor that will be fed at runtime, in effect becoming an “input” node

Abrahams, Sam; Hafner, Danijar; Erwitte, Erik; Scarpinelli, Ariel (2016-07-23). TensorFlow For Machine Intelligence: A hands-on introduction to learning algorithms

Demo

Specify placeholders in our simple math operations

Use a feed dictionary to feed these into TensorFlow operations

Demo


Work with fetches and the feed_dict passed to Session.run()
operations

Variables

The "Best" Regression Line



y



$$y = A_{\text{start}} + B_{\text{start}}x$$

x




Start off with some values for A and B

The "Best" Regression Line



y



$$y = A_{\text{start}} + B_{\text{start}}x$$

x




Calculate the least square error and feed that back

The "Best" Regression Line



y



$$y = A_i + B_i x$$

x




This will give us new values for A and B

The "Best" Regression Line



y



$$y = A_i + B_i x$$

x




Adjust values of A and B by feeding back the error values

The "Best" Regression Line



y




$$y = A_i + B_i x$$

x



Adjust values of A and B by feeding back the error values

The “Best” Regression Line




The “best fit” line is called the regression line

The "Best" Regression Line



y



Regression Line: $y = A + Bx$

x



The values of A and B are updated till we get the best fit line

**Machine learning algorithms
iterate to get closer to the solution**

**The model should have the ability to hold
constantly changing values**

Constants, Placeholders and Variables

Constants

Immutable values which do not change

Placeholders

Assigned once and do not change after

Variables

Are constantly recomputed

Variables

Mutable Tensor values that persist across multiple calls to `Session.run()`

Abrahams, Sam; Hafner, Danijar; Erwit, Erik; Scarpinelli, Ariel (2016-07-23). TensorFlow For Machine Intelligence: A hands-on introduction to learning algorithms

Summary

Implement placeholders and variables in TensorFlow

Make TensorBoards more useful using named scopes

Working with Images


Overview

Representing color and grayscale images as Tensors

Implementing image operations such as transpose, resize, cropping


Image Recognition

Pixels in Images



Copyright: Wikimedia commons, GFDL and cc-by-sa-2.5,2.0,1.0

Pixels in Images



Copyright: Wikimedia commons, GFDL and cc-by-sa-2.5,2.0,1.0

Image Recognition



Images represented as pixels



Identify edges, colors,
shapes



A photo of a horse

Neural networks, specifically convolutional
neural networks (CNNs) work well for
hard image recognition tasks

Image Recognition Using Neural Networks

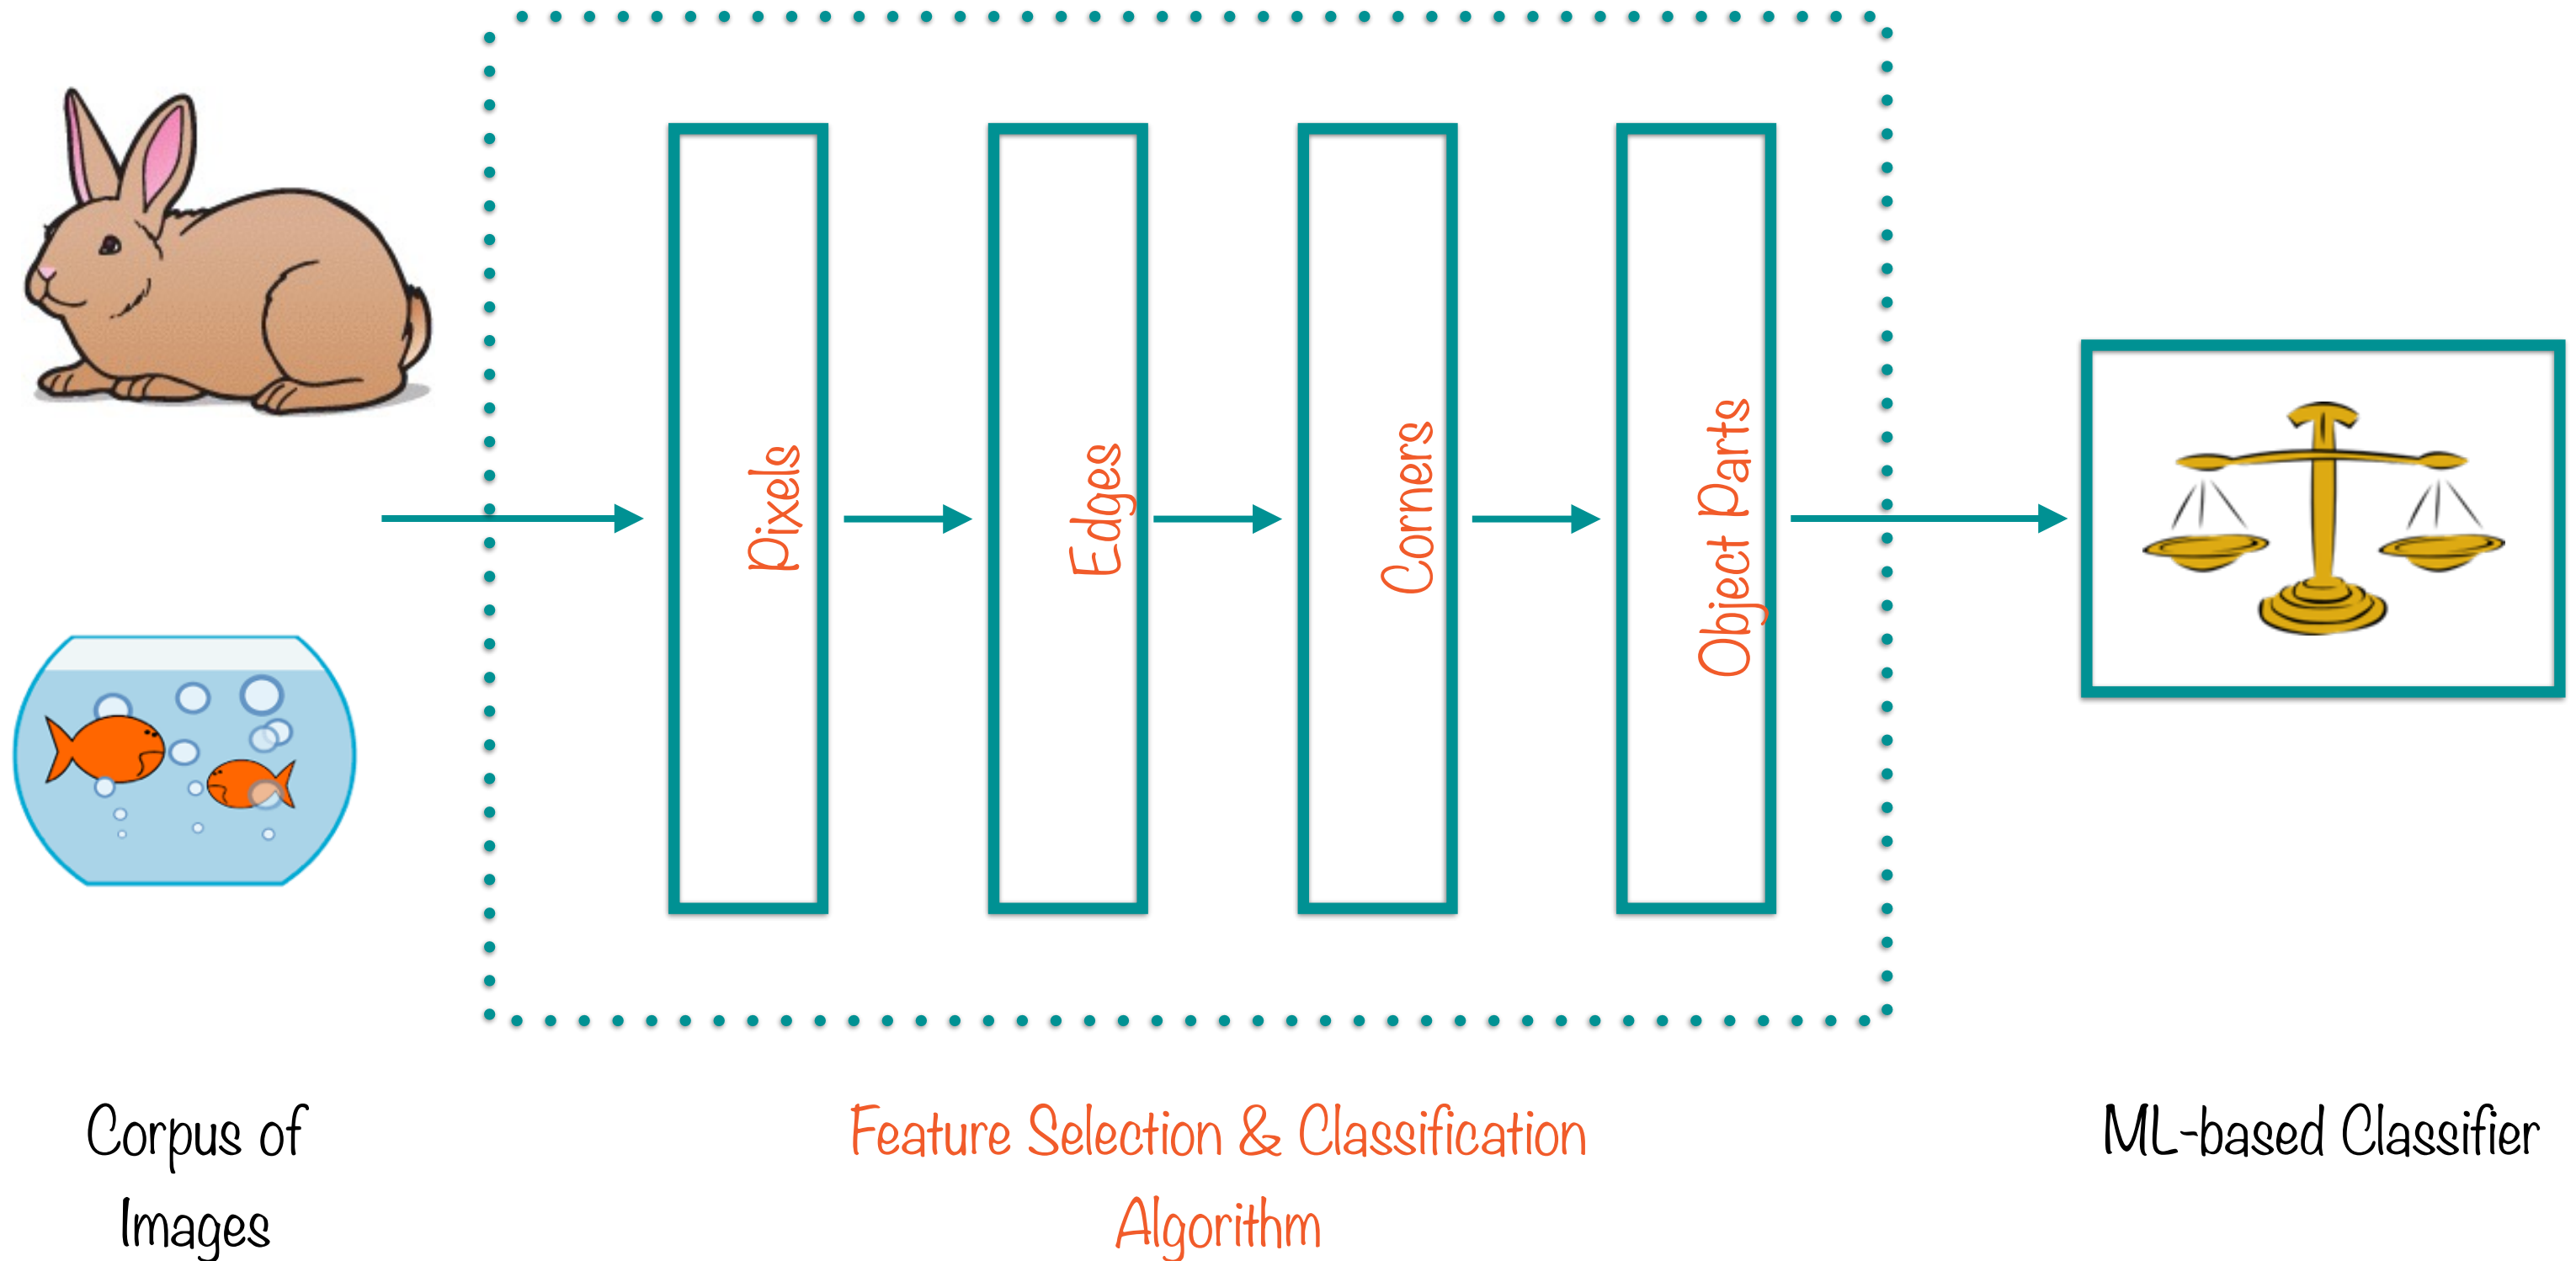


Image Recognition Using Neural Networks

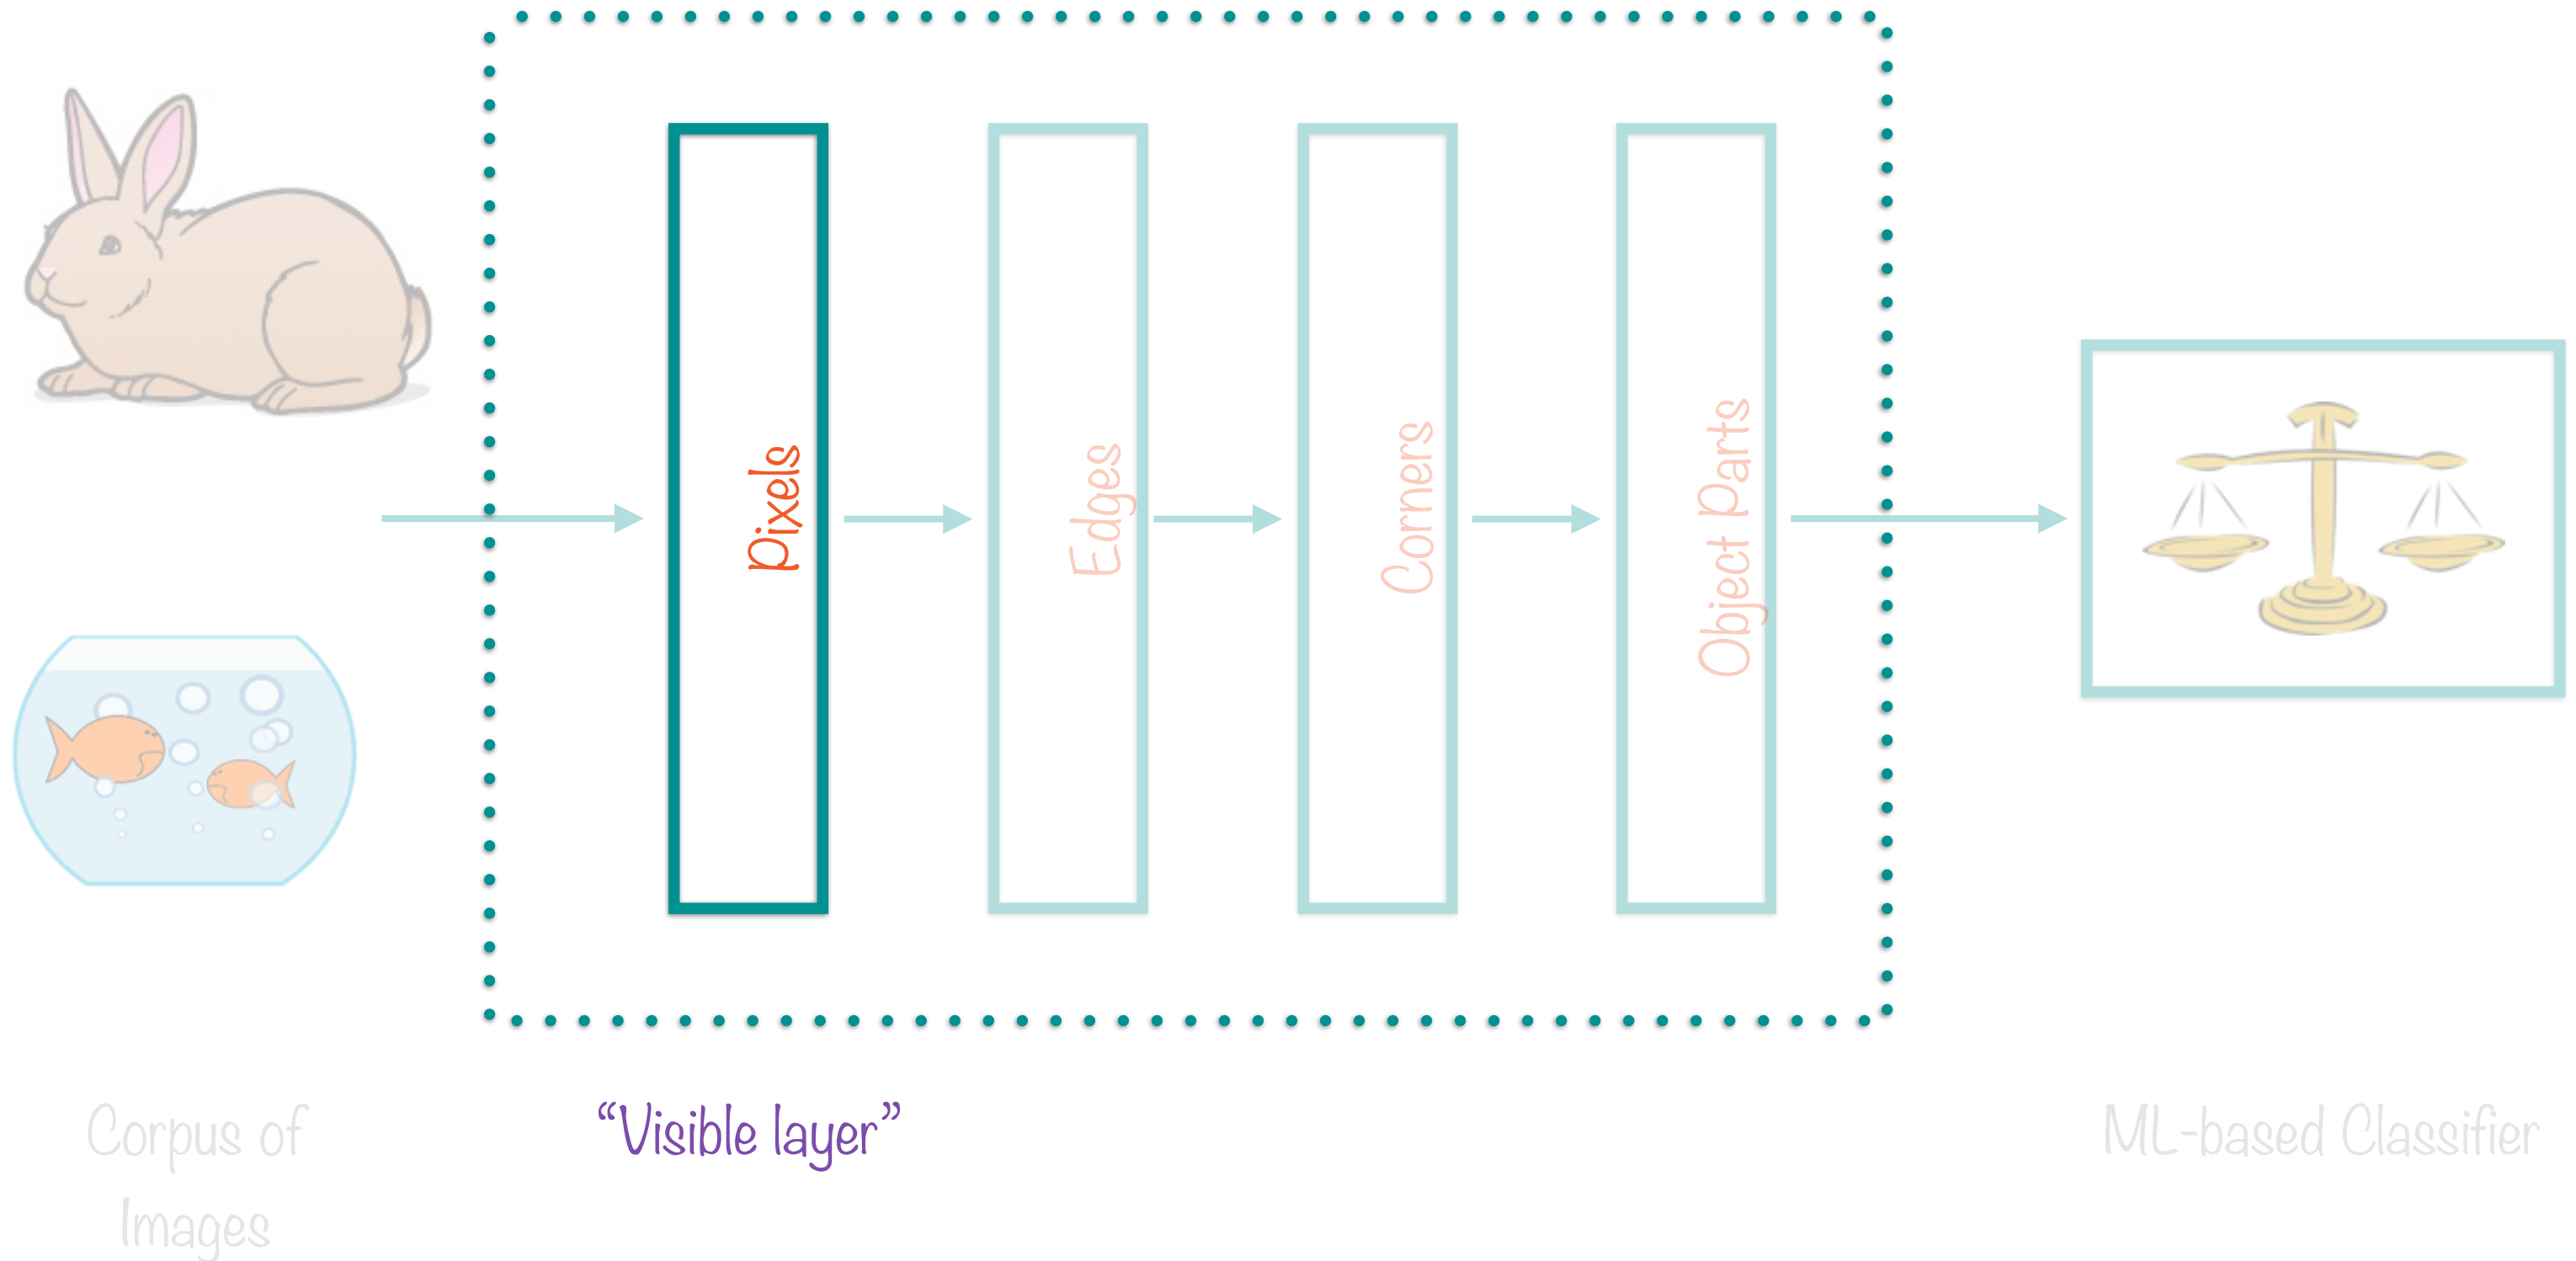

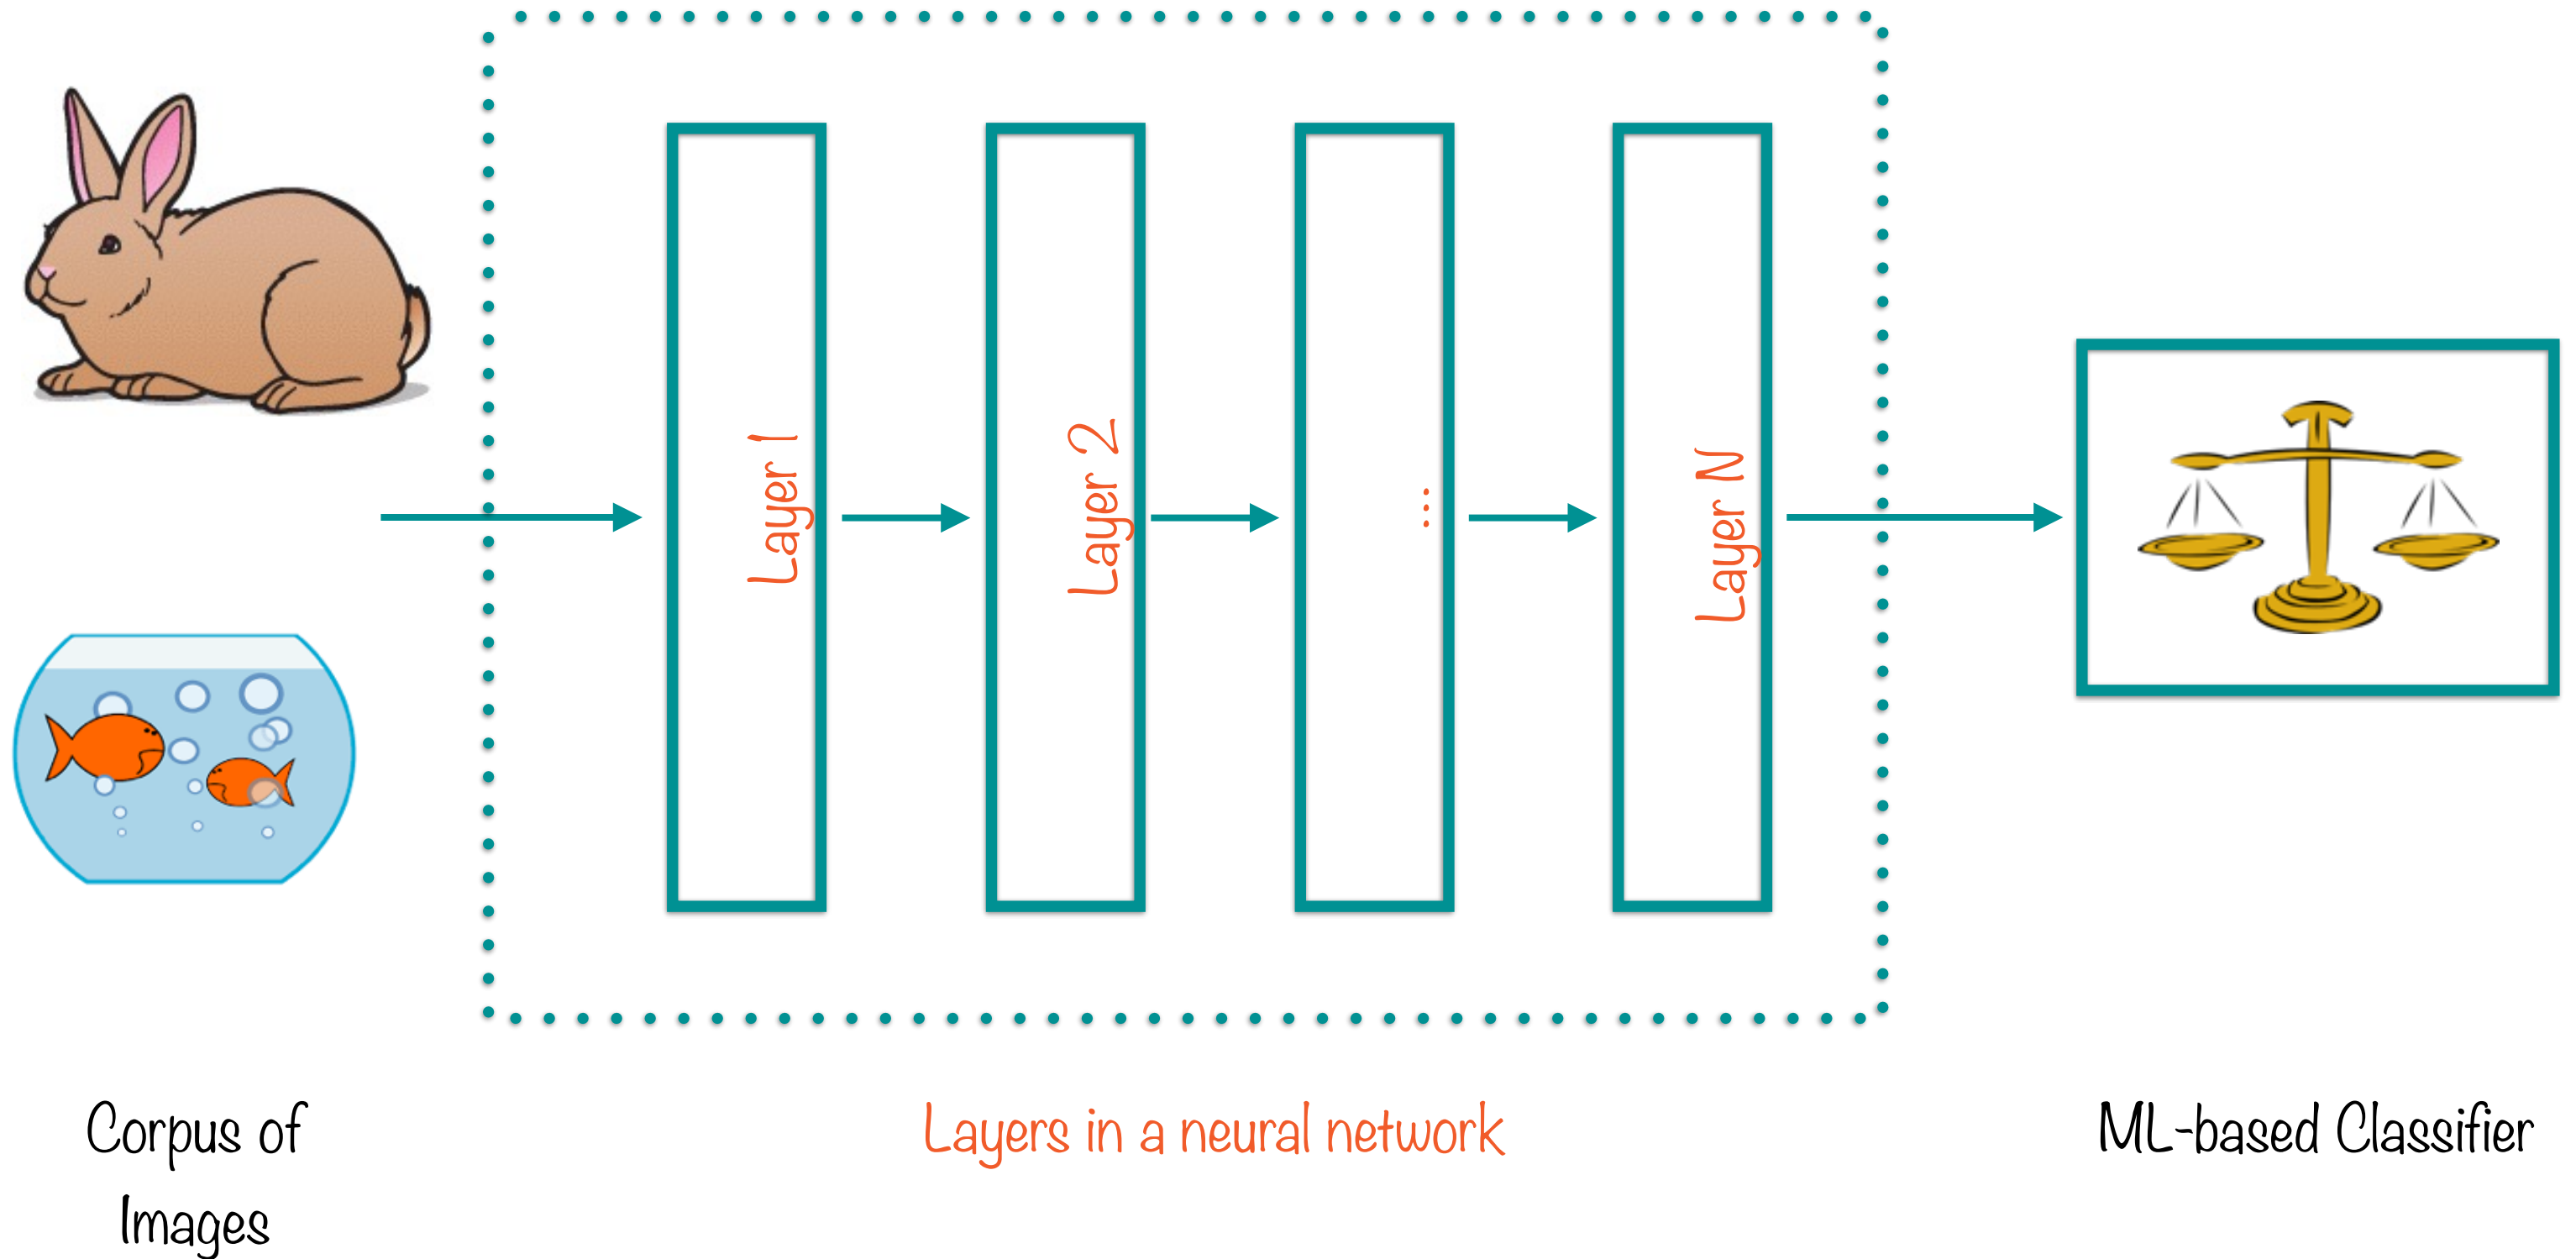


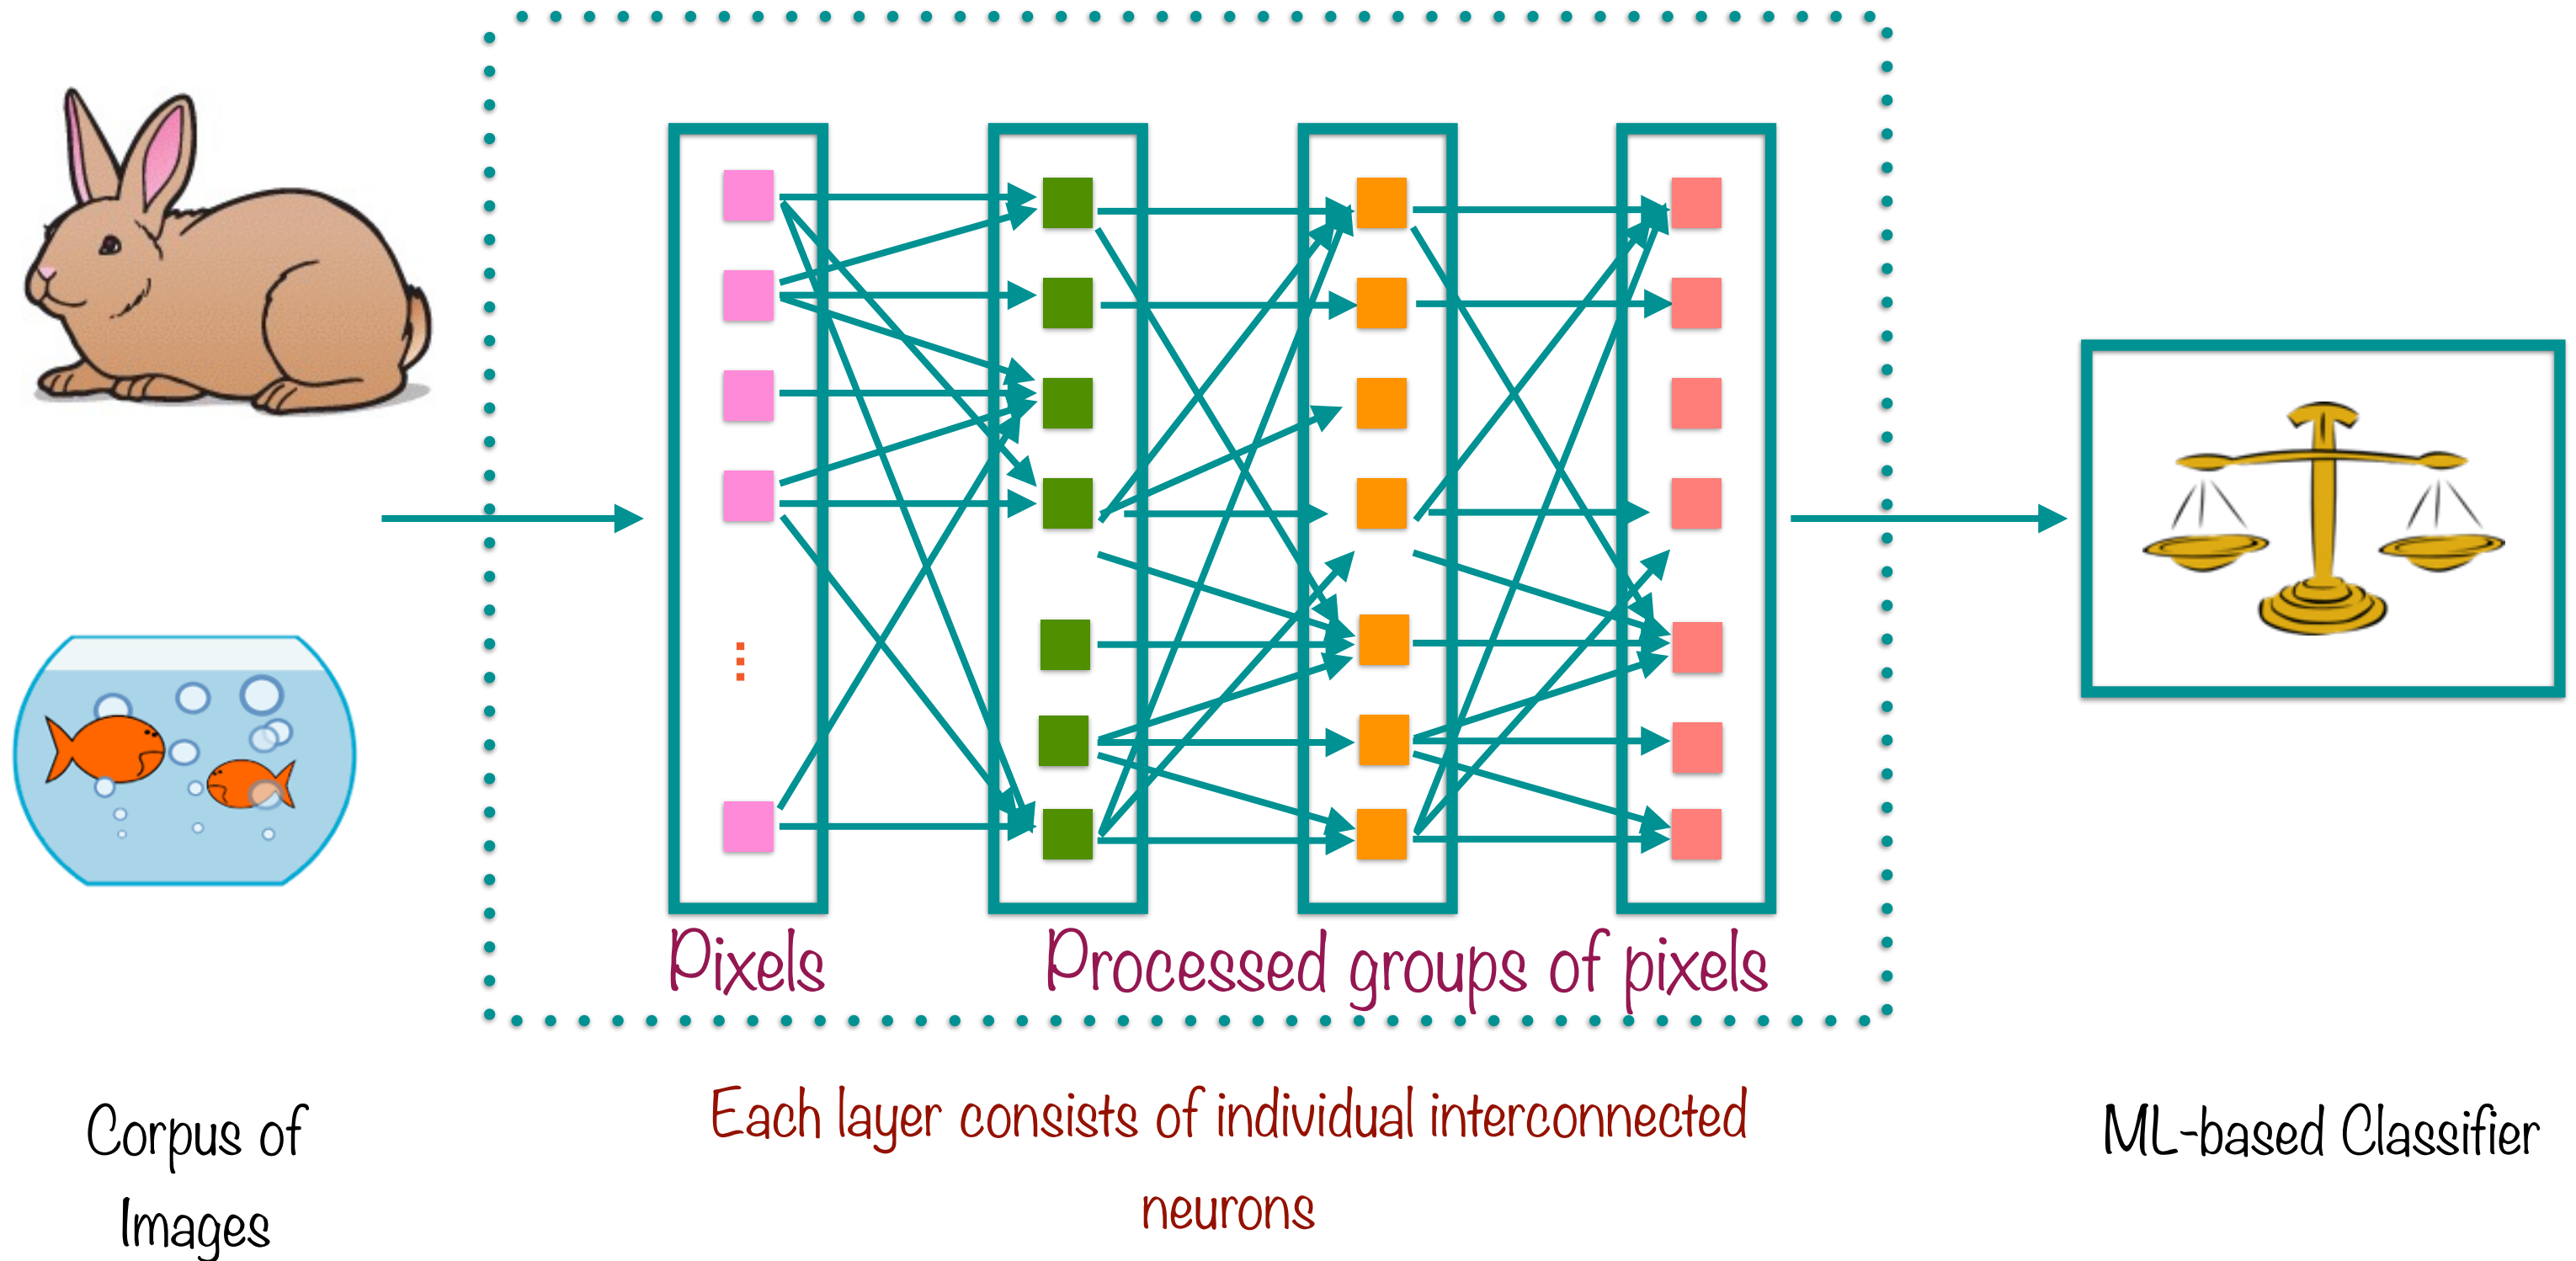
Image Recognition Using Neural Networks



Neural Networks Introduced




Neural Networks Introduced





TensorFlow is optimized at building
neural network solutions for image
recognition

Representing Images as 3-D Tensors

Images as Tensors




Images as Tensors



Each pixel holds a value based on the type of image



RGB Images




RGB values are for color images

R, G, B: 0-255




RGB Images



255, 0, 0




RGB Images



0, 255, 0




RGB Images




0, 0, 255


3 values to represent color, 3
channels

Grayscale Images





Grayscale Images




Each pixel represents only intensity
information

0.0 - 1.0




Grayscale Images



0.5





Grayscale Images



0.5




| value to represent intensity, |
channel

Images as Tensors







Single channel and multi-channel images

Images as Tensors





Images can be represented by a 3-D matrix

Images as Tensors





The **number of channels** specifies the **number of elements** in the 3rd dimension

Images as Tensors



(6, 6, 1)



(6, 6, 3)

Demo

Read in an image using matplotlib and then
transpose it using TensorFlow

Demo

Read in a list of images in TensorFlow using a queue and coordinators

Resize images to be of the same dimensions

Show image summaries in TensorBoard

List of Images as 4-D Tensors

List of Images



TensorFlow usually deals with a **list of images in one 4-D Tensor**

List of Images



The images should all be the same size



List of Images

(10, 6, 6, 3)

The number of channels



List of Images

(10, 6, 6, 3)

The height and width of each image in the list



List of Images

(10, 6, 6, 3)

The number of images

Summary

Understood image representation of color and grayscale images as Tensors

Learnt image transformations such as resize, flip and crop

Worked with multiple images in TensorFlow

Using K-nearest-neighbors for Digit Recognition

Overview

Introduce the MNIST handwritten digit dataset

Understand the K-nearest-neighbors machine learning algorithm

Implement K-nearest-neighbors in TensorFlow to identify handwritten digits from 0 to 9

The MNIST Handwritten Digits Dataset

MNIST Dataset



Handwritten digits database

Large quantity of handwritten digits commonly
used for training image processing systems

MNIST Dataset



Handwritten digits database


Modified National Institute of Standards and
Technology

MNIST Dataset



Each digit is in grayscale


MNIST Dataset



Every image is standardized
to be of size 28×28


$= 784$ pixels

MNIST Dataset



Every pixel holds a **single** value for intensity

MNIST Dataset




0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



Every image has an associated label

MNIST Dataset



5



0



4



1

**MNIST for machine learning is the
equivalent of the “Hello World” for
programming**

The K-nearest-neighbors Algorithm

Types of ML Algorithms



Supervised

Labels associated with the training data
is used to correct the algorithm



Unsupervised

The model has to be set up right to learn
structure in the data

Supervised Learning

Input variable x and output variable y

Learn the mapping function $y = f(x)$

Approximate the mapping function so for new values of x we can predict y

Use existing dataset to **correct** our mapping function approximation





Unsupervised Learning


Only have input data **x** no output data

Model the underlying structure to learn more about data


Algorithms **self discover** the patterns and structure in the data

Training Data

Breathes like a mammal
Gives birth like a mammal




Mammal




**KNN is an supervised learning algorithm
which uses training data to find what is
most similar to the current sample**

K-nearest-neighbors



Uses the entire training dataset as a model

K-nearest-neighbors



Rocket



Buildings



Signal



Pig



Shop

Each element in training data has a **label**

K-nearest-neighbors



Rocket



Buildings



Signal



Pig




Shop



Predictions for a new sample involves figuring out which element in the training data it is **similar** to

The nearest neighbor

K-nearest-neighbors



Rocket



Buildings



Signal




Pig



Shop



K-nearest-neighbors



Rocket



Buildings



Signal




Pig



Shop



K-nearest-neighbors



Rocket



Buildings



Signal




Pig



Shop



K-nearest-neighbors



Rocket



Buildings



Signal




Pig



Shop



K-nearest-neighbors



Rocket



Buildings



Signal



Pig



Shop




K-nearest-neighbors




How do we calculate neighbors of a sample?

K-nearest-neighbors




Distance measures

K-nearest-neighbors




Euclidean distance, Hamming distance, Manhattan distance


K-nearest-neighbors




K-nearest-neighbors



K-nearest-neighbors




Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \text{sqrt}(\text{sum}((x_j - x_{ij})^2))$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\text{sum}((x_j - x_{ij})^2)}$$

Distance Measures



$$\text{EuclideanDistance}(x, x_i) = \text{sqrt}(\text{sum}((x_j - x_{ij})^2))$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$$

Distance Measures

Distance Measure

Euclidean Distance


As the crow flies




L1 Distance

Distance Measure

L1 distance
Snake distance
City block distance
Manhattan distance




L1 Distance



$$5 - 1 = 4$$


L1 Distance



$$5-1=4$$

$$4-0=4$$

L1 Distance



$$5 - 1 = 4$$
$$4 - 0 = 4$$
$$= 8$$

Demo

Handwritten image recognition using the k-nearest-neighbors ML algorithm

- Use the L1 distance measure to find the nearest neighbor
- Measure the accuracy of the algorithm on the test data

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches using the TensorFlow libraries

Running the algorithm

Predict labels for all the test data and measure accuracy

Calculating L1 distance

Find the distance between the test digit and all training digits

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches using the TensorFlow libraries



MNIST Dataset



5



0




4



1

MNIST Dataset



4

Every image is standardized
to be of size 28×28

$= 784$ pixels

Representing Labels



Vector

0	0	0	0	1	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---

Index

0	1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---

Representing Labels



Vector


0	0	0	0	0	1	0	0	0	0
---	---	---	---	---	---	---	---	---	---

Index

0	1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---

5

Representing Images



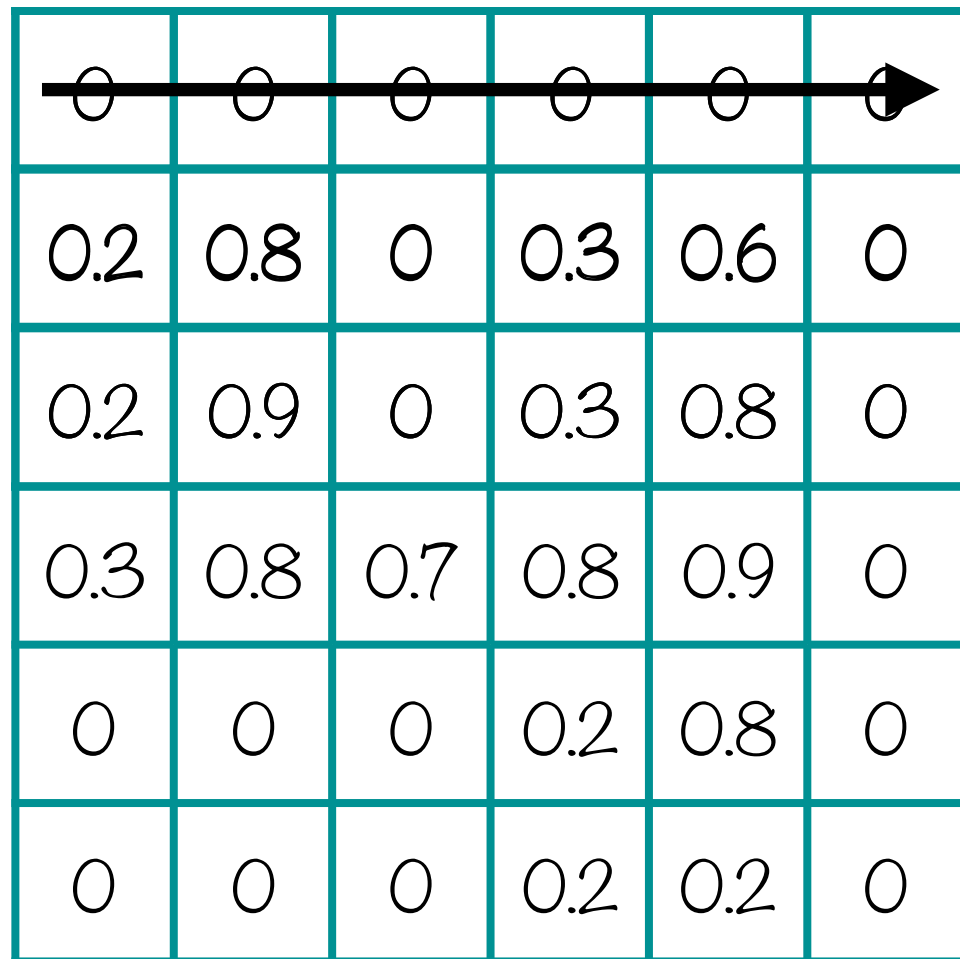
= 784 pixels

Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

= 784 pixels

Representing Images



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0
---	---	---	---	---	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0.2	0.8	0	0.3	0.6	0
---	---	---	---	---	---	-----	-----	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	→
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0.2	0.9	0	0.3	0.8	0	0.2	0.8	0	0.3	0.6	0
---	---	---	---	---	---	-----	-----	---	-----	-----	---	-----	-----	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0	0	0	0.2	0.2	0
---	---	---	---	---	---	----	----	----	----	----	----	---	---	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0	0	0	0.2	0.2	0
---	---	---	---	---	---	----	----	----	----	----	----	---	---	---	-----	-----	---


0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0	0	0	0.2	0.2	0
---	---	---	---	---	---	----	----	----	----	----	----	---	---	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images



= 784 pixels

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches using the TensorFlow libraries

Calculating L1 distance

Find the distance between the test digit and all training digits

L1 Distance

Training

0.2	0.8	0	0.3	0.6	0
-----	-----	---	-----	-----	---

Test

0	0.6	0.2	0.3	0.3	0.1
---	-----	-----	-----	-----	-----

tf.negative()

0.2	0.8	0	0.3	0.6	0
-----	-----	---	-----	-----	---

0	-0.6	-0.2	-0.3	-0.3	-0.1
---	------	------	------	------	------

tf.add()

0.2	0.2	-0.2	0	0.3	-0.1
-----	-----	------	---	-----	------

L1 Distance

tf.add()

0.2	0.2	-0.2	0	0.3	-0.1
-----	-----	------	---	-----	------

tf.abs()

0.2	0.2	0.2	0	0.3	0.1
-----	-----	-----	---	-----	-----

tf.reduce sum()

1.0

tf.add()

0	0.2	0.3	0.6	-0.3	-0.1
---	-----	-----	-----	------	------

tf.abs()

0	0.2	0.3	0.6	0.3	0.1
---	-----	-----	-----	-----	-----

tf.reduce sum()

1.5

tf.add()

-0.2	0.4	-0.2	0	-0.3	-0.1
------	-----	------	---	------	------

tf.abs()

0.2	0.4	0.2	0	0.3	0.1
-----	-----	-----	---	-----	-----

tf.reduce sum()

1.2

L1 Distance

tf.reduce_sum()

1

tf.reduce_sum()

1.5

tf.reduce_sum()

1.2

index = 0

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches


Running the algorithm

Predict labels for all the test data and measure accuracy

Calculating L1 distance

Find the distance between the test digit and all training digits

Representing Labels



`np.argmax()`

Summary

Familiar with the MNIST handwritten digit dataset

Understood the logic behind the K-nearest-neighbors algorithm

Implemented K-nearest-neighbors using L1 distance to identify handwritten digits from 0 to 9